


Deep Learning from Scratch

Session #6: Recurrent Neural Networks



by: Ali Tourani – Summer 2021

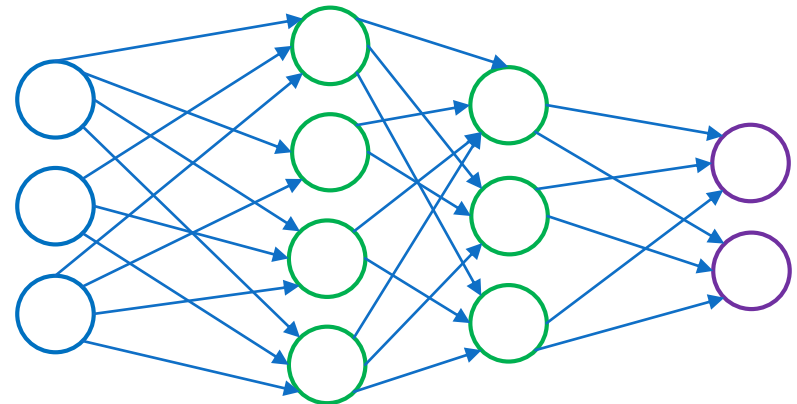
Agenda

- ▶ Sequence Models
- ▶ Recurrent Neural Networks (RNNs)
- ▶ Backpropagation Through Time (BPTT)
- ▶ Long Short-Term Memory (LSTM)
- ▶ Gated Recurrent Unit (GRU)

Sequence Models

Remember our simple ANNs?!

- ▶ Let's call these architectures **Feedforward NNs**
 - ▶ The information is only passed in one direction
 - ▶ The connections between nodes do not form a cycle
 - ▶ **Single-layer Perceptron**
 - ▶ Only comparing the outputs with actual values
 - ▶ Gradient Descent
 - ▶ **Multi-layered Perceptron**
 - ▶ Backpropagation

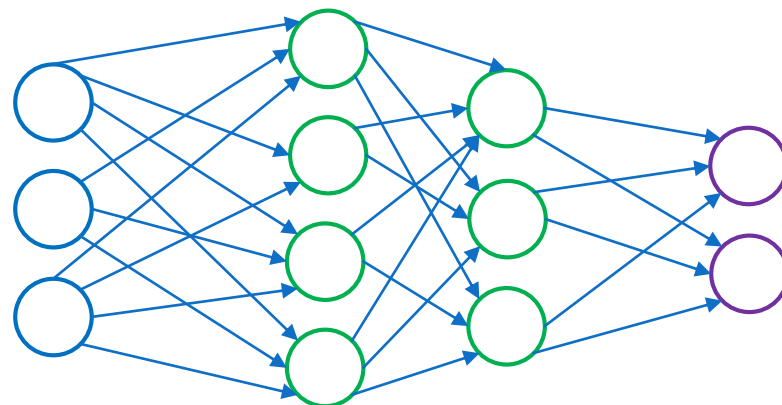
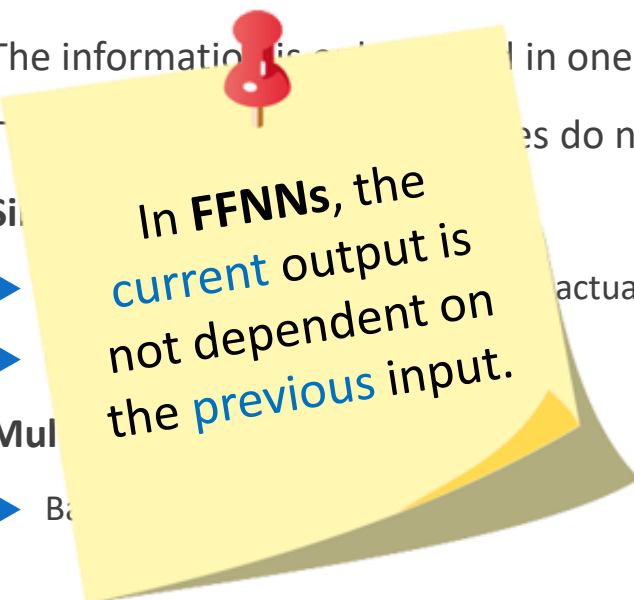


Sequence Models

Remember our simple ANNs?!

▶ Let's call these architectures **Feedforward NNs**

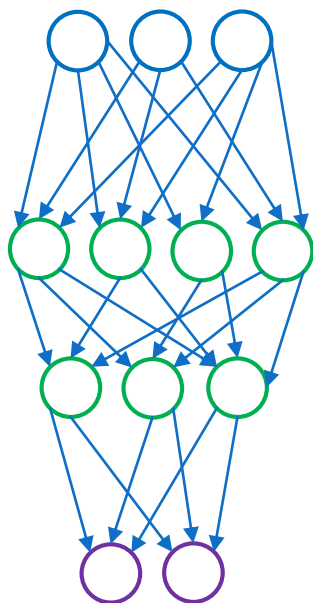
- ▶ The information is sent in one direction
- ▶ The layers do not form a cycle
- ▶ Since the current output is not dependent on the previous input.
- ▶ Multiple layers of hidden units
- ▶ Backpropagation



Sequence Models



Note that ...

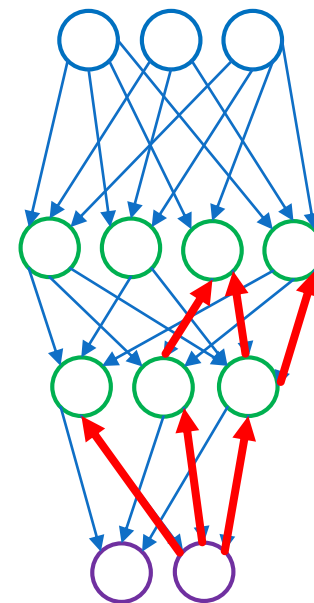


Feed-Forward NN

- An **architecture**
- Data travels from the input layer to hidden layers and finally, the output layer

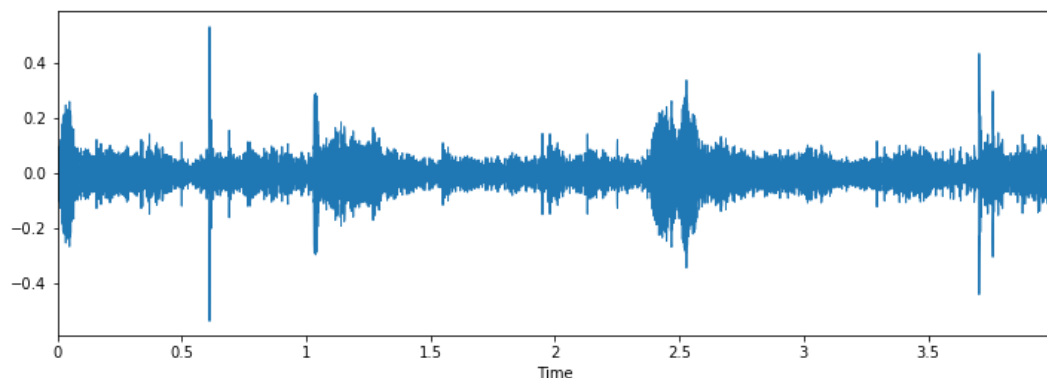
Backpropagation

- A **training algorithm**
- Forwards the values, and only after error calculation, propagates them back to the earlier layers



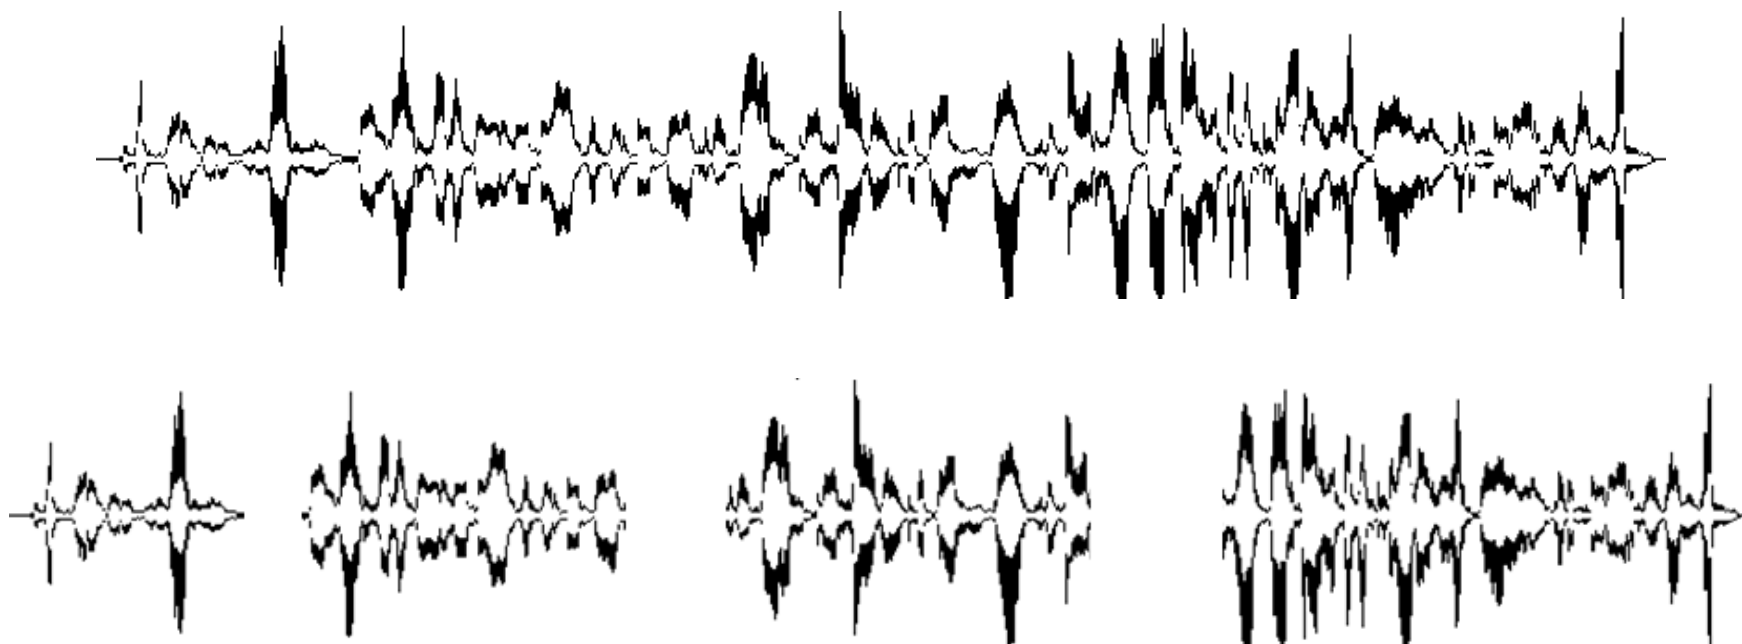
Sequence Models

- ▶ **Sequence Modeling is:**
 - ▶ Predicting **what comes next**, according to the **previous information**
 - ▶ We need to know the prior state of the model to predict its next ones
- ▶ Machine learning models with sequences of data as input/output
- ▶ Samples of sequential data:
 - ▶ Text streams
 - ▶ Audio
 - ▶ Videos
 - ▶ Time-series data



Sequence Models

- ▶ Sequential data can be split into sub-samples - **Audio**



Sequence Models

- ▶ Sequential data can be split into sub-samples - **Video**



Sequence Models

What else about Sequence Modeling?

- ▶ The current output **is not independent** on the previous input
 - ▶ **In contrast with Feedforward Neural Networks (FNNs)**
- ▶ The length of the input is not fixed
 - ▶ We might have a stream of data

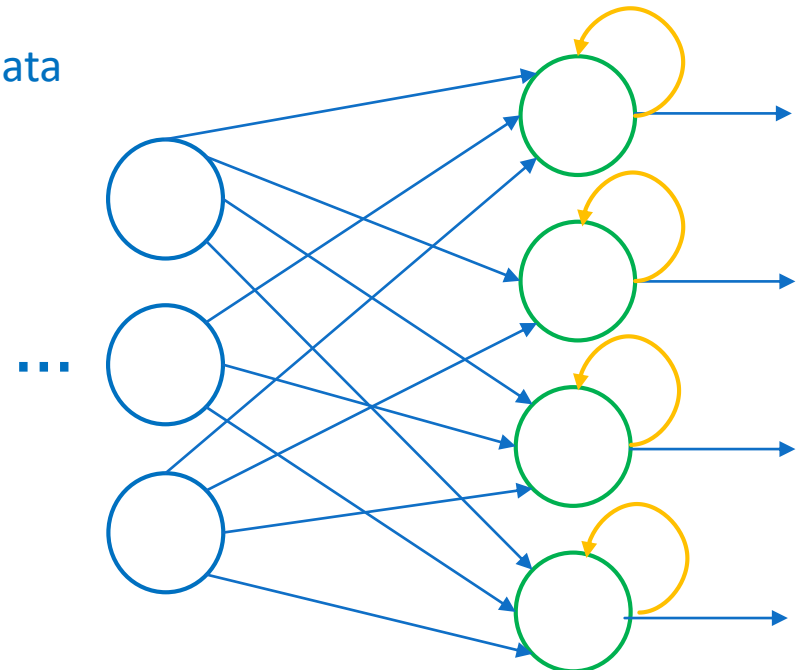
Use case: speech tagging

- ▶ Target: marking up a word in a text based on its **definition** and **context**



Recurrent Neural Networks

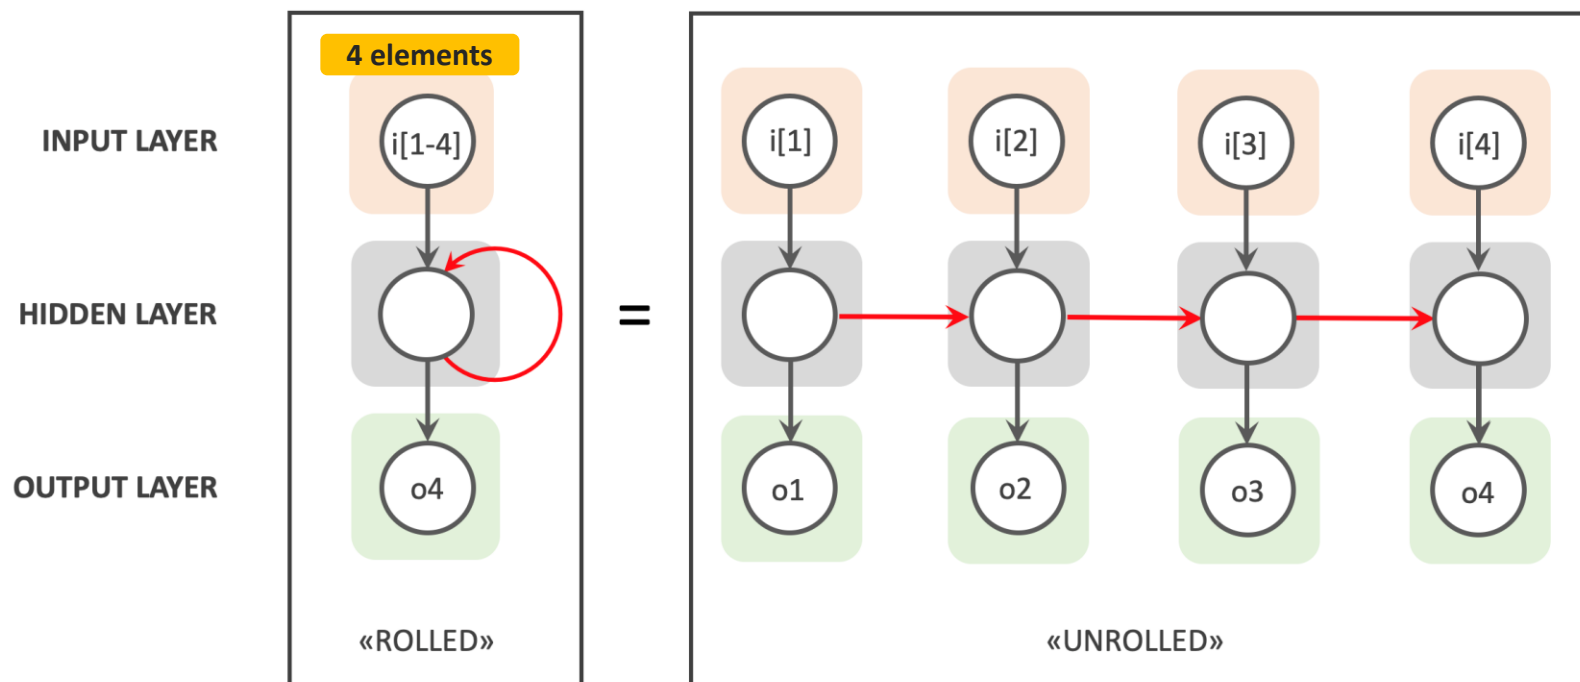
- ▶ A DL algorithm + a type of ANN architecture
 - ▶ The output of step s is fed to the step $s+1$
- ▶ Specialized for processing sequential data
 - ▶ They remember previous inputs
 - ▶ They can share the features
 - ▶ They use historical information
- ▶ Use cases:
 - ▶ Time series predictions
 - ▶ Natural Language Processing (NLP)



Recurrent Neural Networks

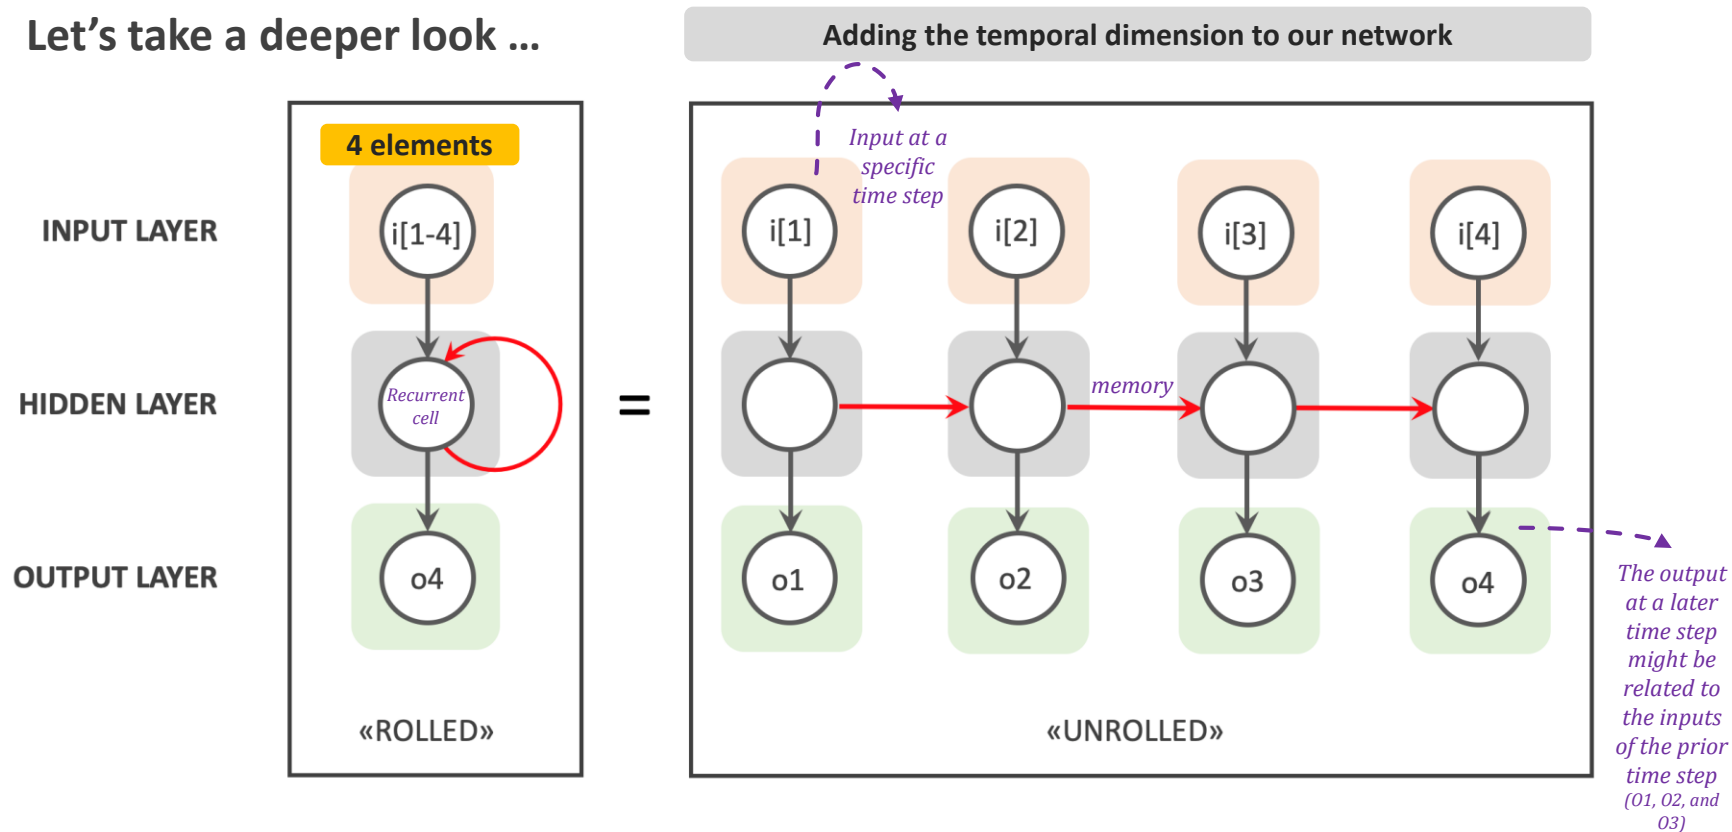
Let's take a deeper look ...

Adding the temporal dimension to our network



Recurrent Neural Networks

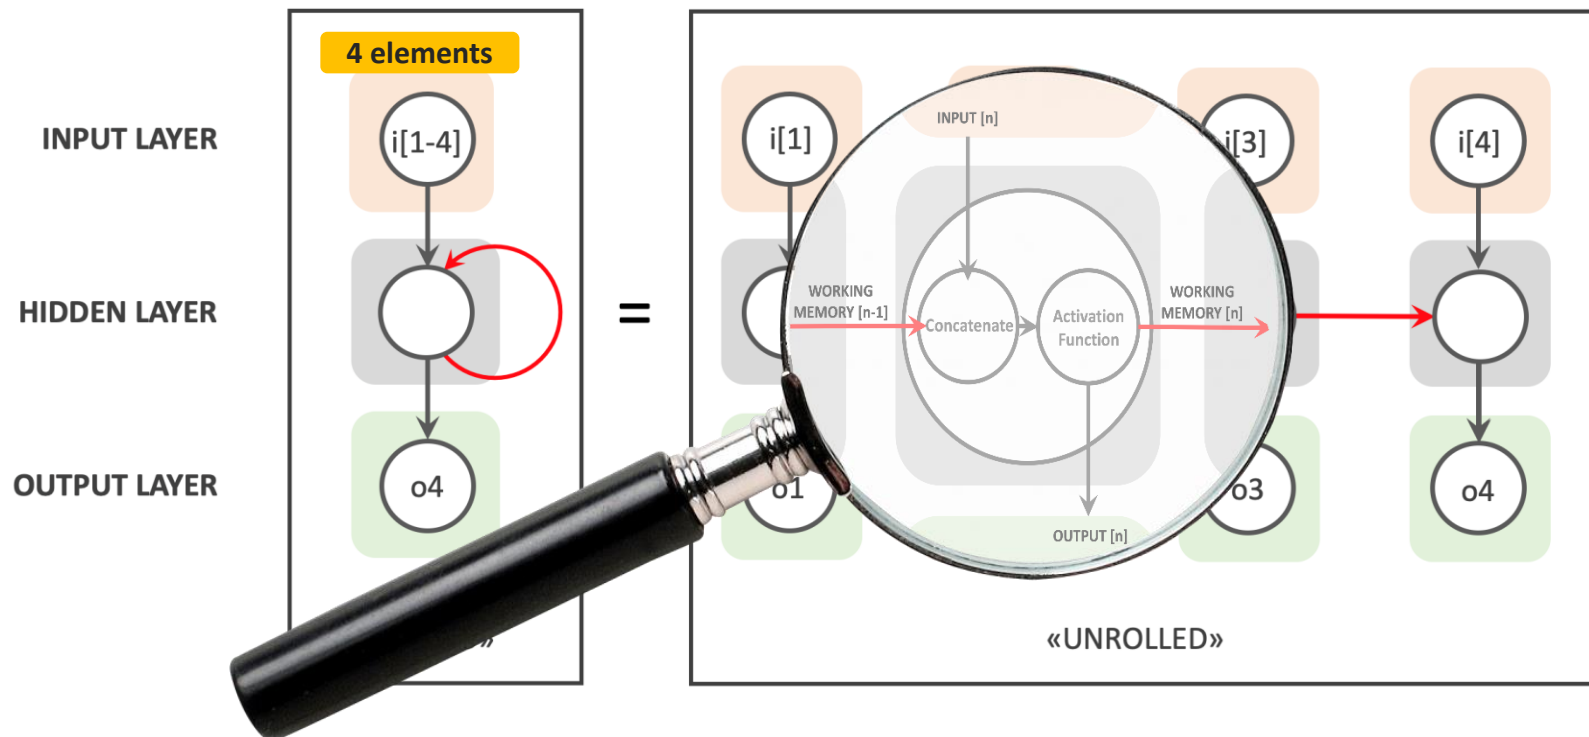
Let's take a deeper look ...



Recurrent Neural Networks

Let's take a deeper look ...

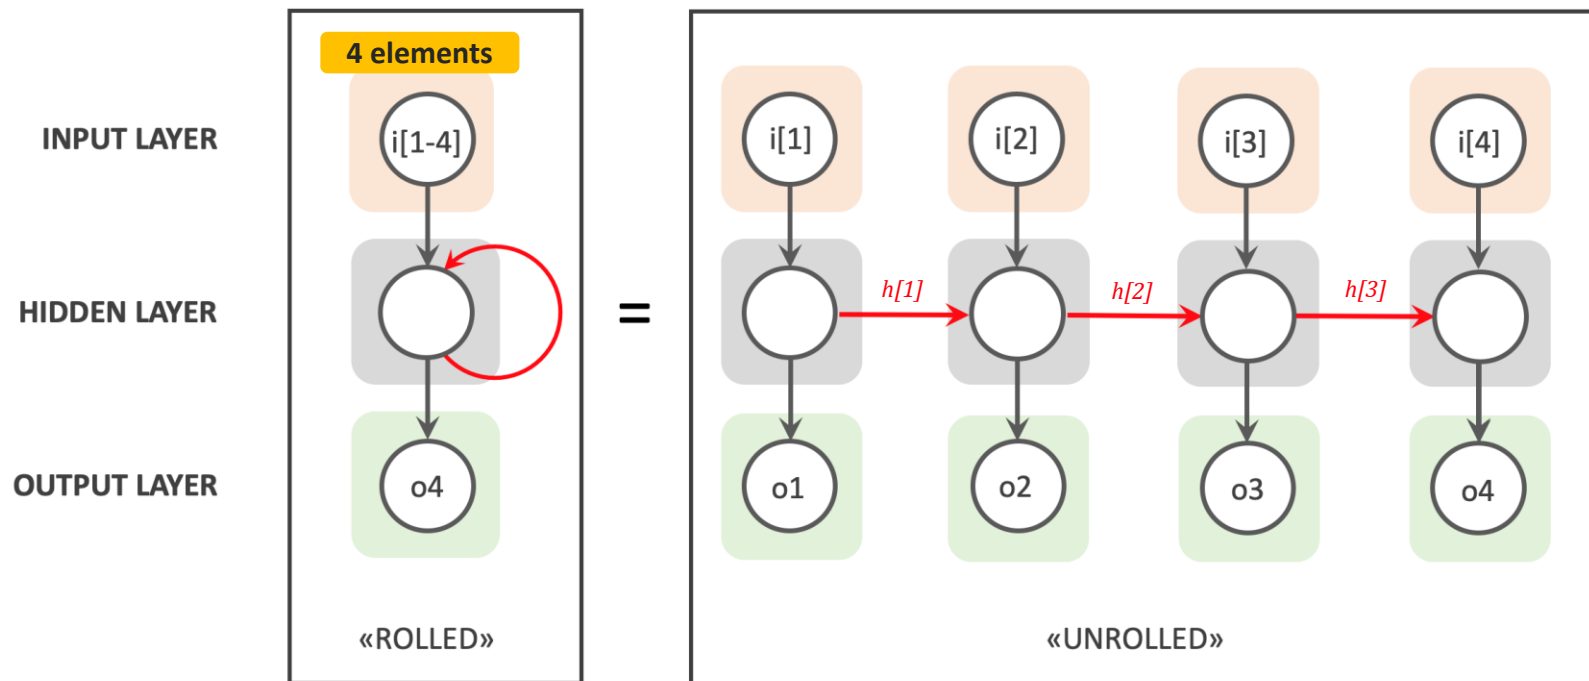
Adding the temporal dimension to our network



Recurrent Neural Networks

Let's take a deeper look ...

$$\hat{o}_t = f(i_t, h_{t-1})$$

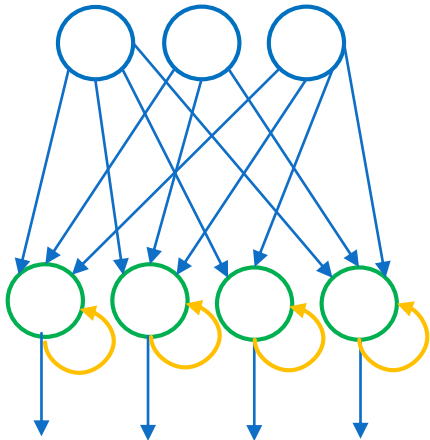


Recurrent Neural Networks



Note that ...

⋮

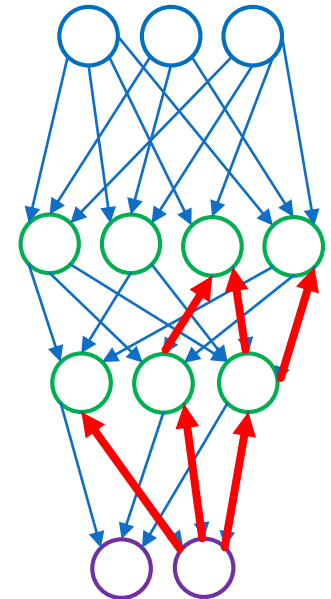


A RNN

- An **architecture**
- Contains weights that are pointed into themselves
- Used for modeling temporal or sequential data

Backpropagation

- A **training algorithm**
- Forwards the values, and only after error calculation, propagates them back to the earlier layers



Recurrent Neural Networks



Important Notes on RNNs

- ▶ RNNs support the processing of sequential data using **loops**
 - ▶ A loop: a **chain** of identical Feedforward ANNs
- ▶ The **loss function** is defined based on **the loss at each time step**
- ▶ **Backpropagation** is done at each point in time (**BPTT**)
 - ▶ Although each loop has its input-output pair, they share the same weights
- ▶ The "working memory" of standard RNNs struggles to retain longer-term dependencies
 - ▶ **Vanishing Gradient problem**

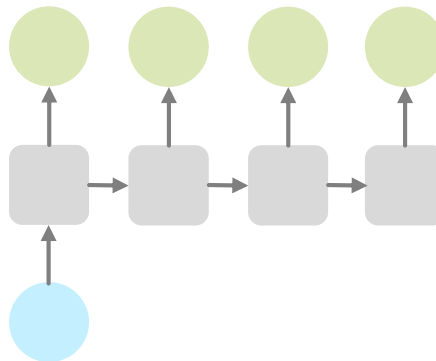


Recurrent Neural Networks



Important Notes on RNNs

- ▶ RNNs can handle **variable-length** sequences while keeping the **order of input data** and **memorizing the dependencies** among them
 - ▶ In contrast with the fixed dimensionality in Feedforward NNs
- ▶ RNNs can **share parameters** (e.g., weights) across the sequence of data



Recurrent Neural Networks



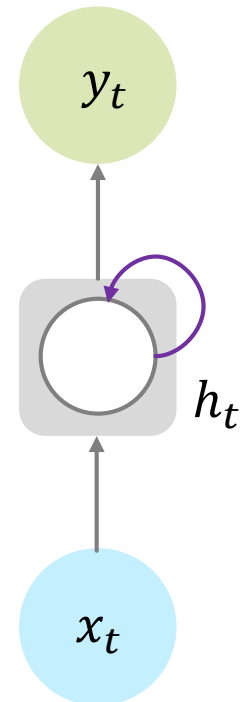
Important Notes on RNNs

- ▶ In RNNs, we can always find a **state** that is updated at each time step

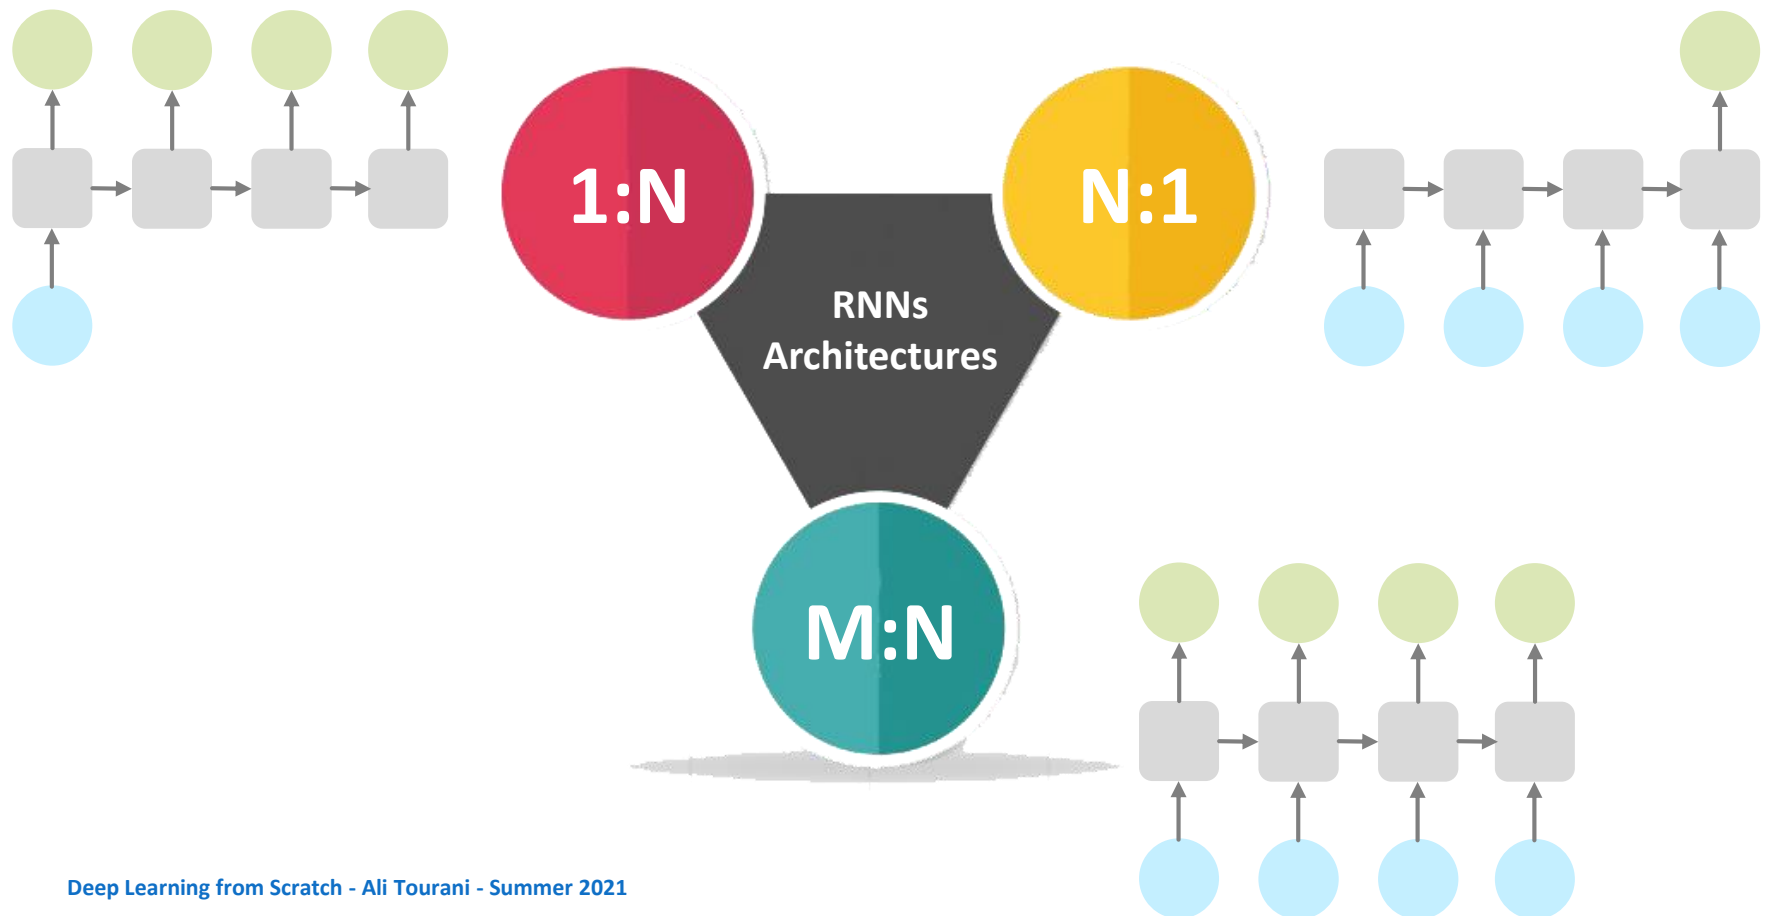
$$h_t = f_W(x_t, h_{t-1})$$

cell state weighted function old state

- ▶ Weights in the weighted function are used for training
- ▶ The same procedure and parameters are used at **every time step**



Recurrent Neural Networks



Recurrent Neural Networks

Use case:

- Image Captioning



A cat is playing
with a green ball

1:N

RNNs
Architectures

N:1

Use case:



- Sentiment Classification

I love this movie. I have
seen it several times



Use case for $M \neq N$:

- Machine Translation

 I always knew what the right path was
 Ho sempre saputo quale fosse la strada giusta

M:N

Use case for $M = N$:

- Named Entity Recognition

My Name is **Ali Person**
 I am from **Iran GPE**
 I know **John Person** very well

Recurrent Neural Networks

Applications of Standard RNNs



Speech Recognition
(human speech to
written text)

Recurrent Neural Networks

Applications of Standard RNNs

Stock Prediction
(determining the future value of a stock)



Recurrent Neural Networks

Applications of Standard RNNs

Enter or paste your text

abc ✓ Spell check

à á â ç é è ê ï ñ ó ô õ ù ú û ü

☐ Auto-translation

French

↔

English

Translate

🇫🇷 🇪🇸 🇩🇪 🇮🇹 🇷🇺 🇨🇳 🇵🇹 🇯🇵 🇧🇪 🇮🇱 🇵🇱 🇹🇷

Online Translation
(translating text/audio data
into another language)

Online translator based on AI technology for French, Spanish, German, Russian and many more languages

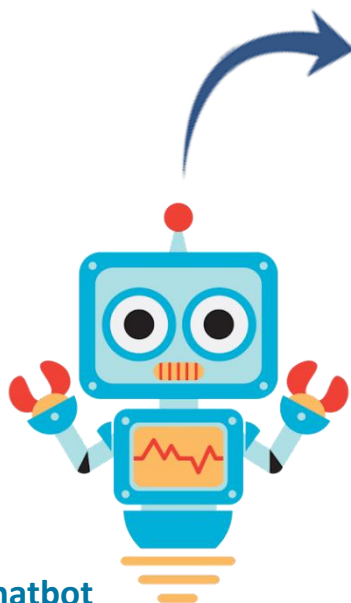
Recurrent Neural Networks

A Simple Application: Chatbot

Your query

How is the
weather in
Roma?

Chatbot



How
is
the
weather
in
Roma
?

Segmentation

How to feed
them to our
RNN?



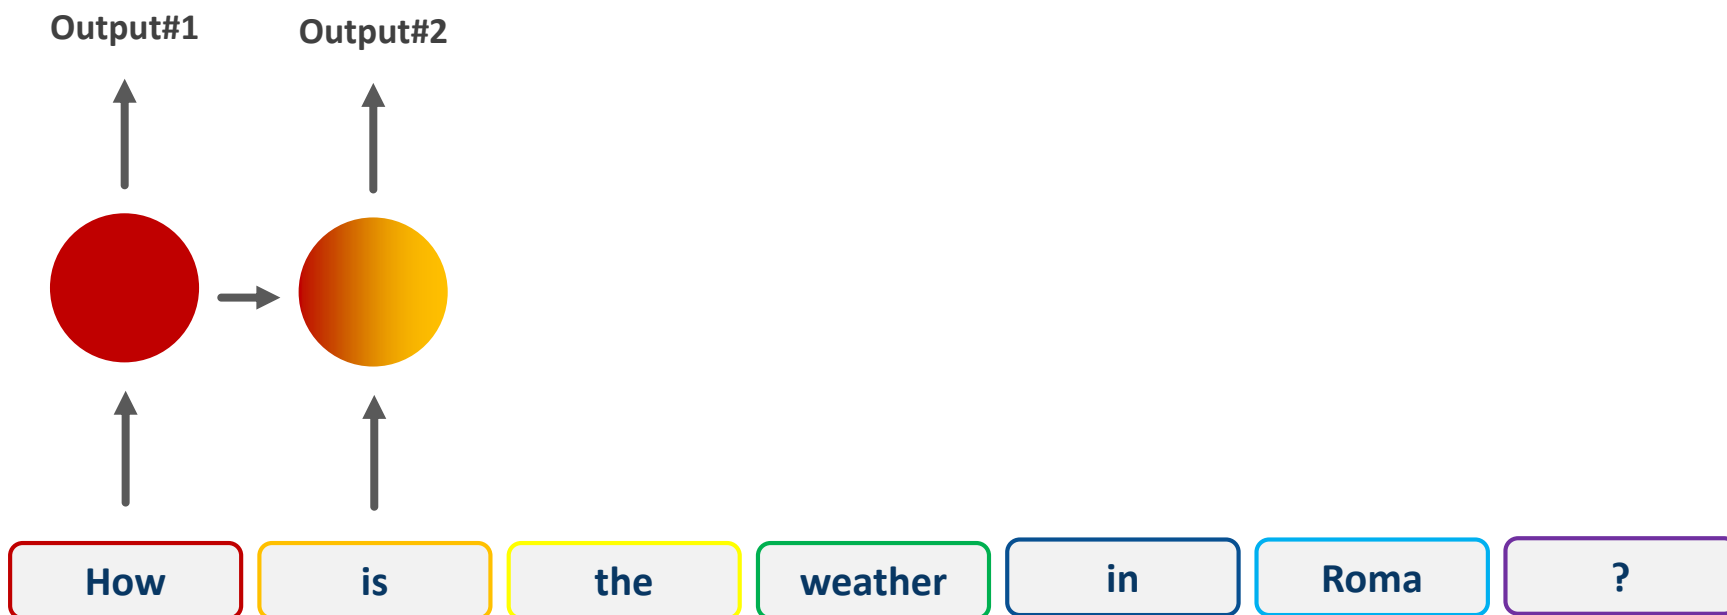
Recurrent Neural Networks

A Simple Application: Chatbot




Recurrent Neural Networks

A Simple Application: Chatbot



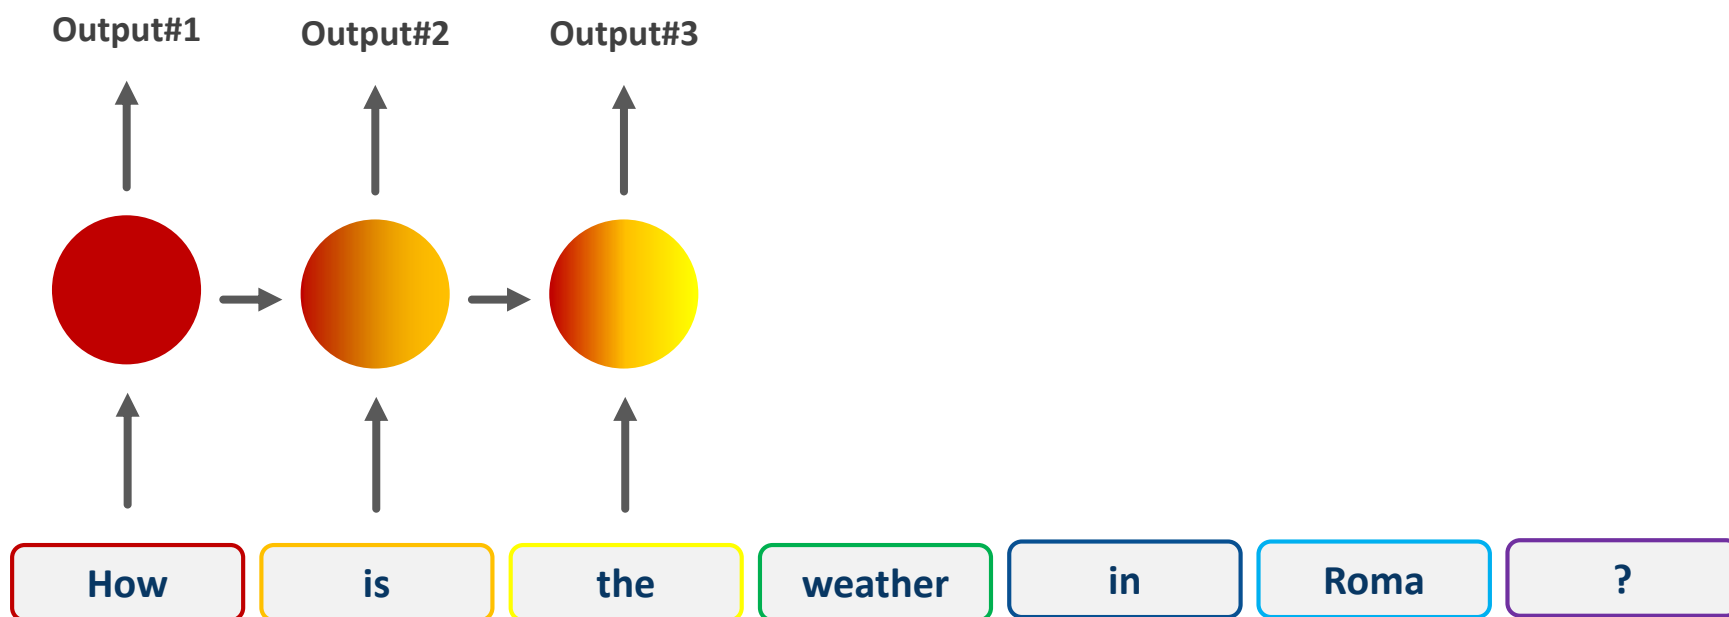
Recurrent Neural Networks

A Simple Application: Chatbot



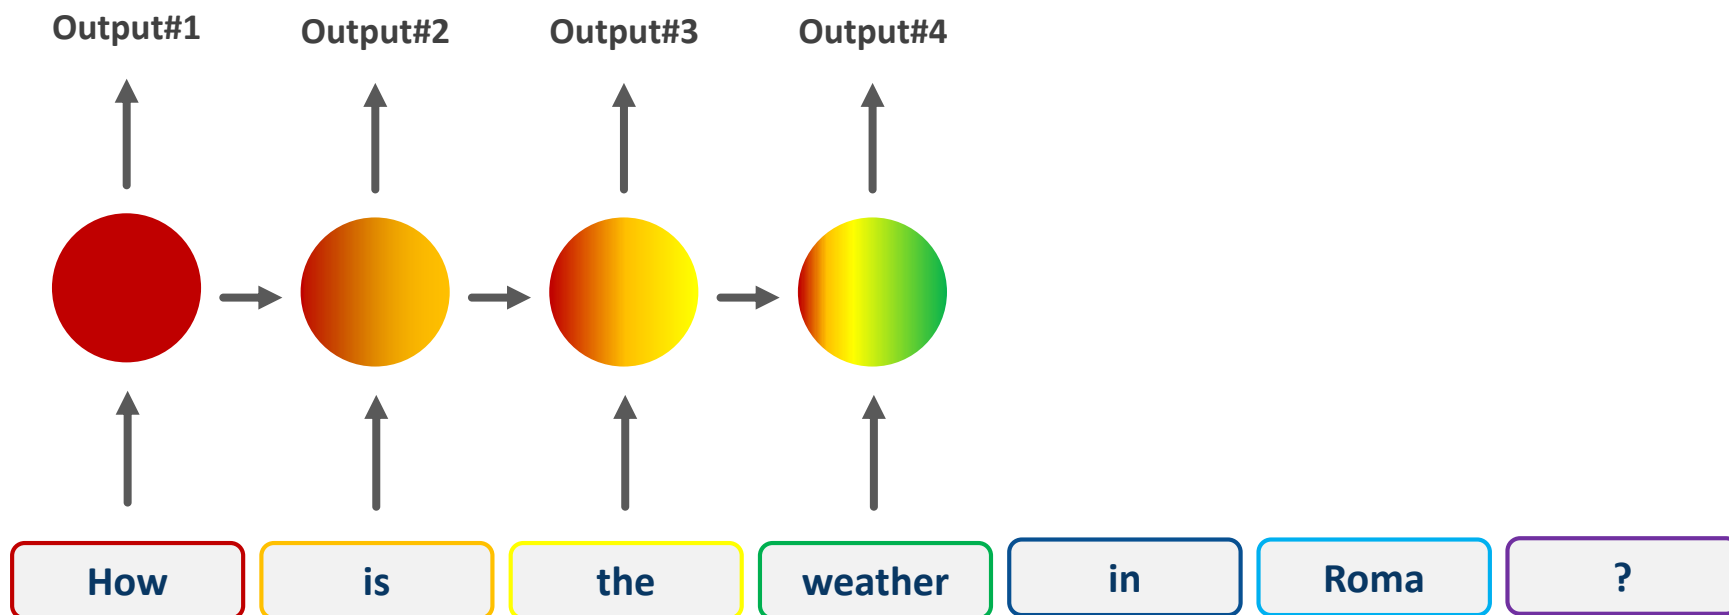
Recurrent Neural Networks

A Simple Application: Chatbot



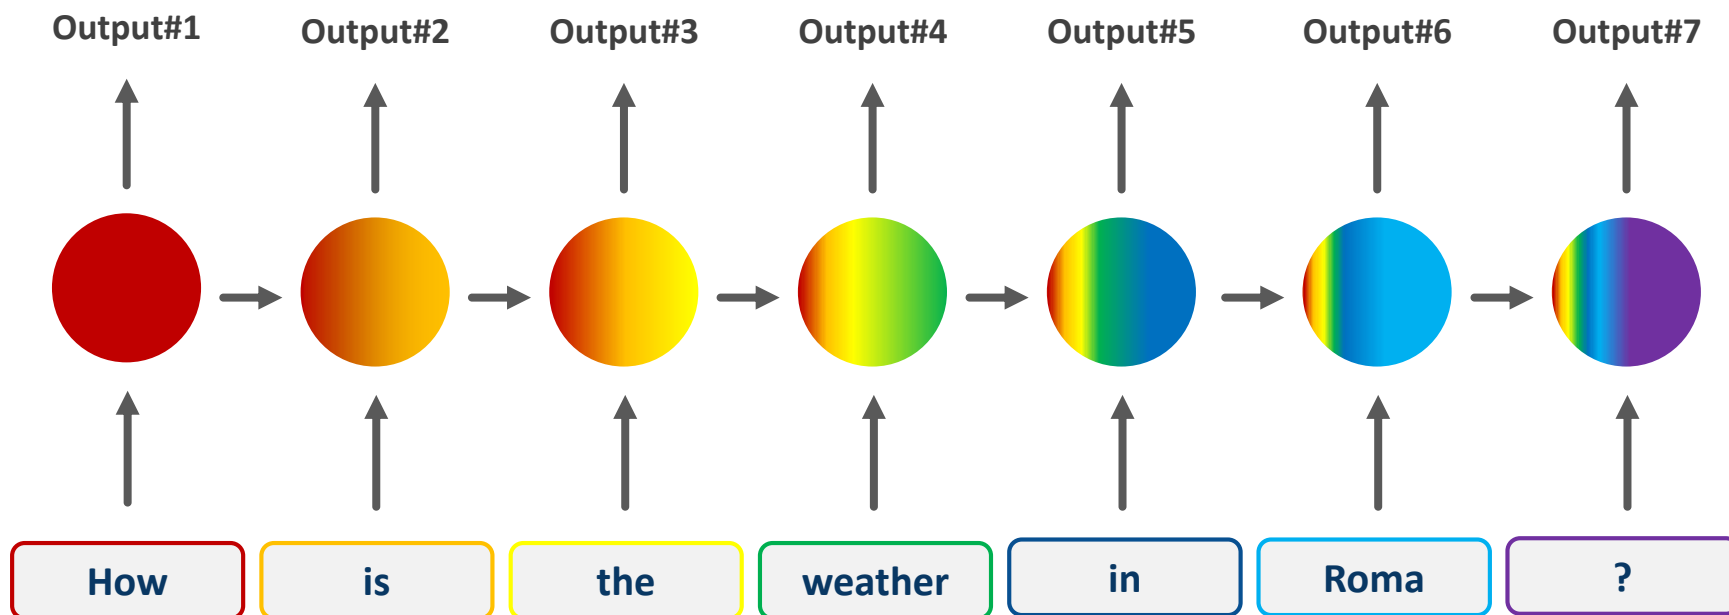
Recurrent Neural Networks

A Simple Application: Chatbot




Recurrent Neural Networks

A Simple Application: Chatbot



Recurrent Neural Networks

A Simple Application: Chatbot



Now the RNN holds the information of all the previous steps

Recurrent Neural Networks

A Simple Application: Chatbot

How is the
weather in
Roma?

Your query


Output#7



28

°C | °F

Precipitation: 0%
Humidity: 44%
Wind: 19 km/h



Chatbot

Recurrent Neural Networks

A Simple Application: Word Prediction

We decided to go on a short trip this _____.


Given

Want to know

- ▶ How to represent the concepts behind words to our ANN?
 - ▶ Surely, ANNs cannot understand the meaning of the words!
 - ▶ We have to find a way to make them numerical
 - ▶ Goal: applying mathematical operations

["We", "to", ...]

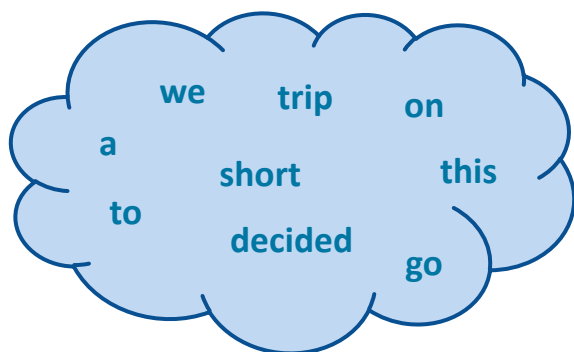
[0.1, 0.9, ...]



Recurrent Neural Networks

A Simple Application: Word Prediction

We decided to go on a short trip this _____.



Vocabulary

Diagram showing the flow from Vocabulary to Indexing to Embedding. A curved arrow points from the Vocabulary cloud to the Indexing table, and another curved arrow points from the Indexing table to the Embedding list.

Word	Index
a	00001
decided	00002
go	00003
...	...

Indexing

"a" = [1, 0, 0, 0, 0, ...]
"decided" = [0, 1, 0, 0, 0, ...]
"go" = [0, 0, 1, 0, 0, ...]
.....

Embedding

Recurrent Neural Networks

alitourani Update README.md 0976f6d yesterday History

File	Commit	Time
README.md	Update README.md	yesterday
_RecurrentNeuralNetwork-RNN-AliTourani-DeepLearnin...	feat: add thumbnail for RNNs	yesterday

README.md

Recurrent Neural Networks (RNNs)

In contrast with Feed-Forward NNs in which the information is only passed in one direction, RNNs can easily handle sequential data processing. These architectures of Artificial Neural Networks (ANNs) can remember previous inputs, share the features across the network, and use historical information.

The diagram illustrates an unrolled RNN. At the bottom, there is a sequence of input boxes: 'How', 'is', 'the', 'weather', 'in', 'Roma', and '!'. Above these inputs are seven colored circles representing hidden states, connected by horizontal arrows from left to right. Each hidden state is connected to an output label above it by a vertical arrow. The output labels are 'Output#1' through 'Output#7'. The first four hidden states are fully visible and colored (red, orange, yellow, green), while the last three are partially obscured by a teal callout box. The callout box contains the text 'Full code on GitHub' and a hand cursor icon pointing at it.

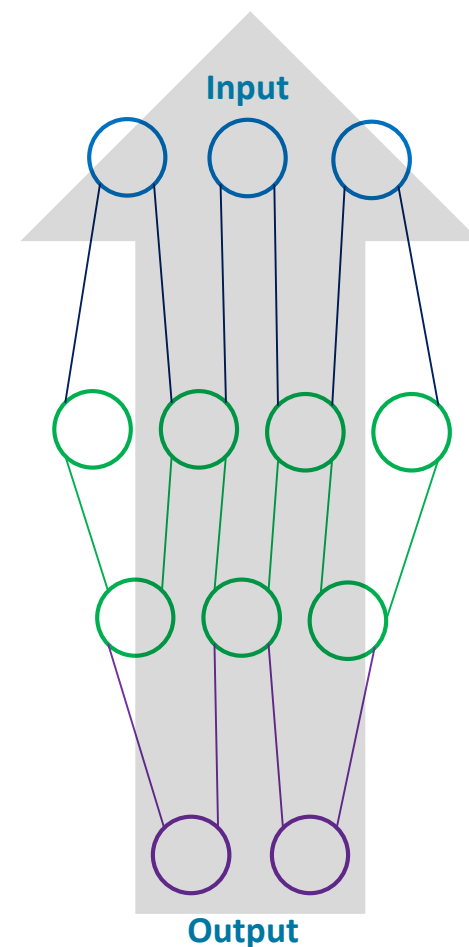
Backpropagation Through Time

Backpropagation in Feedforward ANNs

- ▶ Recall: [Session#1](#)
- ▶ GDA algorithm
 - ▶ The gradient of the loss with respect to each weight parameter
 - ▶ Shifting parameters to minimize final loss

$$\Delta w_i(t) = \mu \left(-\frac{\delta J(W)}{\delta w} \right)$$

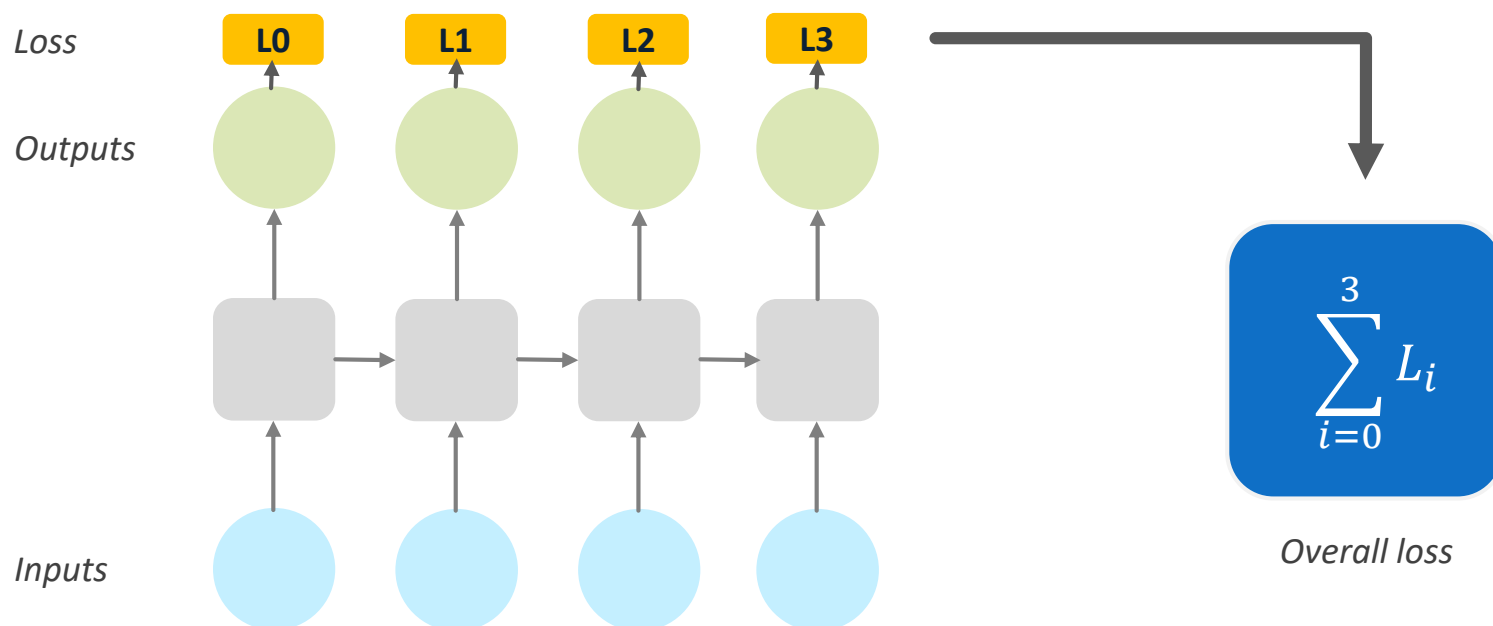
$$\frac{\delta J(W)}{\delta w_1} = \frac{\delta J(W)}{\delta y} * \frac{\delta y}{\delta b} * \frac{\delta b}{\delta a} * \frac{\delta a}{\delta w_1}$$



Backpropagation Through Time

Backpropagation in RNNs


- In the **forward pass**, loss values are calculated at each time step.



Backpropagation Through Time

Backpropagation in RNNs

- In the **backward pass**, the overall loss spreads errors in each time step.



We need to unroll all input time steps, calculate and accumulate errors, and then roll back and update weights

Backpropagation Through Time

Backpropagation in RNNs

- ▶ Flowing gradients in a repetitive manner cause some problems:
 - ▶ An input sequence with thousands of time steps means *thousands of derivatives* for a single weight update!
 - ▶ Many **matrix multiplications** for updating the weights
 - ▶ Adding **many computation factors** while moving towards the initial state
 - ▶ We may face one of these problems:
 1. **Exploding gradient:** accumulation of large error gradients → *huge updates*
 2. **Vanishing gradient:** minimal gradients → *preventing the update process*



Backpropagation Through Time

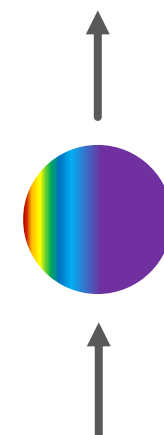
Vanishing Gradient Problem

- ▶ Recall the GDA and backpropagation
- ▶ If the number of hidden layers is huge, the gradient **diminishes dramatically** as it is propagated backward
- ▶ The error is so tiny when it reaches the layers close to the input
 - ▶ A smaller error means less impact on the learning process → **Vanishing Gradients**
- ▶ The network captures only short-term dependencies



Vanishing gradients challenges the network, so it is unclear which direction the parameters should move to improve the cost function.

Output #n

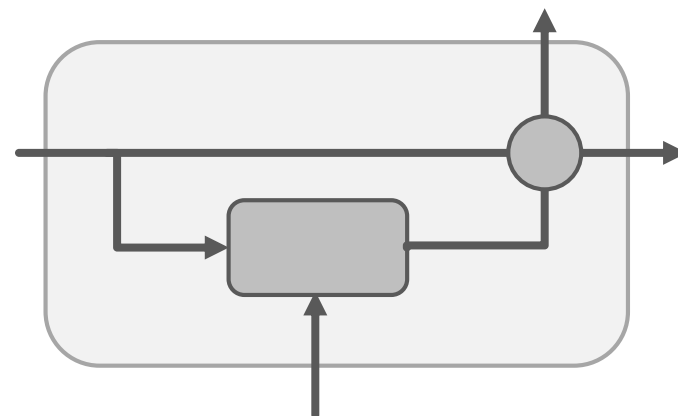
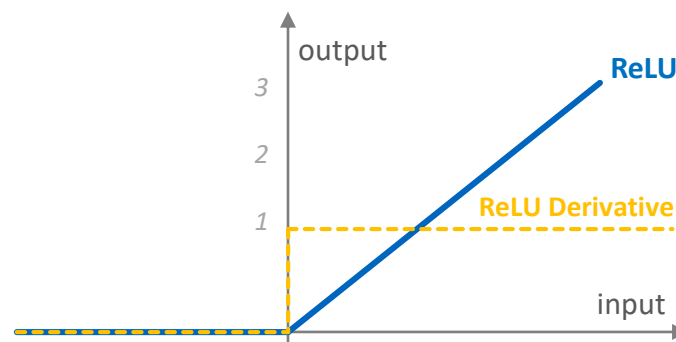


Input #n

Backpropagation Through Time

Vanishing Gradient Problem

- ▶ We can resolve this issue by:
 - ▶ Smartly selecting the AFs of the network
 - ▶ **ReLU** is a great choice!
 - ▶ Smartly initializing the parameters
 - ▶ Trying to prevent the weights from shrinking to zero
 - ▶ Using **Gated Cells**
 - ▶ Using a complex recurrent unit with gates to enable controlling the data

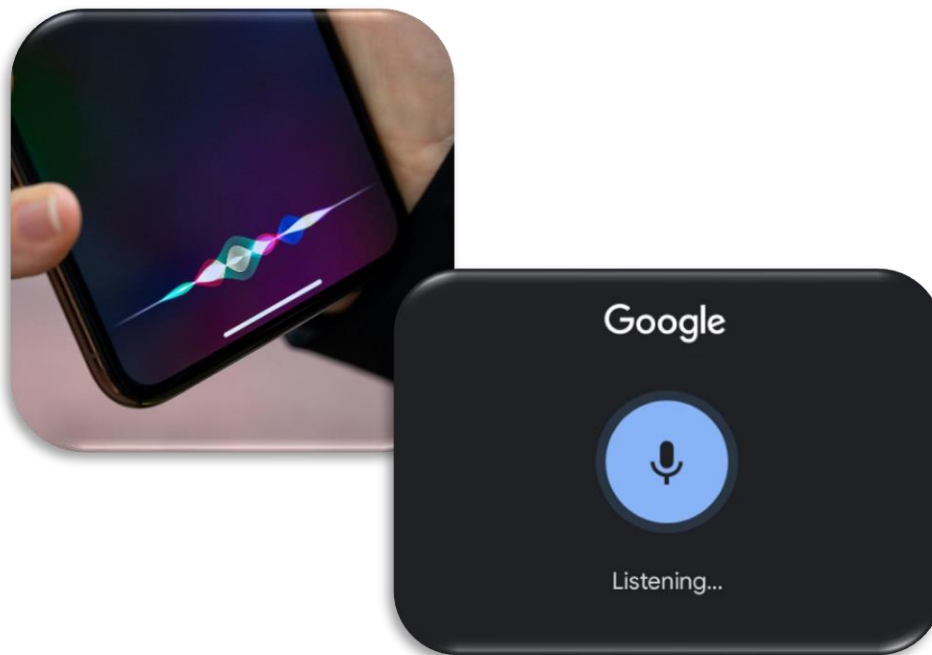


Long Short-Term Memory

- ▶ A popular deep learning algorithm for [Sequence Models](#)
- ▶ It uses a **Gated Cell** to track information that flows in the network throughout many time steps

Use cases:

- ▶ Apple's Siri
- ▶ Google's voice search
- ▶ Time-series predictions
- ▶ Text classification



Long Short-Term Memory

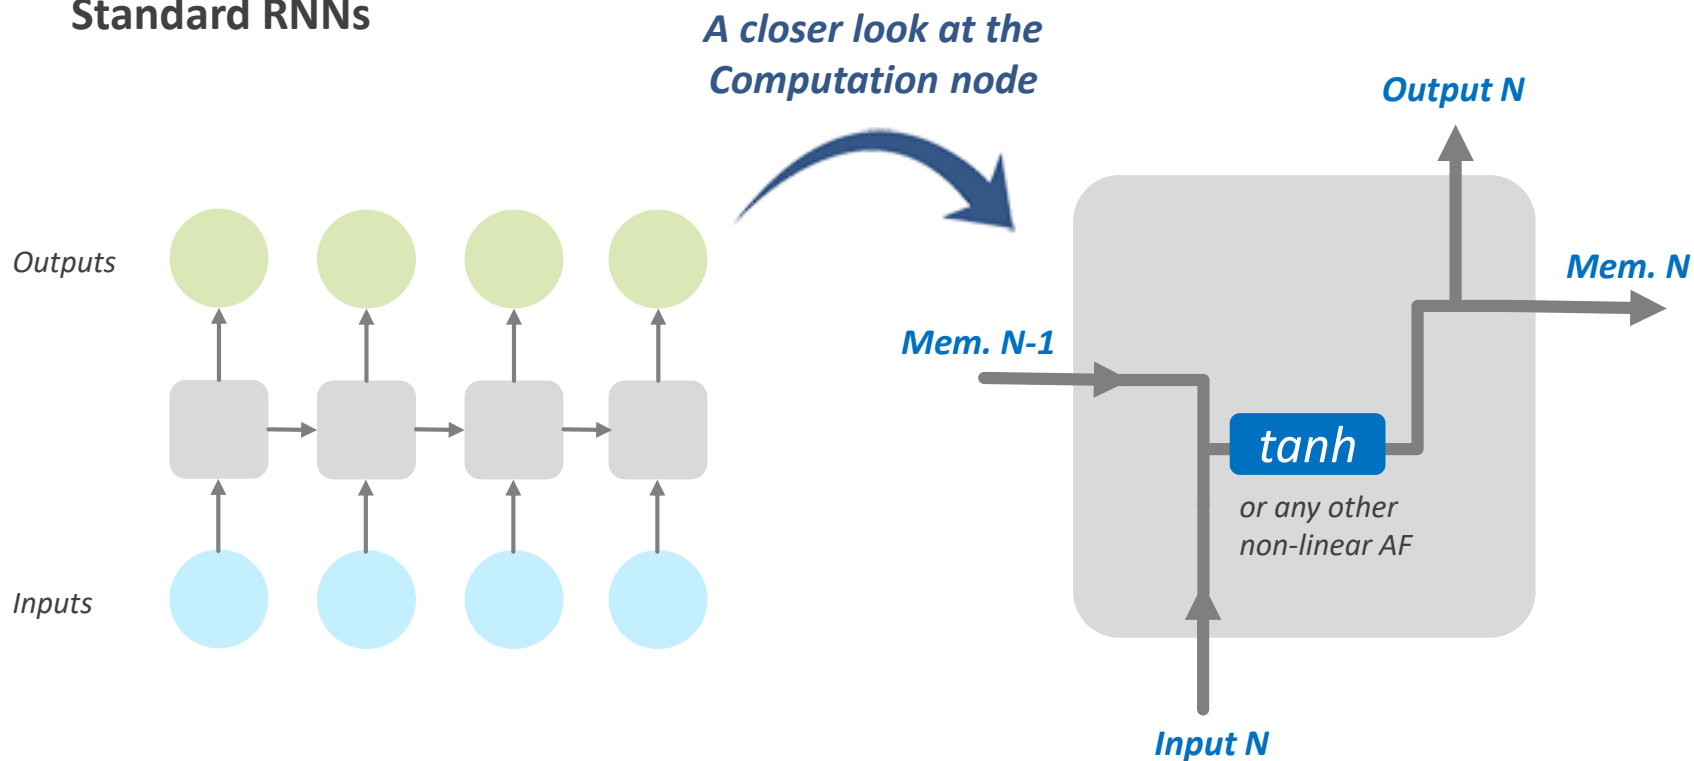
What is LSTM?

- ▶ Traditional RNNs are not good at capturing long-range dependencies
 - ▶ The main reason is the **Vanishing Gradient problem**
 - ▶ They may stop the ANN from being mature through training
- ▶ **LSTMs** are able to remember the input over a longer period
- ▶ How does it reflect the inputs?
 - ▶ Through passing the two things to the next time step:
 - ▶ **Cell state** (the long-term memory)
 - ▶ **Hidden state** (an output of a cell that is being updated at every step)



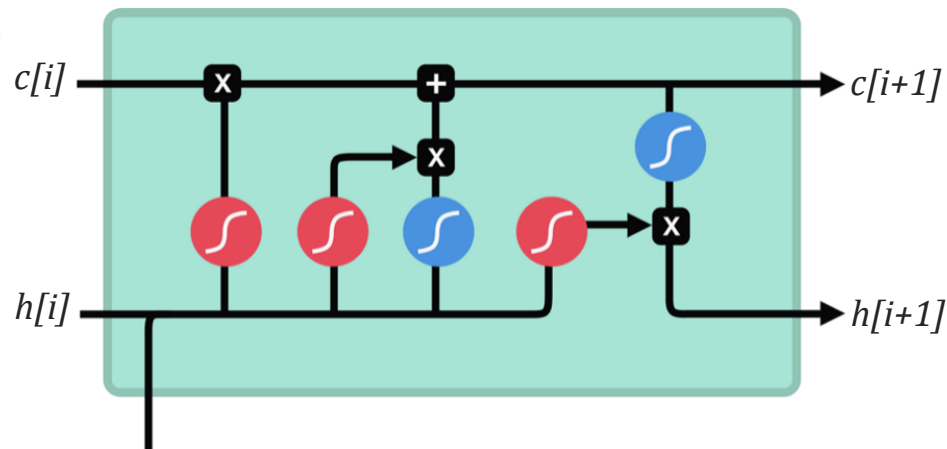
Long Short-Term Memory

Standard RNNs



Long Short-Term Memory

LSTMs' repeating components contain Computational Blocks to control the flow of information.



sigmoid



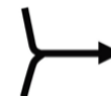
tanh



pointwise
multiplication



pointwise
addition



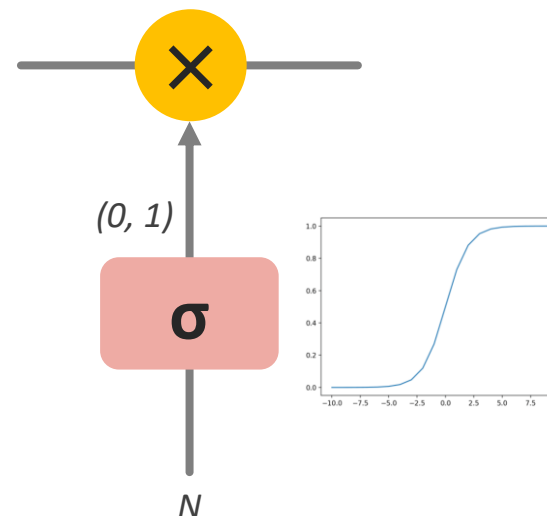
vector
concatenation

Long Short-Term Memory

Memory Management

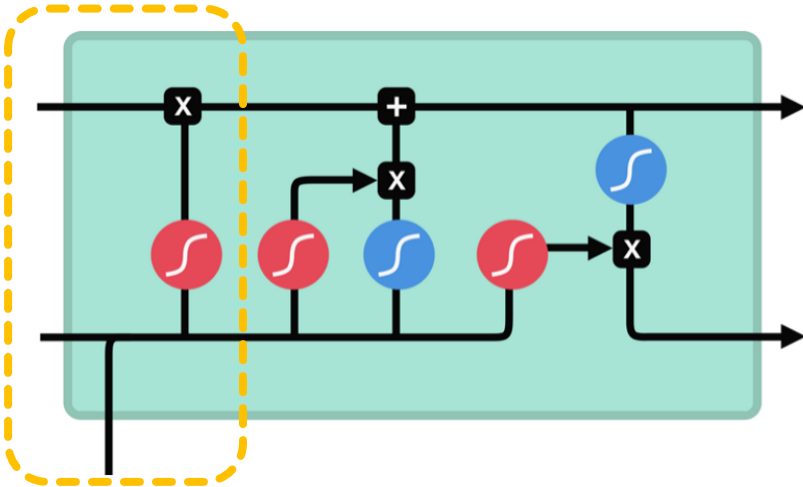
- ▶ LSTMs use a **Gating Mechanism** for:
 - ▶ Controlling the gradient propagation
 - ▶ Keeping, updating, ignoring, or forgetting information in the memory cell

*A Gate lets the information pass via **an ANN layer** (e.g., Sigmoid) and a **pointwise multiplication***



Long Short-Term Memory

Using this block, LSTM can **forget** irrelevant parts of the Previous State (PS) through passing the PS through a Sigmoid AF for filtration



sigmoid



tanh



**pointwise
multiplication**



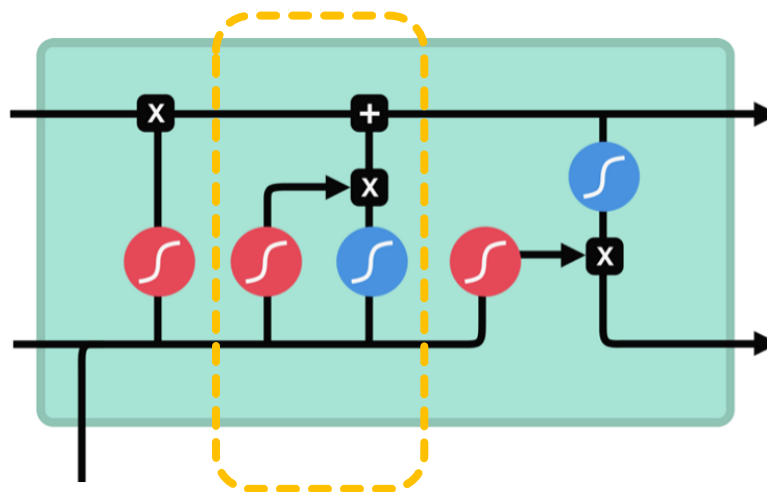
**pointwise
addition**



vector concatenation

Long Short-Term Memory

Using this block, LSTM can **store** relevant information into the cell state and keeping them in the memory



sigmoid



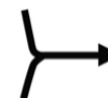
tanh



pointwise
multiplication



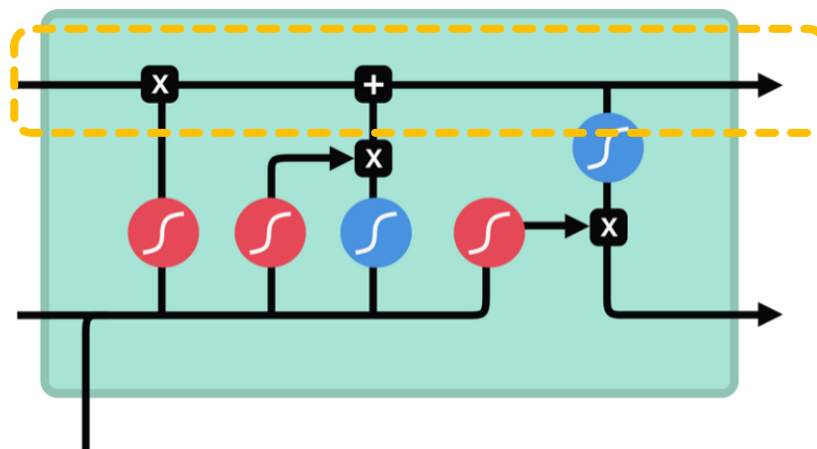
pointwise
addition



vector
concatenation

Long Short-Term Memory

Using this block, LSTM can **update** the selected values of information in the cell state



sigmoid



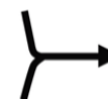
tanh



pointwise
multiplication

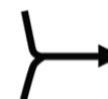


pointwise
addition



vector
concatenation

Long Short-Term Memory



**vector
concatenation**

Long Short-Term Memory



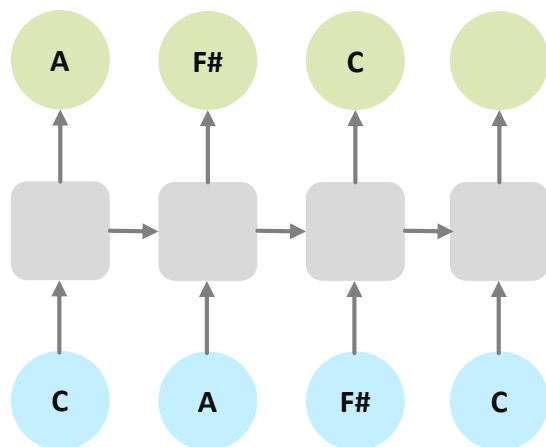
Important Notes on LSTMs

- ▶ They have a **separate Cell State** along with the output in each step
- ▶ They use **gates** for information flow management and control
- ▶ They can provide a **Backpropagation through Time (BPTT)** process with uninterrupted gradient flow
 - ▶ An improved training process and efficient updating of weights
- ▶ They can easily capture **long-range dependencies**
- ▶ In contrast with simple RNNs, LSTMs can grab information from the distant past to predict the current/future state:

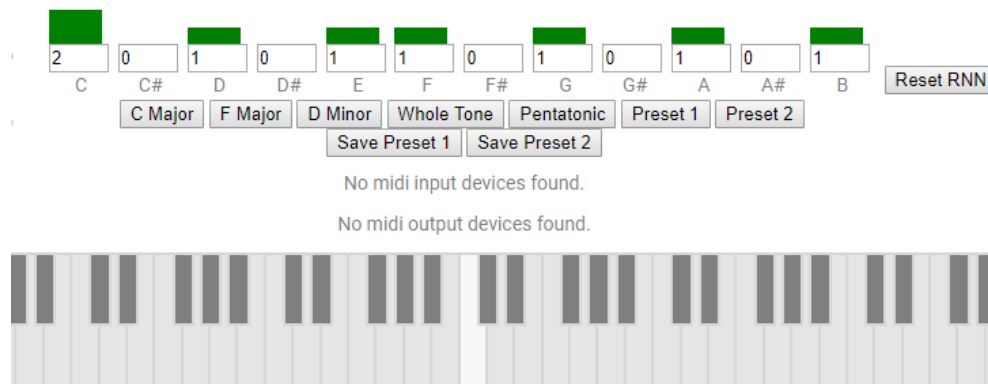
Football is my favorite sport. I have seen many matches so far, and that is why I always dream to be a _____.

Long Short-Term Memory

Applications of LSTMs

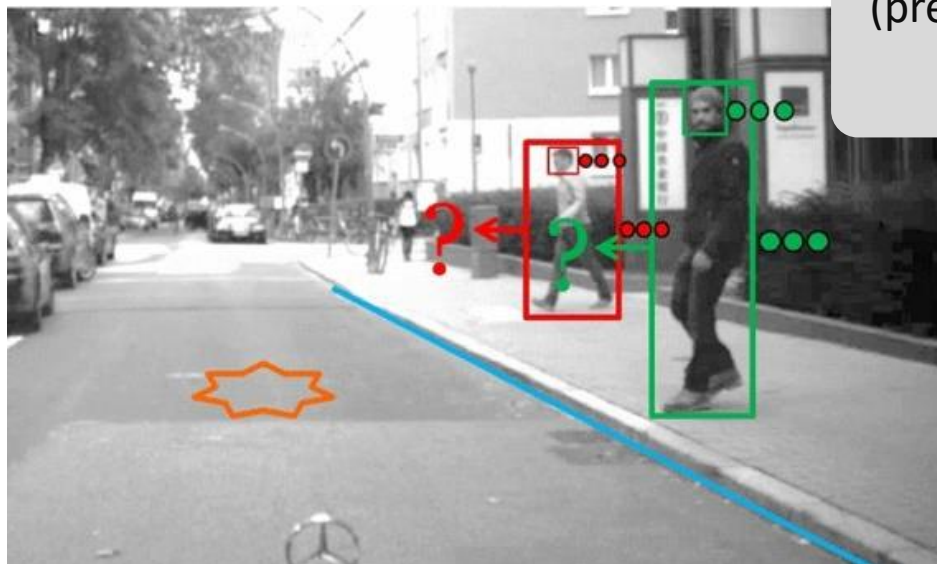


Music Generation
(generating the next character in a given sheet)

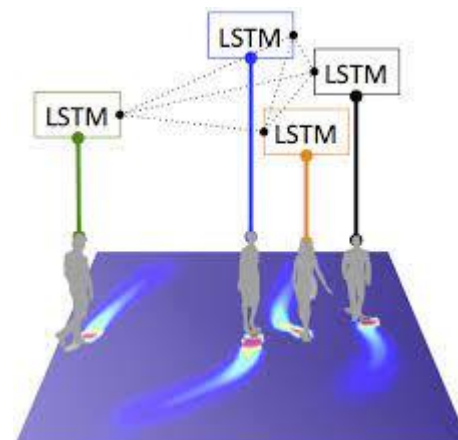


Long Short-Term Memory

Applications of LSTMs





Trajectory Prediction
(predicting the next location of
an object in the scene)

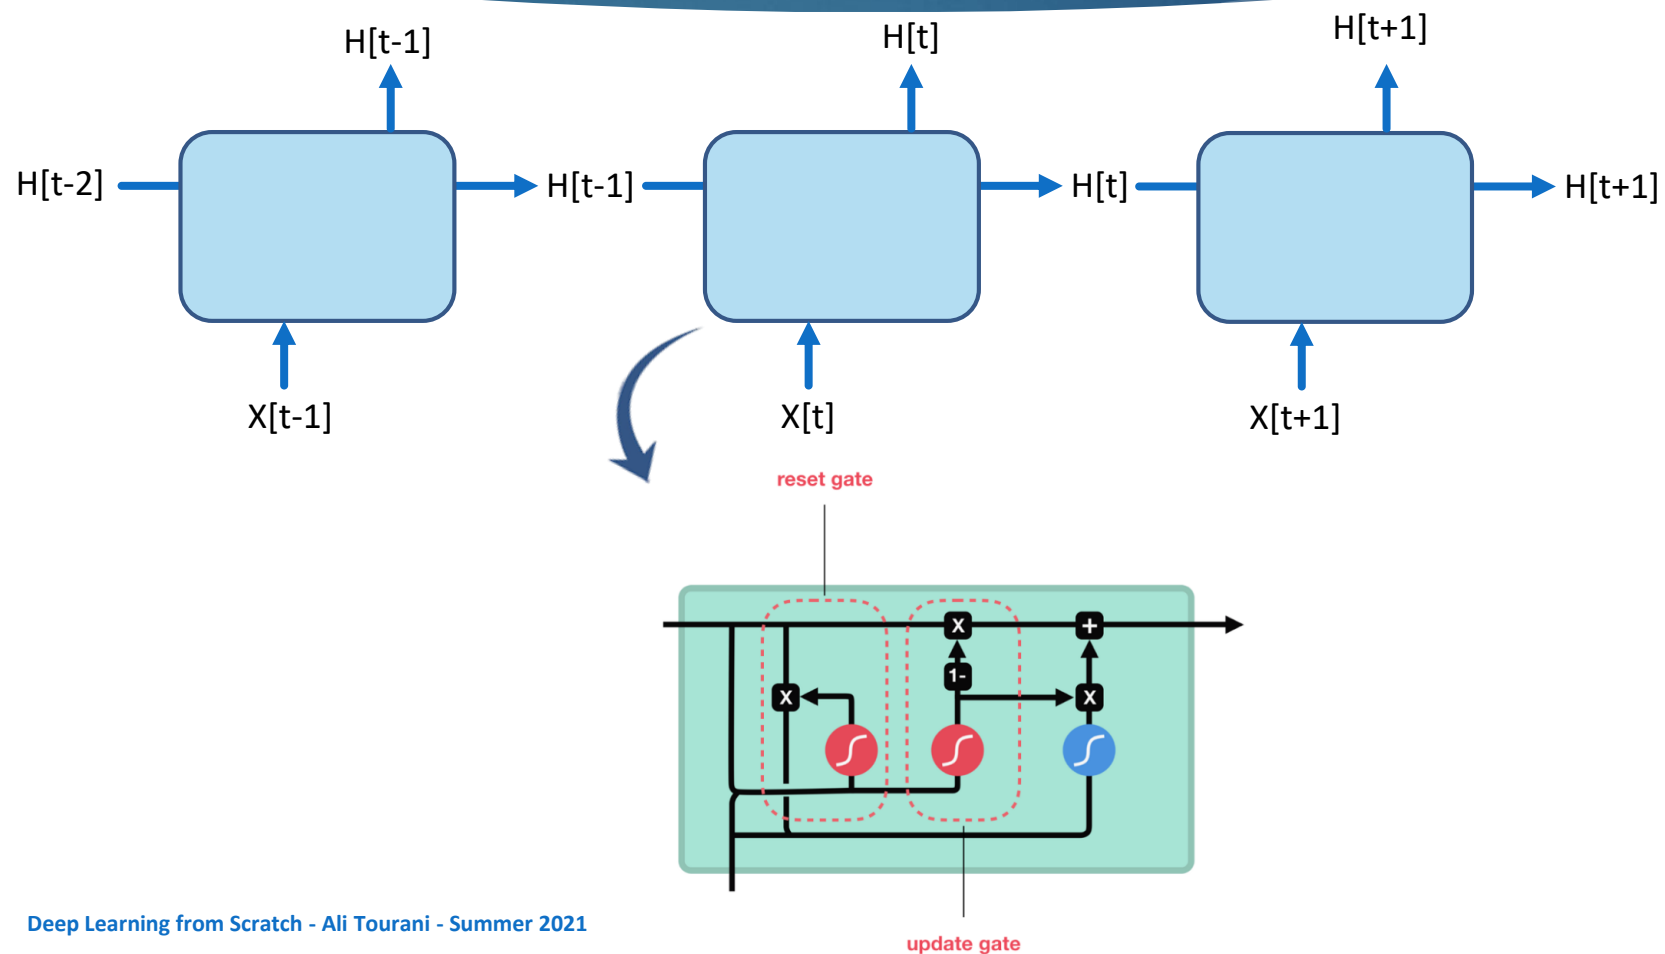


Gated Recurrent Unit

- ▶ One of the most recent RNN approaches, AKA **GRU**
- ▶ Very similar to LSTM, but with a much simpler architecture
 - ▶ GRUs do not contain a **Cell State (long-term memory)**
- ▶ Equipped with two fundamental gates
 - ▶ A **Reset Gate (short-term memory)** and an **Update Gate (long-term memory)**

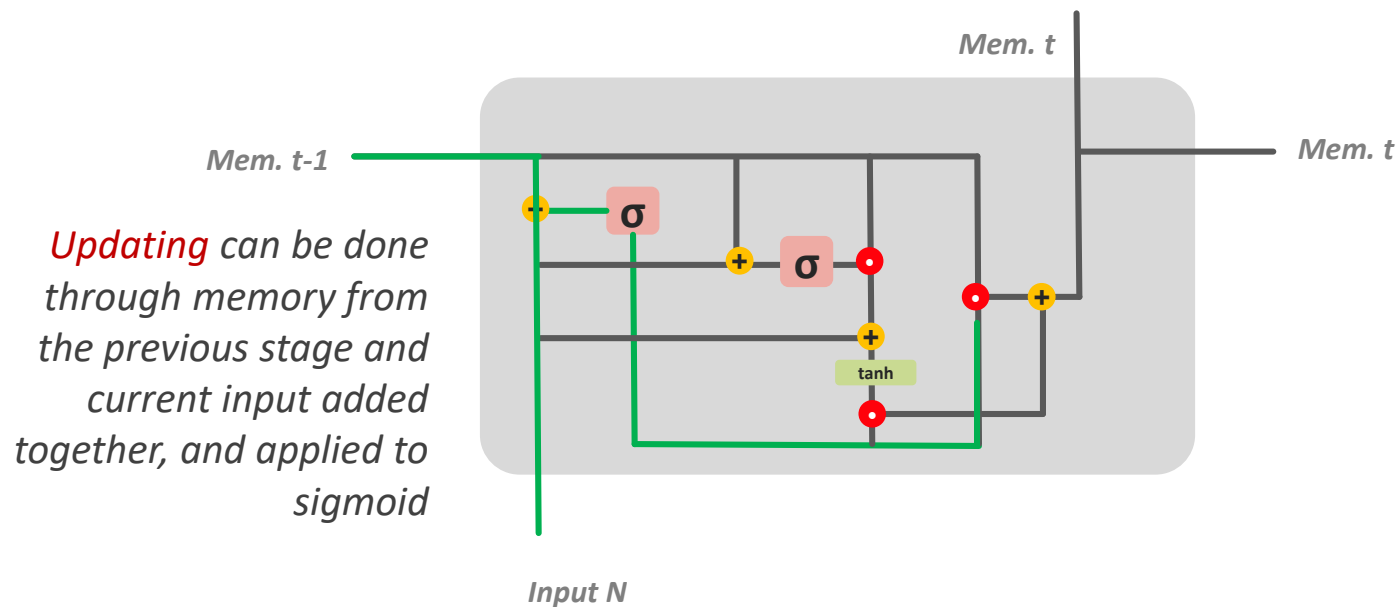


Gated Recurrent Unit




Gated Recurrent Unit

A closer look at GRUs



Gated Recurrent Unit

A closer look at GRUs



References

- ▶ <http://www.IntroToDeepLearning.com>
- ▶ <https://towardsdatascience.com/sequence-models-and-recurrent-neural-networks-rnns-62cadeb4f1e1>
- ▶ <https://towardsdatascience.com/introduction-to-sequence-modeling-problems-665817b7e583>
- ▶ <https://www.bouvet.no/bouvet-deler/explaining-recurrent-neural-networks>
- ▶ <https://machinelearningmastery.com/how-to-fix-vanishing-gradients-using-the-rectified-linear-activation-function/>

Questions?

