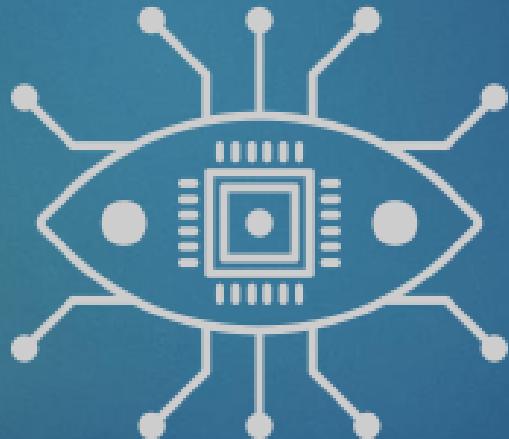


Deep Learning from Scratch

Session #7: Convolutional Neural Networks



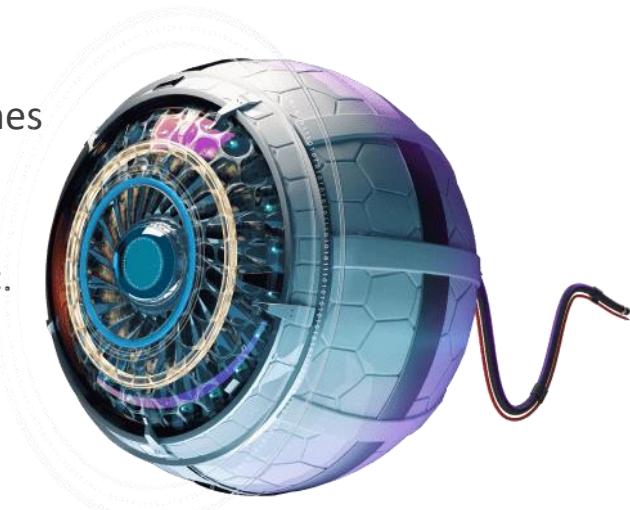
by: Ali Tourani – Summer 2021

Agenda

- ▶ Computer Vision
- ▶ Feature Extraction
- ▶ Classification vs. Regression
- ▶ Convolution
- ▶ Convolutional Neural Networks
- ▶ CNNs Layers

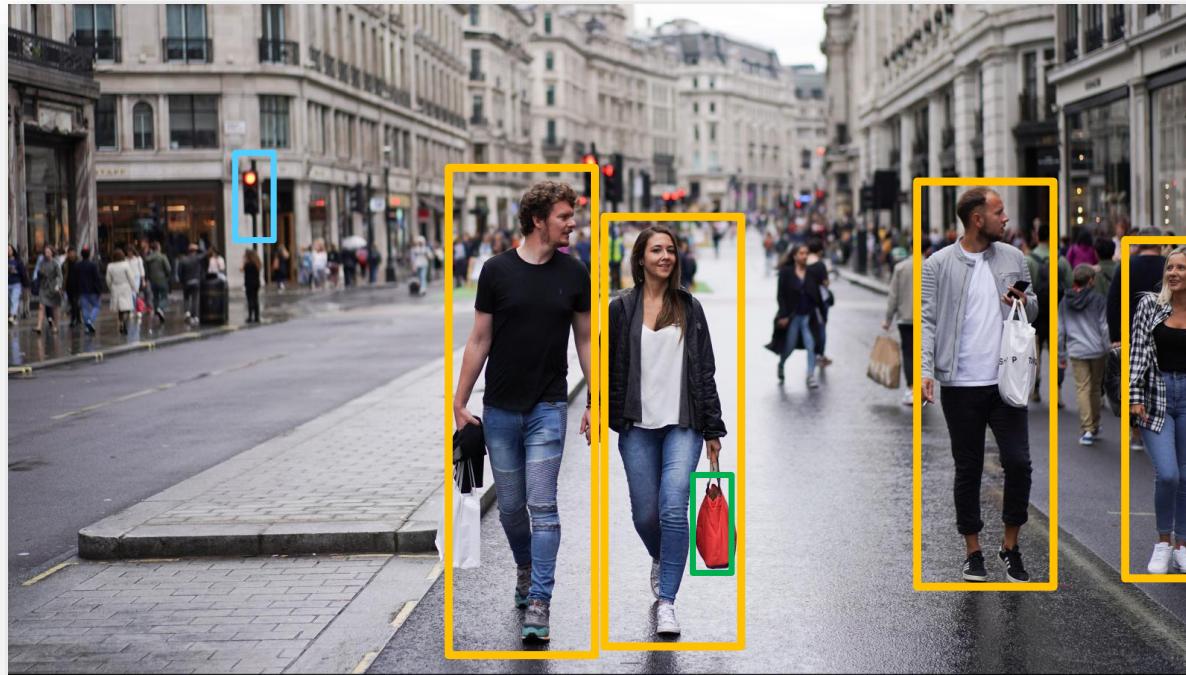
Computer Vision

- ▶ A very popular field of AI and an **interdisciplinary scientific field**
- ▶ Enabling computers to understand from digital images or videos
 - ▶ Taking action based on the information obtained from **visual inputs**
- ▶ Data are captured from **cameras** or generated in computers
- ▶ Modelling **image processing** using **ML techniques**
 - ▶ Applying ML to recognize patterns from images/frames
- ▶ What actions to do?
 - ▶ Distinguishing between objects, classifying them, etc.



Computer Vision

By utilizing Computer Vision algorithms we can **identify objects**



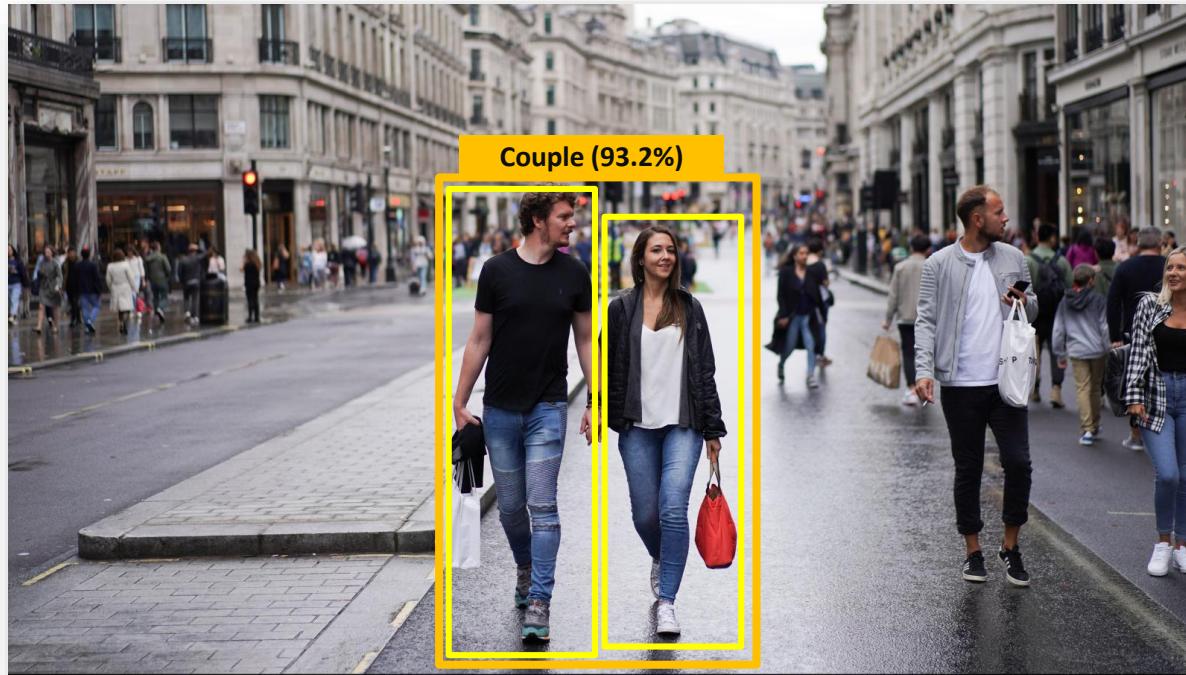
Computer Vision

By utilizing Computer Vision algorithms we can **predict future actions/events**



Computer Vision

By utilizing Computer Vision algorithms we can **extract info and discover patterns**



Computer Vision



Computer Vision

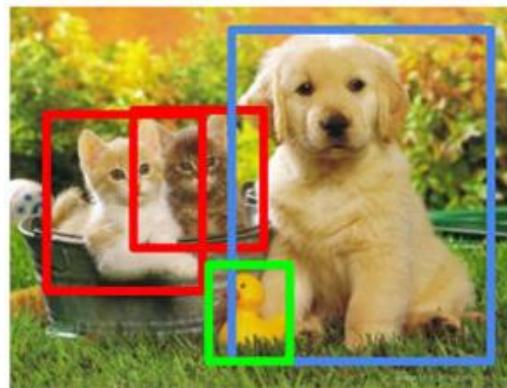
Image Classification



Computer Vision

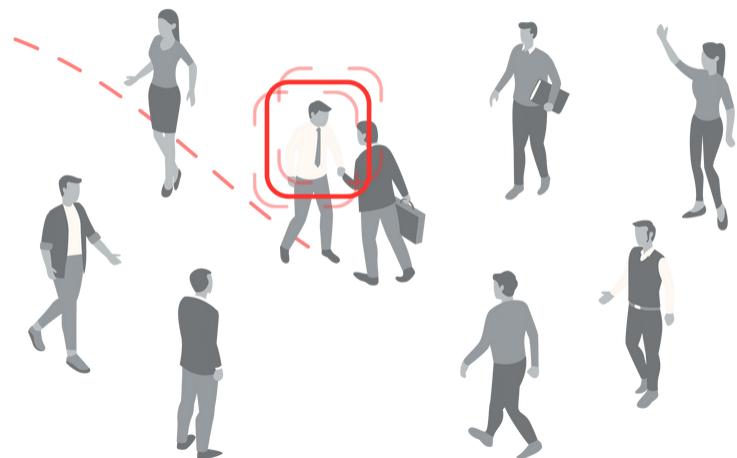
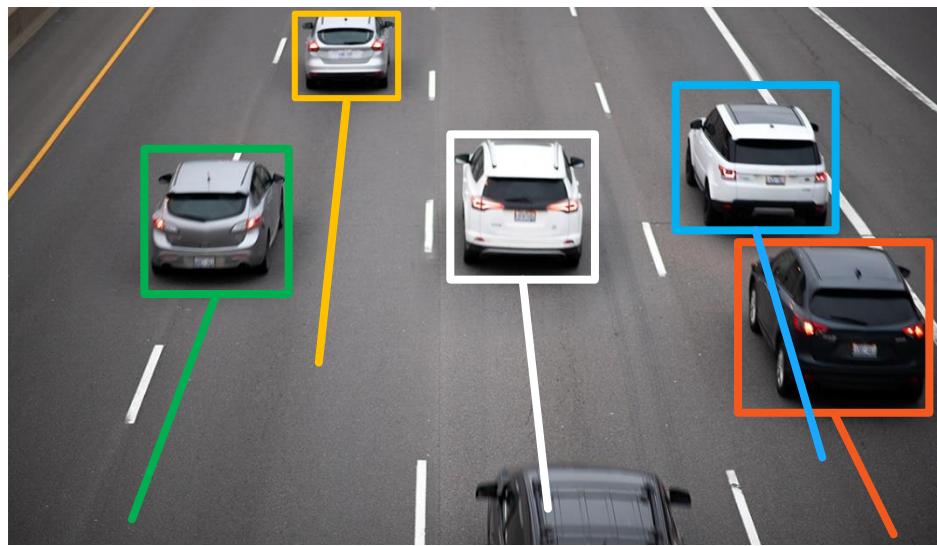
Object Detection

Color	Class
•	Person
•	Car
•	Cat
•	Dog
•	Duck



Computer Vision

Object Tracking



Computer Vision

Image Retrieval

Query



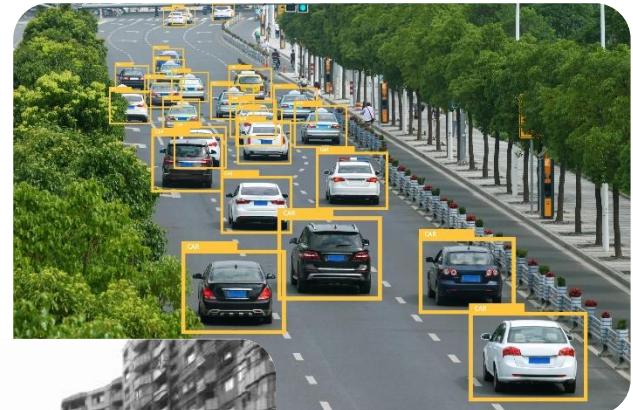
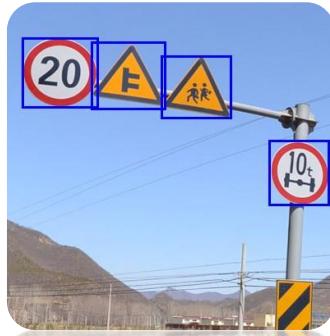
Answer



Computer Vision

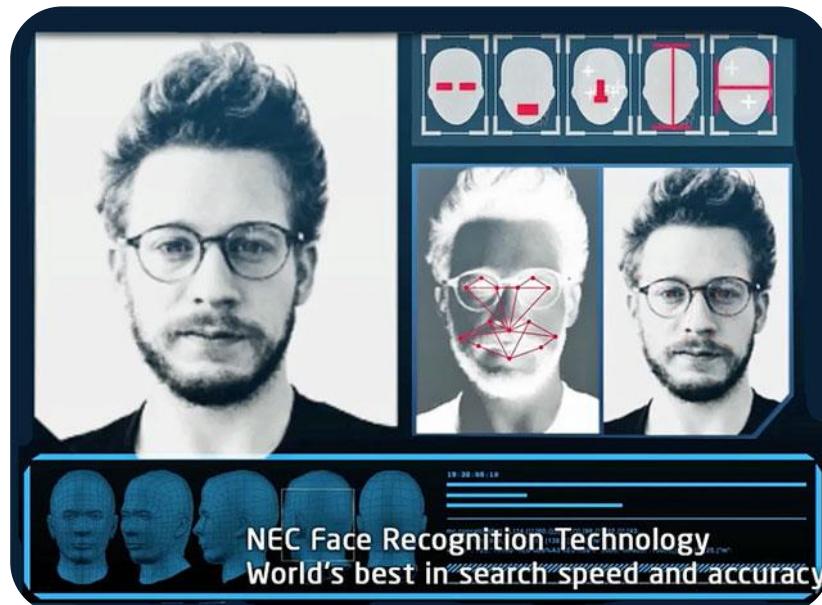
Applications of CV – Autonomous Vehicles

- ▶ Vehicle Detection
- ▶ Pedestrian Detection
- ▶ Traffic Sign Recognition
- ▶ Lane Markers Detection



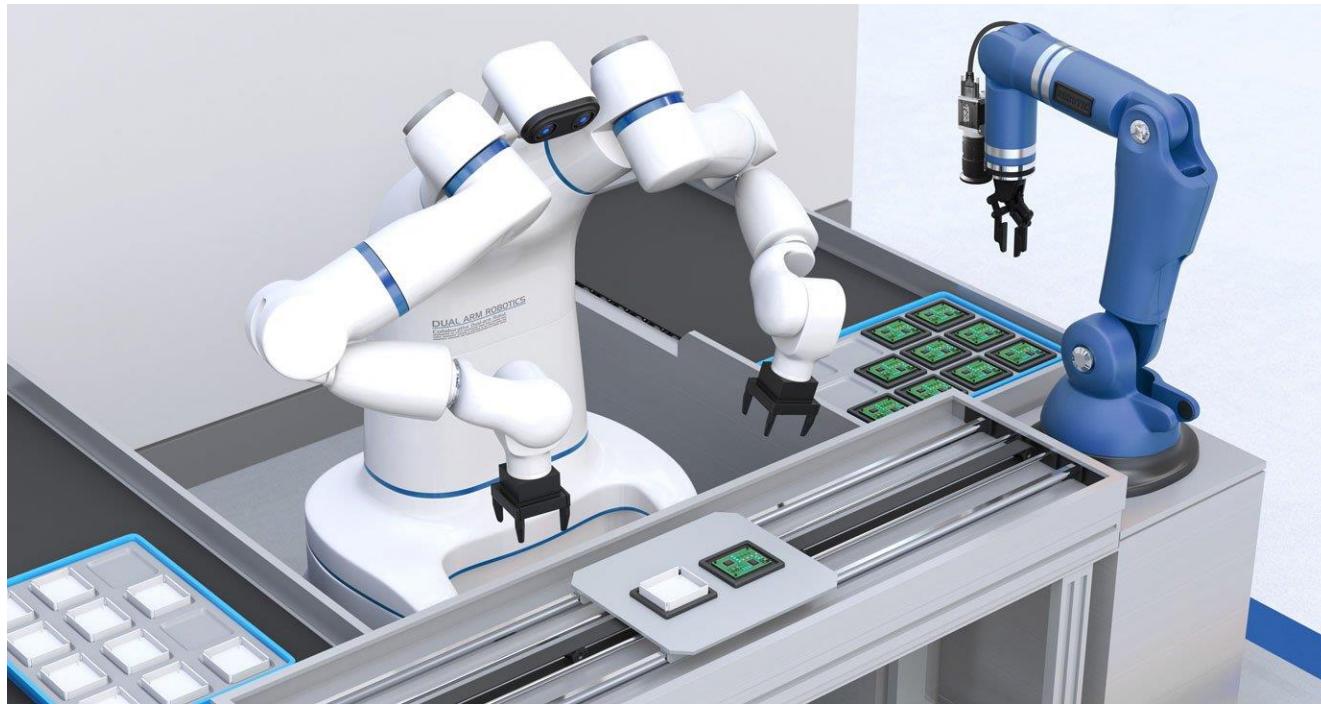
Computer Vision

Applications of CV – Face Recognition



Computer Vision

Applications of CV – Robotics



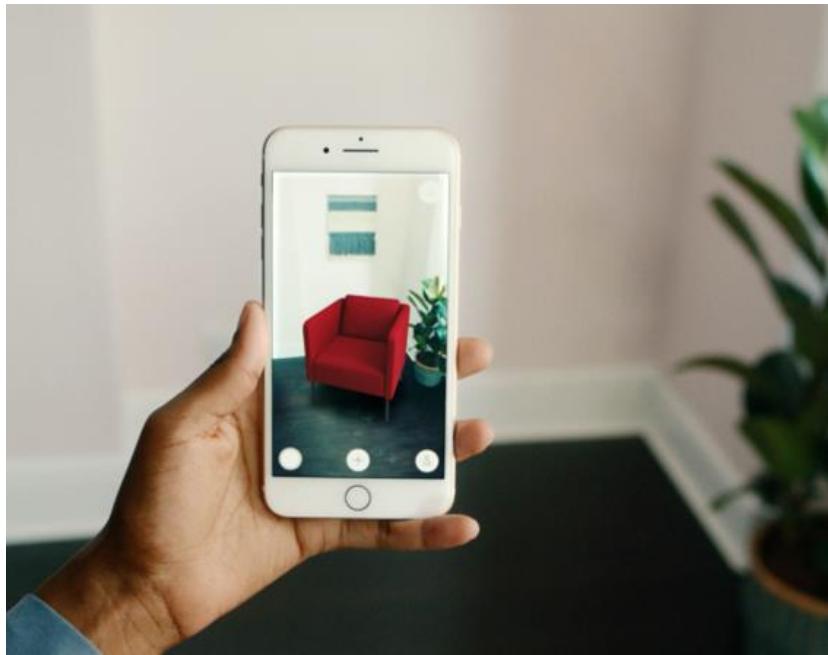
Computer Vision

Applications of CV – Google Translate



Computer Vision

Applications of CV – Mobile Computing



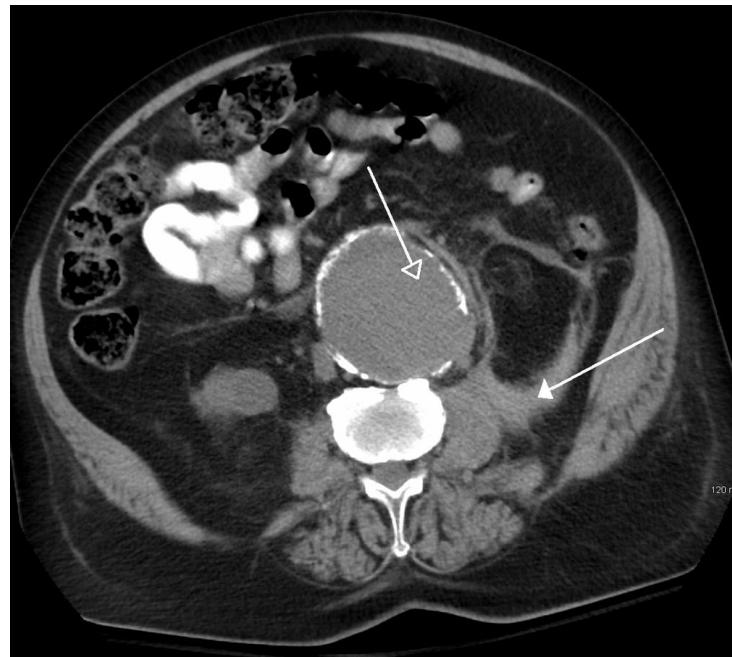
Computer Vision

Applications of CV – Abandoned Object Detection



Computer Vision

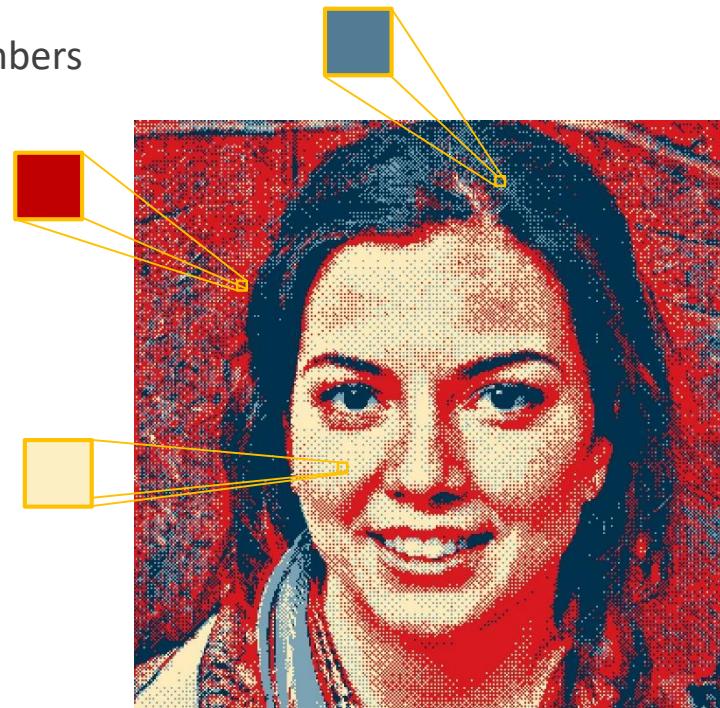
Applications of CV – Medicine



Feature Extraction

How do computers understand digital images?

- ▶ Digital images
 - ▶ Representations of real image as a set of numbers
 - ▶ Divided into small areas, AKA pixels
 - ▶ Pixels hold a small set of numbers for:
 - ▶ Color
 - ▶ Brightness
 - ▶ Intensity of light



Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



Binary



Grayscale

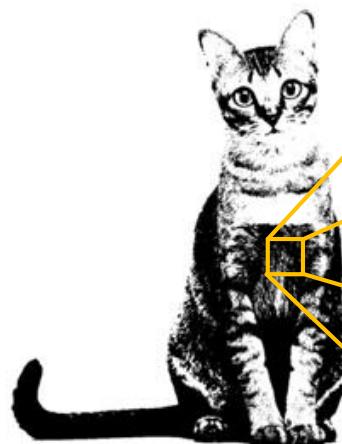


Color

Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



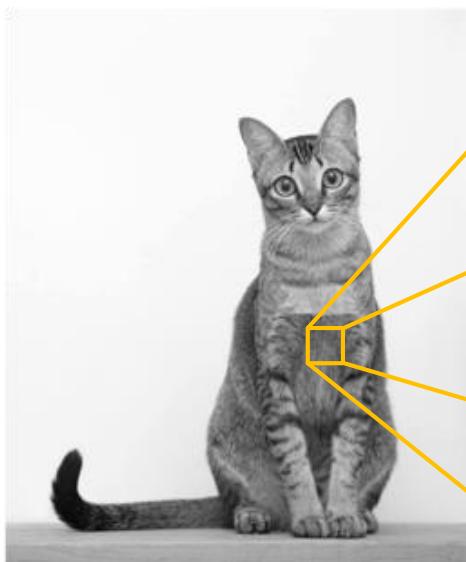
Binary image

1	1	1	1	1	0	1	1	0	0
1	1	1	0	0	0	0	0	1	1
1	1	1	1	0	0	0	0	1	1
1	1	1	0	0	0	0	0	1	1
1	1	0	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	1	1
1	1	1	1	1	0	0	0	0	0
1	1	1	0	0	0	0	0	0	0

Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



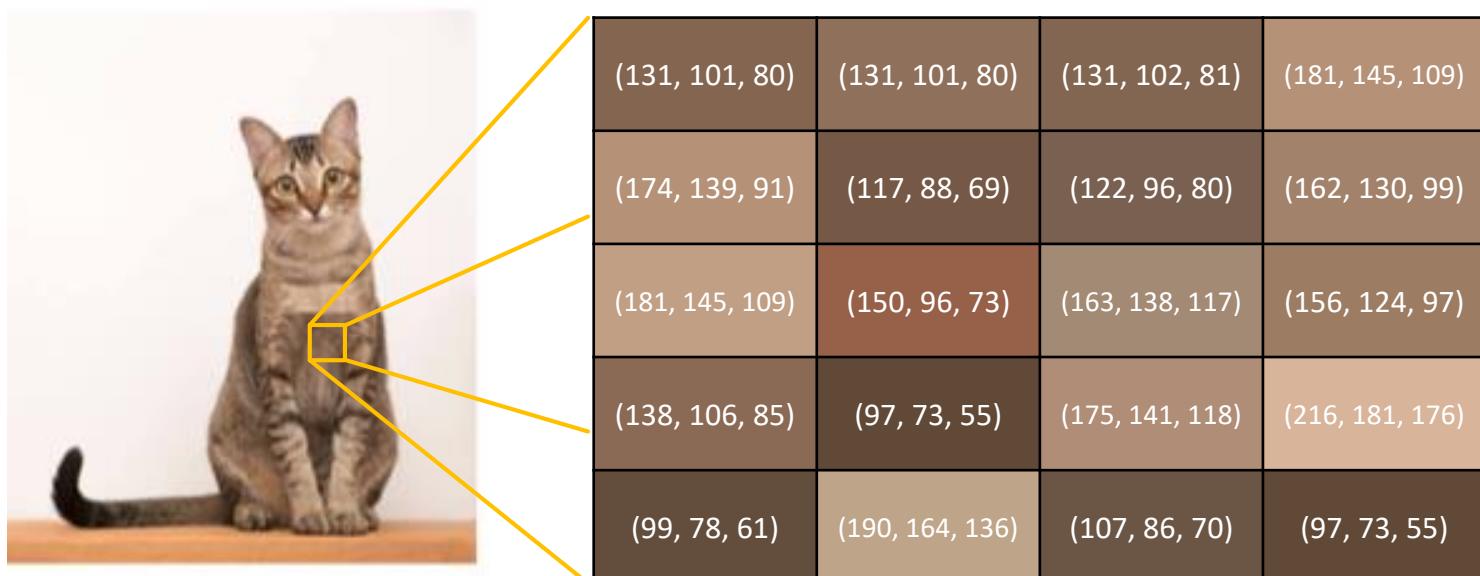
Grayscale image

121	99	78	69	129	194	93	182	81	42
78	92	188	96	95	189	204	219	94	67
105	182	79	94	79	199	219	241	201	95
185	178	98	65	132	14	18	192	182	90
98	182	201	92	94	182	194	182	71	95
84	95	81	95	102	201	29	95	185	206
96	47	173	82	96	195	181	148	103	97
98	57	142	95	84	197	61	184	73	93

Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images

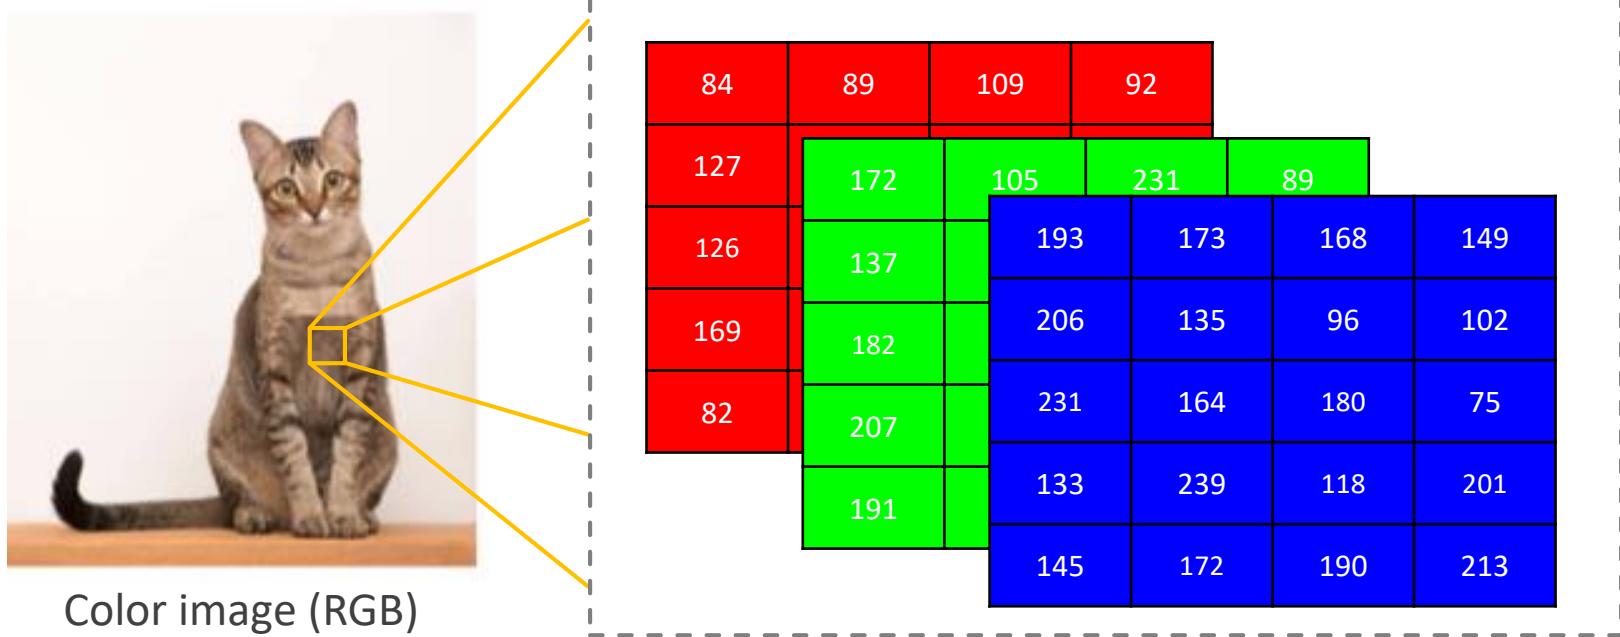


Color image (RGB)

Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



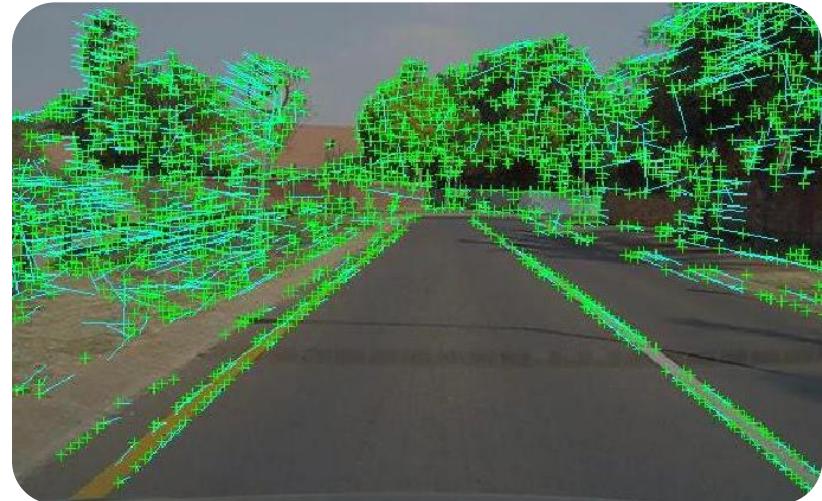
Feature Extraction

What are features of an image?

- ▶ Visual parts or patterns in an image
- ▶ Useful to identify it and extract information from it

Where to find the features?

- ▶ Corners
- ▶ Edges
- ▶ Regions of Interest (RoI) points
- ▶ Etc.



Feature Extraction

Where to find these features?



Feature Extraction

Where to find these features?



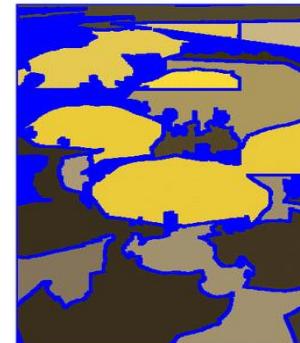
Corners

Feature Extraction

Where to find these features?



Regions



Feature Extraction

Looking for features using the pixel representation

- ▶ Low-level features
 - ▶ Capturing local appearance/texture statistics of objects
 - ▶ Highly effective in classification tasks
- ▶ High-level features
 - ▶ Semantic concepts that can be interpreted in a scene
 - ▶ The probability of observing an object in an image

Hierarchy of Features

Low-level features

- Algorithms: SIFT, SURF, etc.
- Interest points, edges, etc.



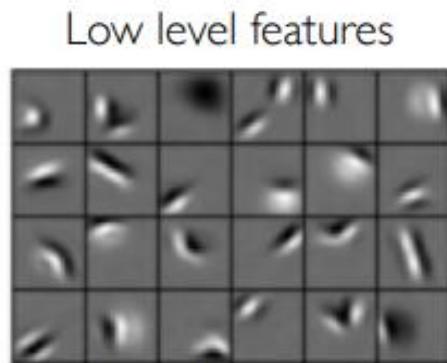
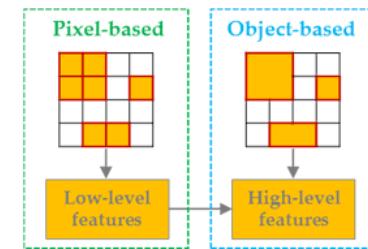
High-level features

- Mainly used in CV
- Birds, leaves

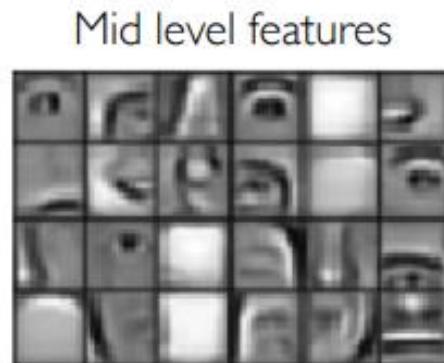
Feature Extraction

Looking for features using the pixel representation

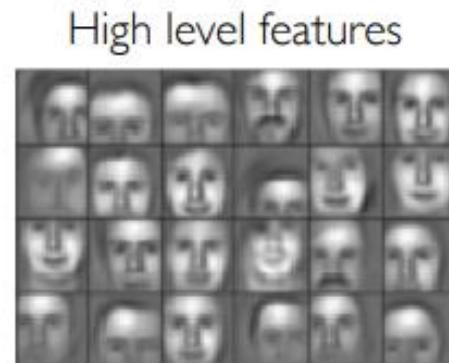
- ▶ Low-level and High-level features
 - ▶ Recall: Session#1 - Basics



Low level features
Edges, dark spots



Mid level features
Eyes, ears, nose



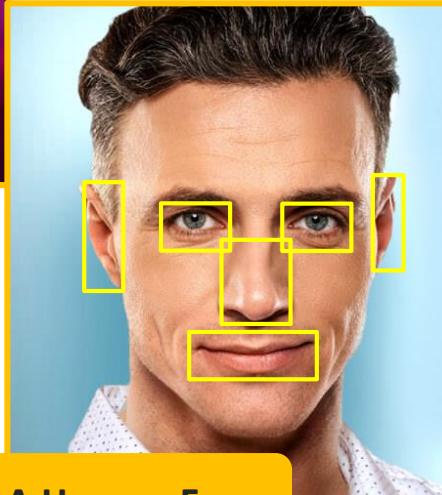
High level features
Facial structure

Feature Extraction

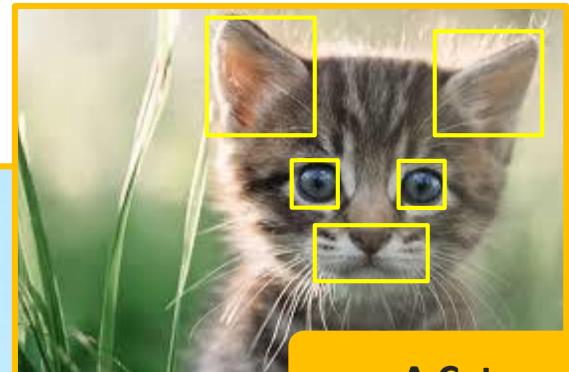
What are the high-level features that classify these images?



A Car



A Human Face



A Cat

Feature Extraction

What are the high-level features that classify these images?



Playing football

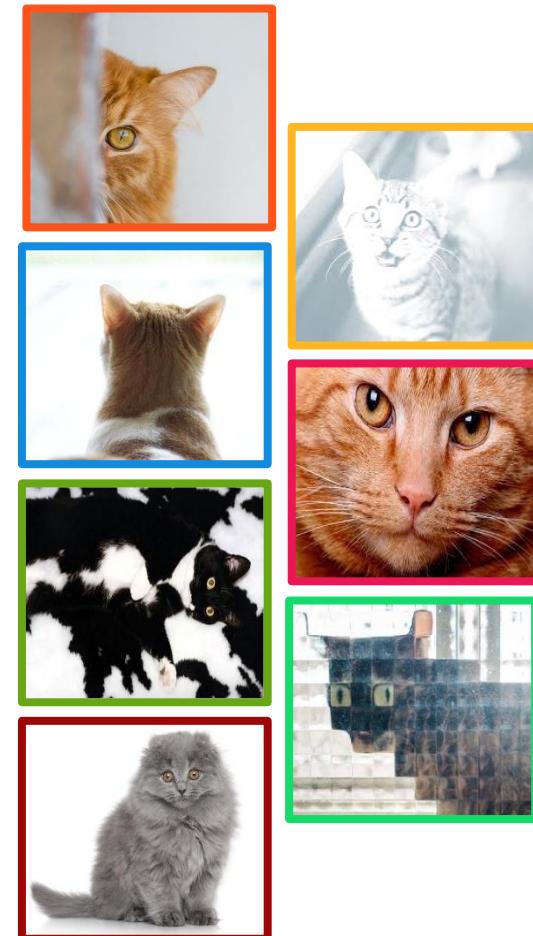


Playing basketball

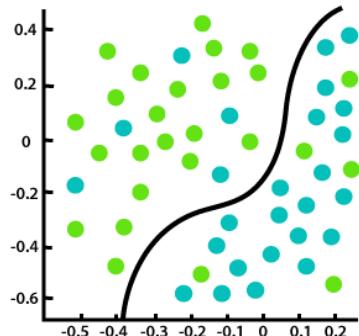
Feature Extraction



What are the challenges?



Classification vs. Regression



Classification

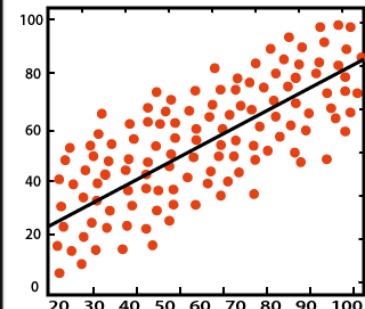
- Predicting a label
- Outputs are discrete
- Outputs = labels
- Input values can be continuous

Use case: forecasting the weather (classes: sunny, rainy, cloudy, snowy)

Regression

- Predicting a quantity
- Outputs are continuous (e.g. an integer or floating point value)
- Outputs can be converted into labels to present a classification problem

Use case: predicting the value of a car two years later



Classification vs. Regression

Classification in Computer Vision



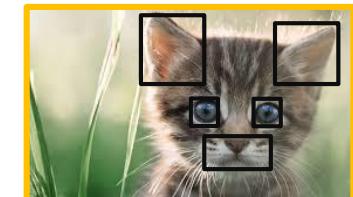
Input image



Pixel representation

121	99	78	69	129	194	93
78	92	188	96	95	189	204
105	182	79	94	79	199	219
185	178	98	65	132	14	18
98	182	201	92	94	182	194
84	95	81	95	102	201	29
96	47	173	82	96	195	181
98	57	142	95	84	197	61

....



Unique features of a cat

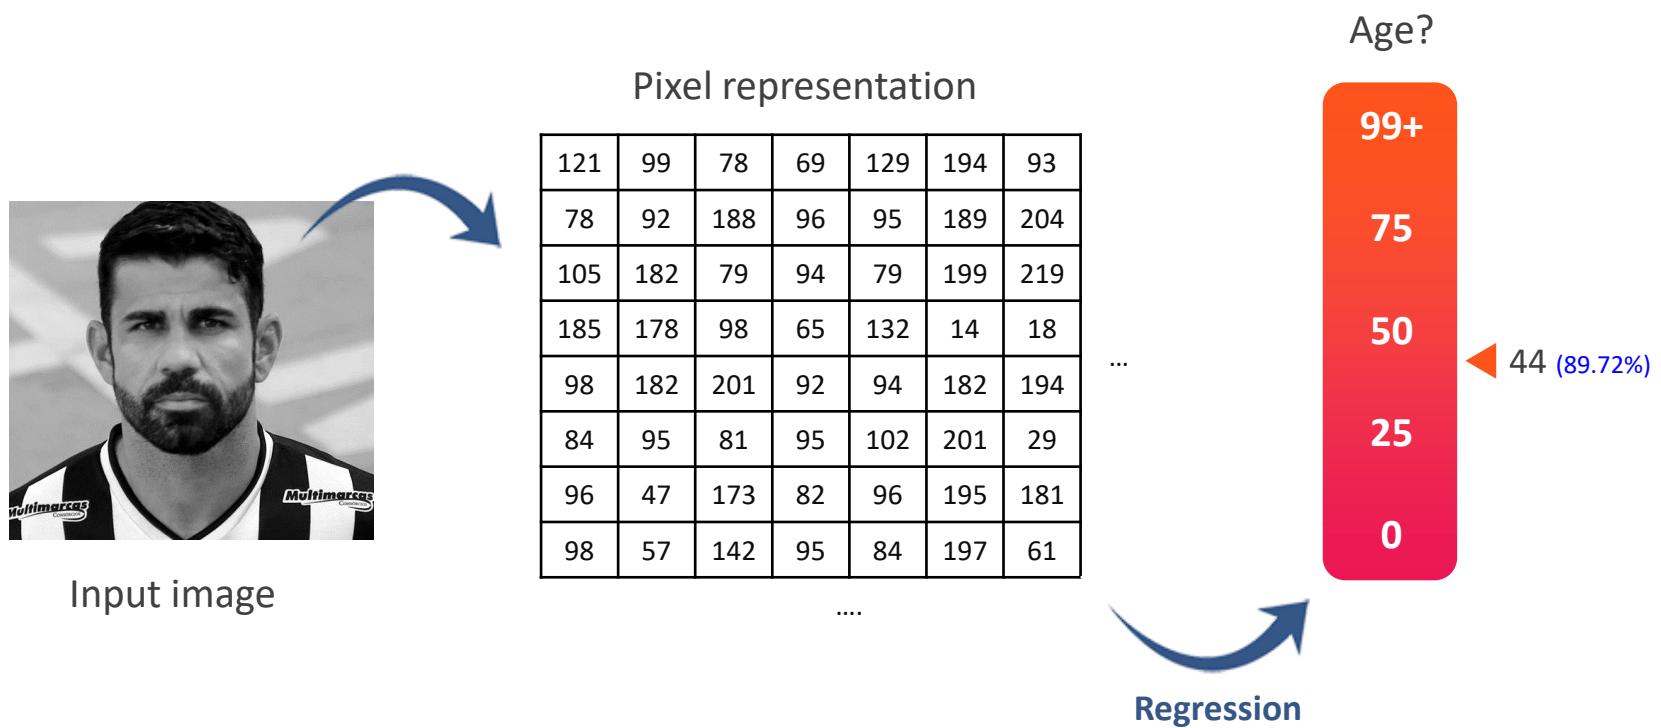
Class	Probability	
Cat	0.85	✓
Dog	0.09	✗
Mouse	0.06	✗



Classification

Classification vs. Regression

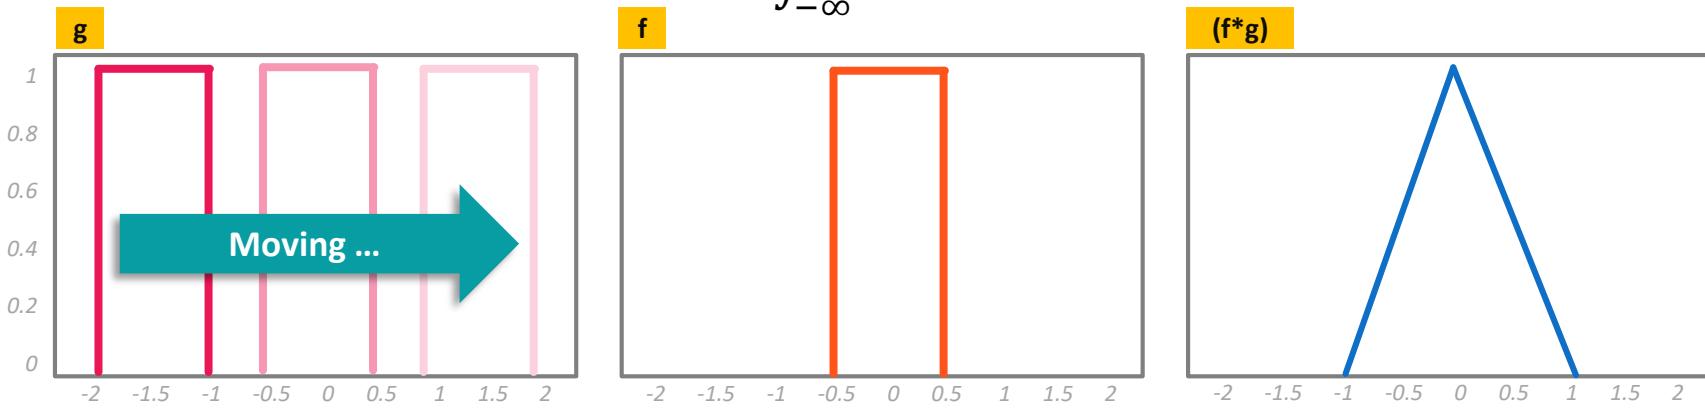
Regression in Computer Vision



Convolution

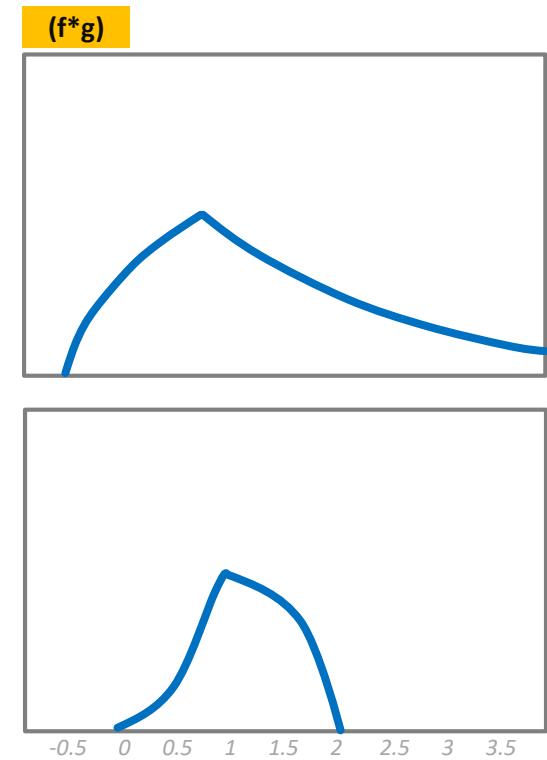
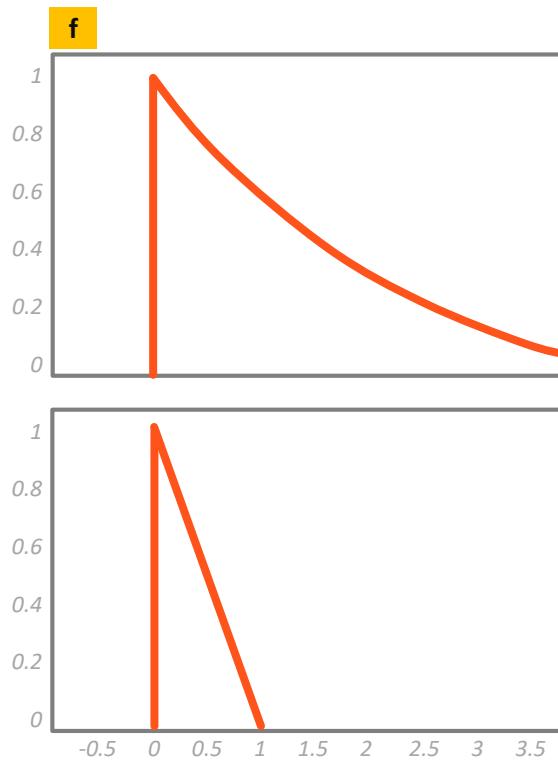
- ▶ A mathematical operation on **two functions (f and g)** to produce a **third one**
- ▶ **Result:** Identifying how f can modify the shape of g
- ▶ **Definition:** an integral that expresses the amount of overlap of f as it is shifted over g

$$[f * g](t) \equiv \int_{-\infty}^{+\infty} f(\tau)g(t - \tau). d\tau$$



Convolution

$$[f * g](t) \equiv \int_{-\infty}^{+\infty} f(\tau)g(t - \tau). d\tau$$



Convolution

What about applying convolution on images?

I(0,0)	I(1,0)	I(2,0)	I(3,0)	I(4,0)	I(5,0)	I(6,0)
I(0,1)	I(1,1)	I(2,1)	I(3,1)	I(4,1)	I(5,1)	I(6,1)
I(0,2)	I(1,2)	I(2,2)	I(3,2)	I(4,2)	I(5,2)	I(6,2)
I(0,3)	I(1,3)	I(2,3)	I(3,3)	I(4,3)	I(5,3)	I(6,3)
I(0,4)	I(1,4)	I(2,4)	I(3,4)	I(4,4)	I(5,4)	I(6,4)
I(0,5)	I(1,5)	I(2,5)	I(3,5)	I(4,5)	I(5,5)	I(6,5)
I(0,6)	I(1,6)	I(2,6)	I(3,6)	I(4,6)	I(5,6)	I(6,6)

Input image



H(0,0)	H(1,0)	H(2,0)
H(0,1)	H(1,1)	H(2,1)
H(0,2)	H(1,2)	H(2,2)

Filter/Kernel/Mask

O(0,0)				

Output image
(feature map)

Convolution

What about applying convolution on images?

1	0	1	0	1
1	1	1	0	0
1	0	1	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} ? \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image
(feature map)

Convolution

What about applying convolution on images?

1 <small>$\times 1$</small>	0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	0	1
1 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1 <small>$\times 1$</small>	0	0
1 <small>$\times 1$</small>	0 <small>$\times 0$</small>	1 <small>$\times 1$</small>	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 \\ \text{[empty]} \\ \text{[empty]} \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

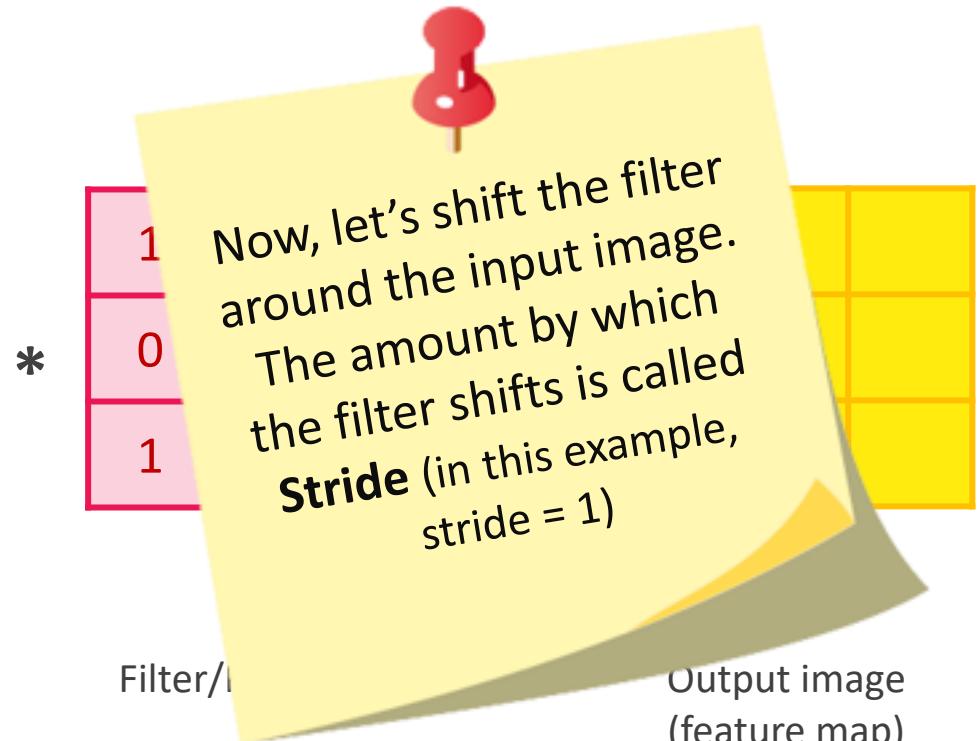
Output image
(feature map)

Convolution

What about applying convolution on images?

1 $\times 1$	0 $\times 0$	1 $\times 1$	0	1
1 $\times 0$	1 $\times 1$	1 $\times 1$	0	0
1 $\times 1$	0 $\times 0$	1 $\times 1$	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)



Convolution

What about applying convolution on images?

1	0 <small>$\times 1$</small>	1 <small>$\times 0$</small>	0 <small>$\times 1$</small>	1
1	1 <small>$\times 0$</small>	1 <small>$\times 1$</small>	0 <small>$\times 1$</small>	0
1	0 <small>$\times 1$</small>	1 <small>$\times 0$</small>	1 <small>$\times 1$</small>	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 & 2 \\ & \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image
(feature map)

Convolution

What about applying convolution on images?

1	0	1 _{x1}	0 _{x0}	1 _{x1}
1	1	1 _{x0}	0 _{x1}	0 _{x1}
1	0	1 _{x1}	1 _{x0}	0 _{x1}
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 & 2 & 3 \\ & & \\ & & \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image
(feature map)

Convolution

What about applying convolution on images?

1	0	1	0	1
1	1	1	0	0
1	0	1	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 & 2 & 3 \\ 4 & 3 & 4 \\ 4 & 4 & 3 \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image
(feature map)

Convolution

What about applying convolution on images?

1 × 1	1 × 0	1 × 1	0	0
0 × 0	1 × 1	1 × 0	1	0
0 × 1	0 × 0	1 × 1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved feature

1	1 × 1	1 × 0	0 × 1	0
0	1 × 0	1 × 1	1 × 0	0
0	0 × 1	1 × 0	1 × 1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	

Convolved feature

1	1	1 × 1	0 × 0	0 × 1
0	1	1 × 0	1 × 1	0 × 0
0	0	1 × 1	1 × 0	1 × 1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved feature

1	1	1	0	0
0 × 1	1 × 0	1 × 1	1	0
0 × 0	0 × 1	1 × 0	1	1
0 × 1	0 × 0	1 × 1	1	0
0	1	1	0	0

Image

4	3	4

Convolved feature

1	1	1	0	0
0	1 × 1	1 × 0	1 × 1	0
0	0 × 0	1 × 1	1 × 0	1
0	0 × 1	1 × 0	1 × 1	0
0	1	1	0	0

Image

4	3	4

Convolved feature

1	1	1	0	0
0	1	1 × 1	1 × 0	0 × 1
0	0	1 × 0	1 × 1	1 × 0
0	0	1 × 1	1 × 0	0 × 1
0	1	1	0	0

Image

4	3	4

Convolved feature

Convolution

What about applying convolution on images?



Input image

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9



Simple box blur

-1	-1	-1
2	2	2
-1	-1	-1



Horizontal lines

Convolution

What about applying convolution on images?



Input image

$$\begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline -1 & 8 & -1 \\ \hline -1 & -1 & -1 \\ \hline \end{array}$$



Edge detection

$$\begin{array}{|c|c|c|} \hline 0 & -1 & 0 \\ \hline -1 & 4 & -1 \\ \hline 0 & -1 & 0 \\ \hline \end{array}$$

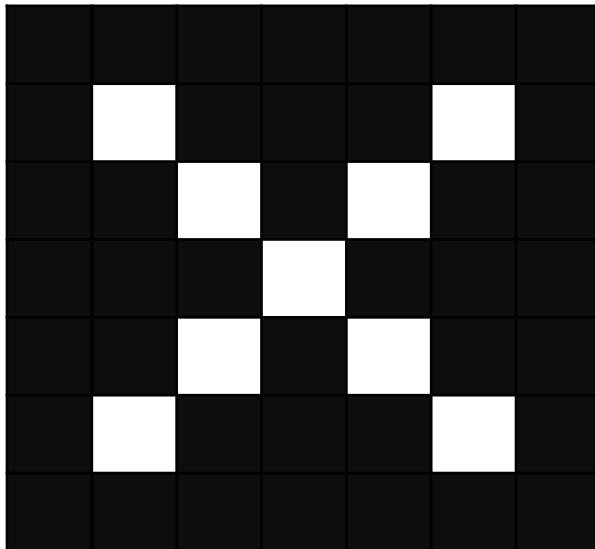


The Laplacian operator

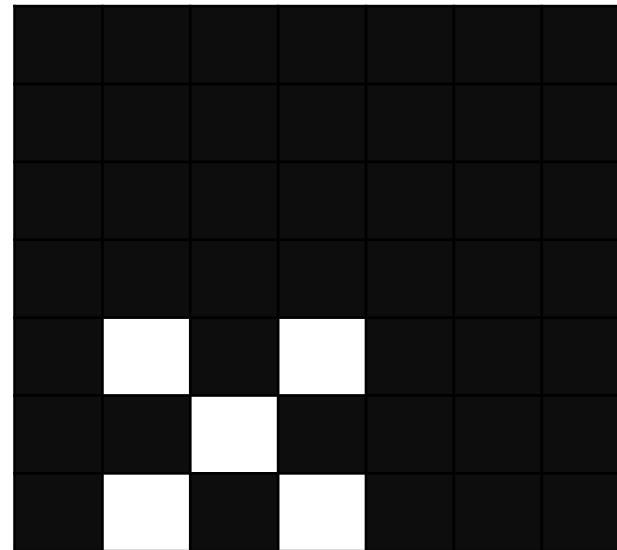
Convolution

Why is it important to use Convolution?

- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



The letter “X” in a binary image

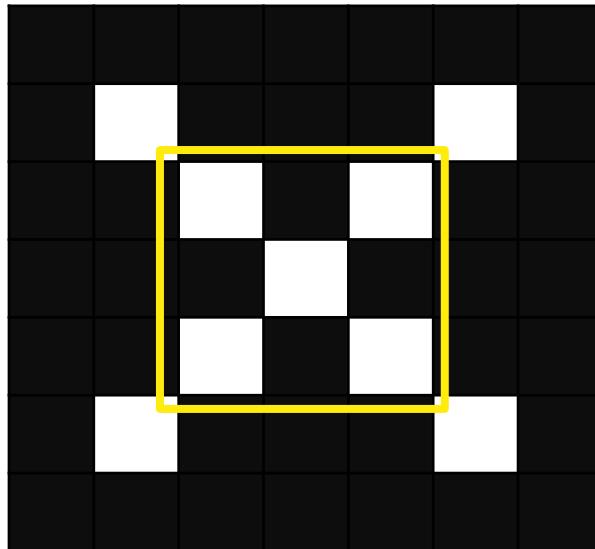


The scaled version

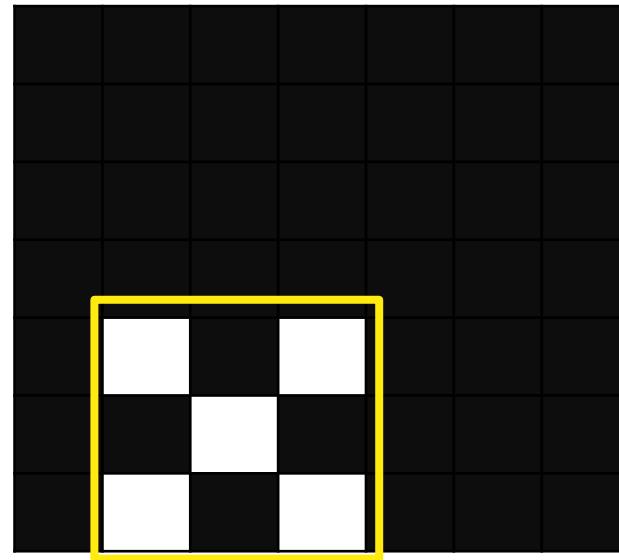
Convolution

Why is it important to use Convolution?

- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



The letter “X” in a binary image



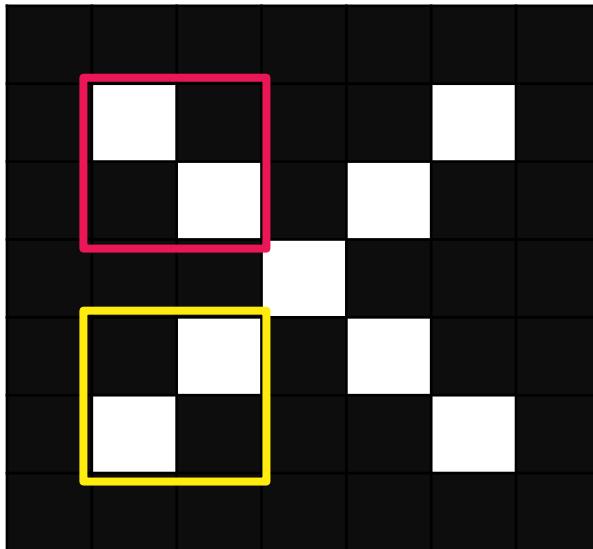
The scaled version

The same features!

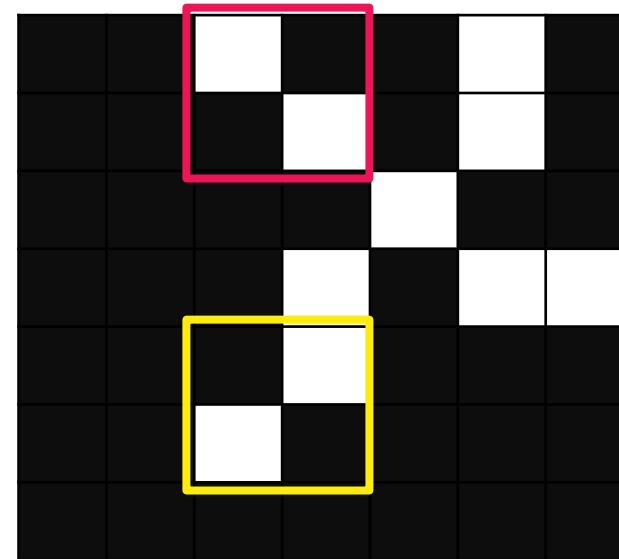
Convolution

Why is it important to use Convolution?

- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



The letter “X” in a binary image



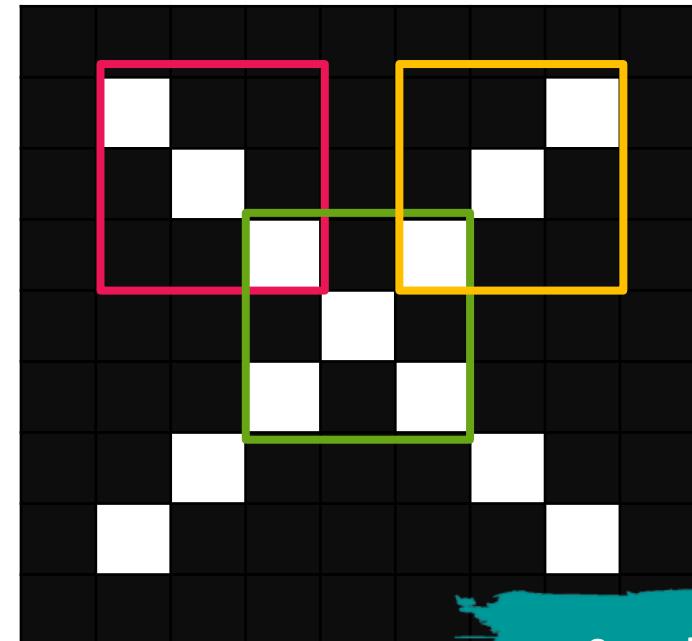
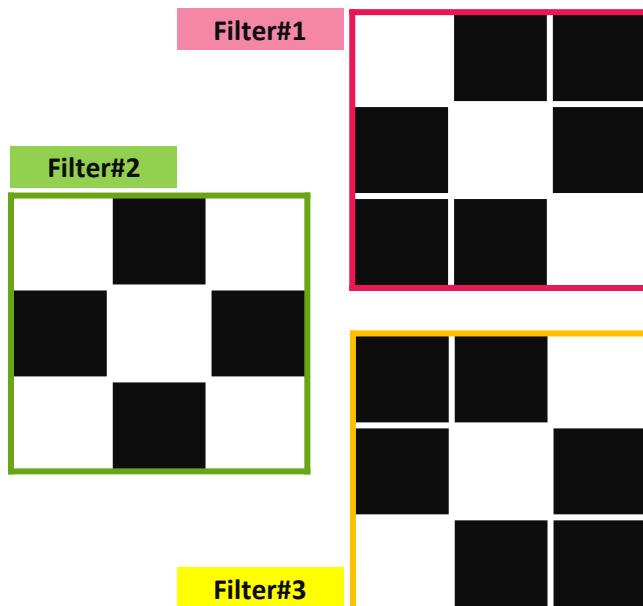
The scaled version

The same features!

Convolution

Why is it important to use Convolution?

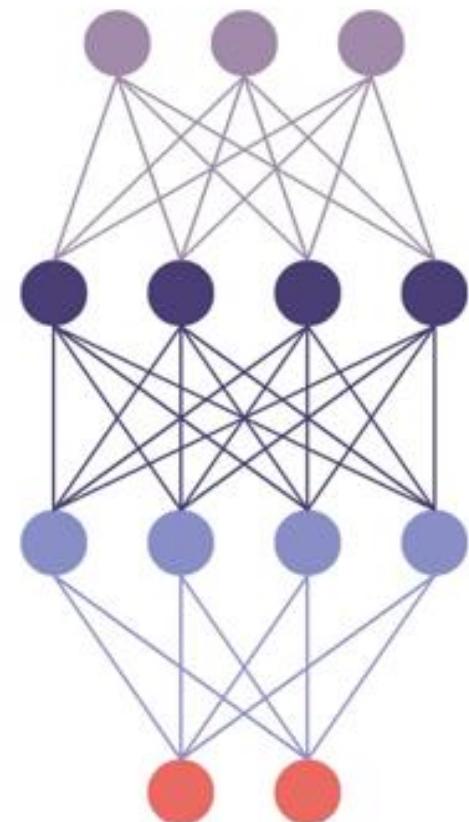
- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



Sample filters

Convolutional Neural Networks

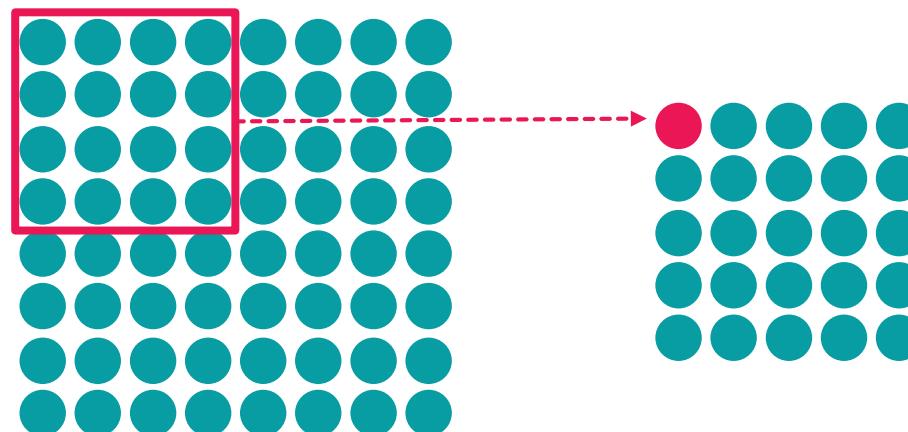
- ▶ Fully-connected Neural Networks (Dense NNs)
 - ▶ Recall: [Session#1 - Basics](#)
- ▶ Can we use them for processing mage inputs?
 - ▶ Regular NNs do not **scale well** to full images
 - ▶ We need many **calculations** and parameters!
 - ▶ We will lose the **spatial information**
 - ▶ **Reason:** flattening the 2D image into 1D vector
 - ▶ **Example:** the nearness of pixels to others



Convolutional Neural Networks

How to preserve the spatial structure?

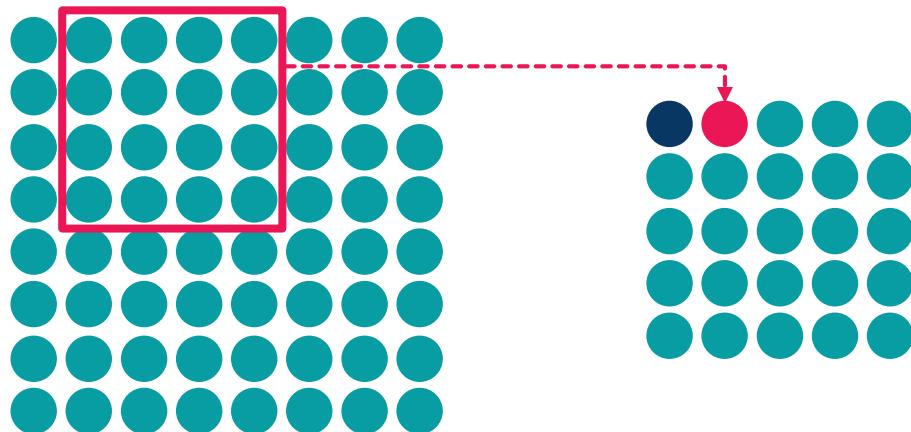
- ▶ Connecting **slices of the input**, instead of connecting all input values to neurons in the hidden layer
 - ▶ Each patch (slice) to a single neuron in the next layer
- ▶ This approach preserves the spatial information + visual features



Convolutional Neural Networks

How to preserve the spatial structure?

- ▶ Connecting **slices of the input**, instead of connecting all input values to neurons in the hidden layer
 - ▶ Each patch (slice) to a single neuron in the next layer
- ▶ This approach preserves the spatial information + visual features

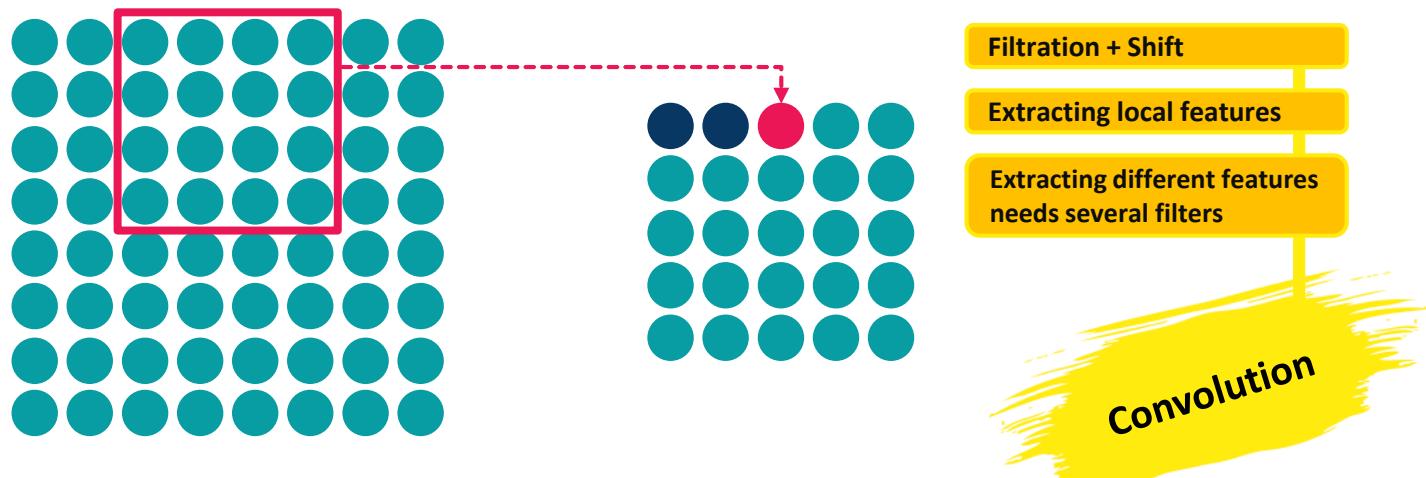


Simply, a sliding window can do the rest!

Convolutional Neural Networks

How to preserve the spatial structure?

- ▶ Connecting **slices of the input**, instead of connecting all input values to neurons in the hidden layer
 - ▶ Each patch (slice) to a single neuron in the next layer
- ▶ This approach preserves the spatial information + visual features

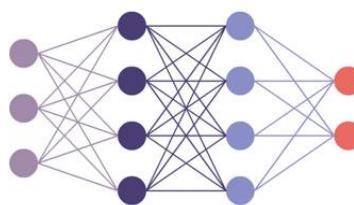


Convolutional Neural Networks

- ▶ CNNs (*ConvNets*) utilize the mathematical foundation of convolution for **computer vision** tasks
- ▶ Inspired by the organization of the Visual Cortex
- ▶ Much lower **pre-processing** comparing to other classification algorithms
 - ▶ CNNs can automatically learn the filters and characteristics
- ▶ The overall process in a CNN:



Input image



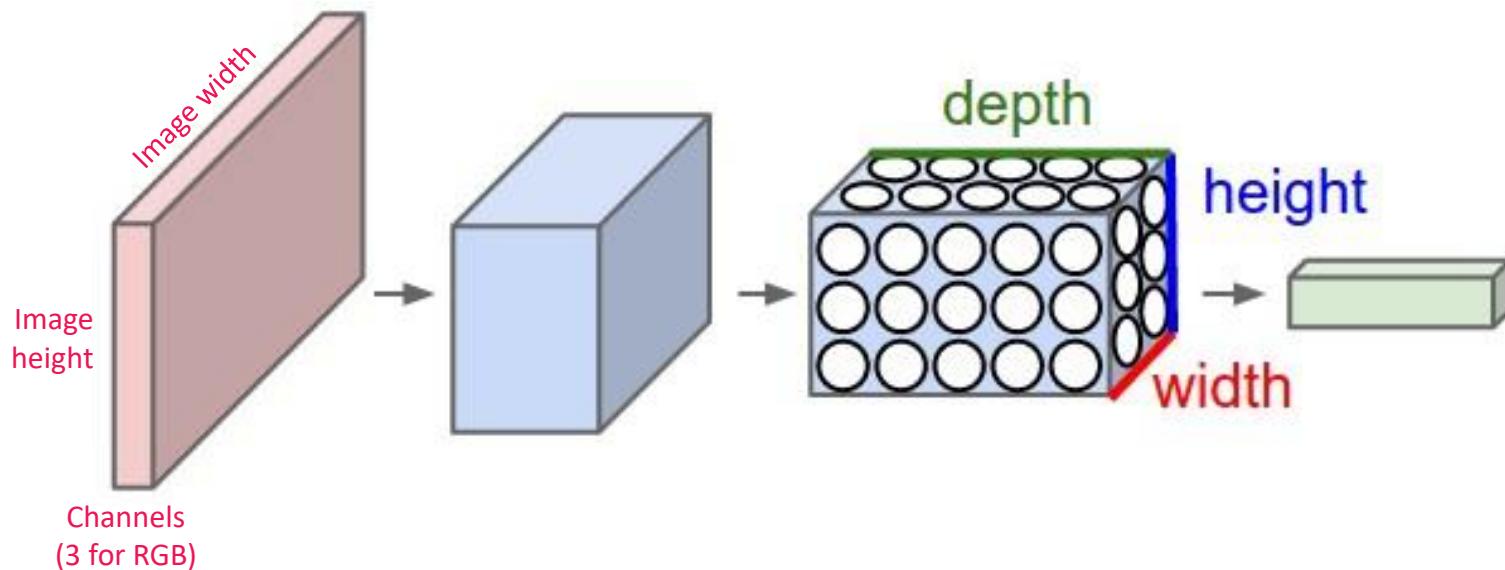
Learnable weights
and biases



Differentiation of
objects in the image

Convolutional Neural Networks

- ▶ In contrast with **regular NNs**, the layers of a CNN have neurons arranged in **three dimensions**: **width, height, depth**
 - ▶ **Note:** “depth” is the third dimension of an activation volume
 - ▶ **Note:** each layer converts a 3D input to a 3D output of neurons’ AFs



Convolutional Neural Networks

Roles of CNNs

- ▶ Assuming that the inputs are images → more efficient network
- ▶ Reducing images into a form that is **easier for processing**
- ▶ Preserving the **critical features** of images for prediction
- ▶ Reducing the **complexity** of the image classification task

Popular Architectures of CNNs **(Click to see the papers)**

AlexNet

GoogLeNet

VGGNet

LeNet

ResNet

Convolutional Neural Networks

21 lines (16 sloc) | 1.1 KB

Convolutional Neural Networks (CNNs)

CNNs (ConvNets) utilize the mathematical foundation of convolution for computer vision tasks. Inspired by the organization of the Visual Cortex, CNNs can provide much lower pre-processing comparing to other classification algorithms. These networks can automatically learn the filters and characteristics of the input images. In contrast with regular NNs, the layers of a CNN have neurons arranged in three dimensions: width, height, depth.

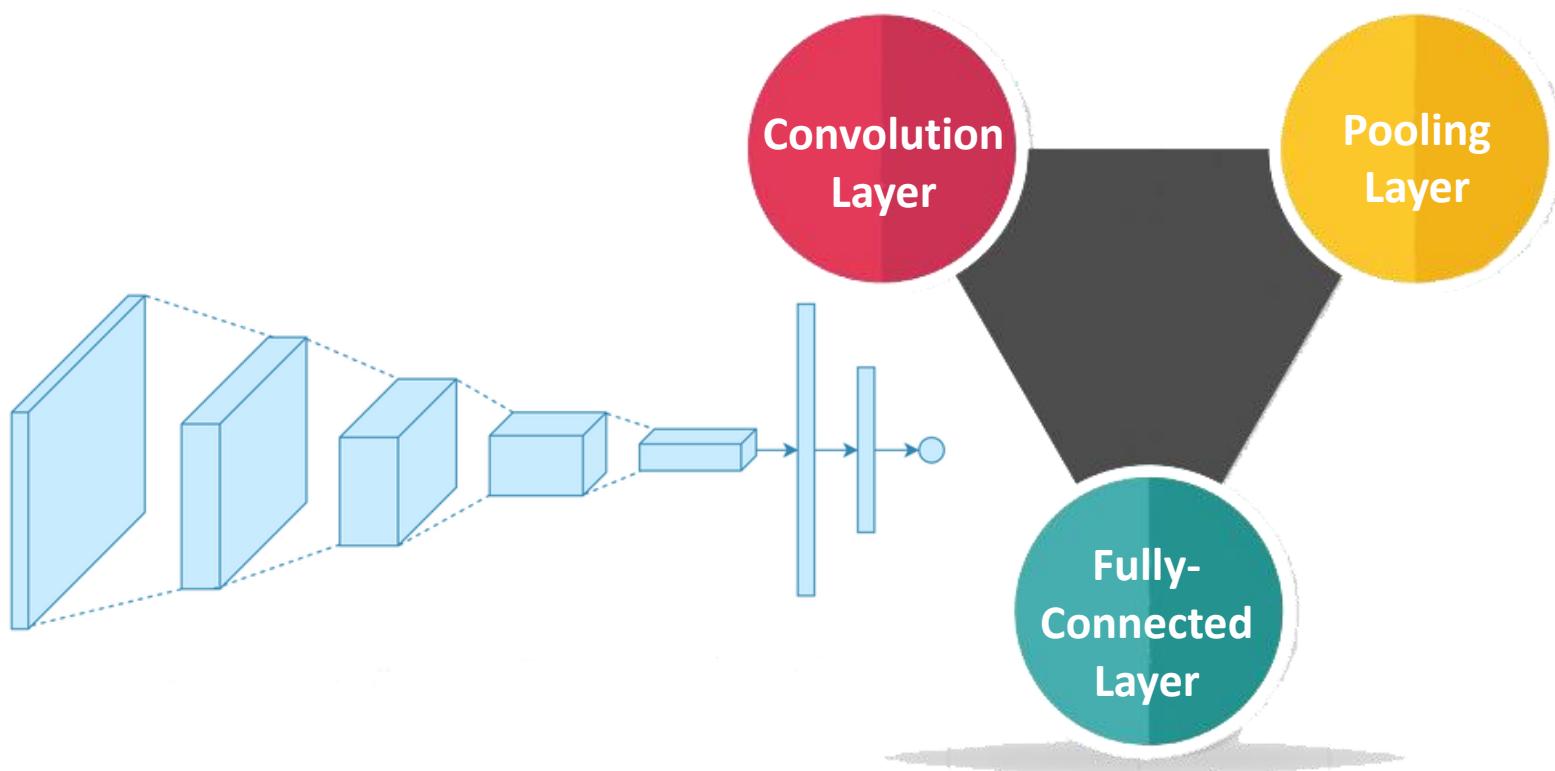
The diagram illustrates a sequence of four layers in a CNN. The first layer is labeled 'Input' and shows a stack of blue and green 3D blocks. The second layer is labeled 'Conv + ReLU' and shows a stack of green 3D blocks. The third layer is labeled 'Pooling' and shows a stack of red 3D blocks. The fourth layer is labeled 'Flattening' and shows a single yellow 3D block. Below the diagram is a teal button with white text.

Full code on GitHub

Codes

A large black outline of a hand cursor is positioned over the teal button labeled 'Full code on GitHub', pointing towards it.

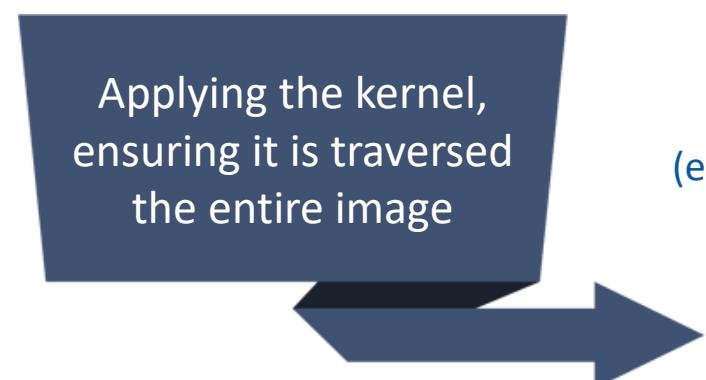
CNNs Layers



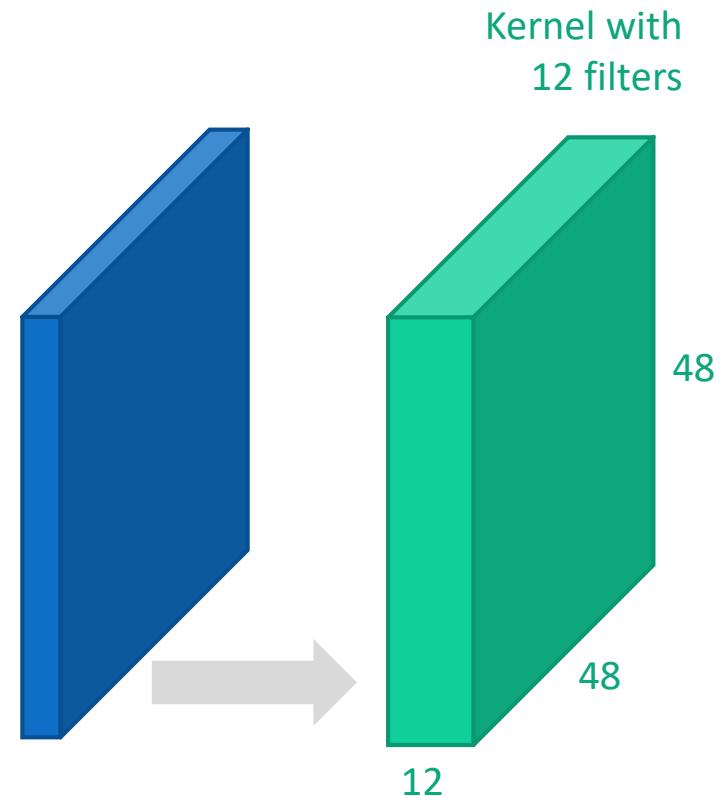
CNNs Layers

Convolution Layer (Kernel)

- ▶ **Output:** Feature maps
- ▶ Extracting high-level features (e.g., edges)
- ▶ Specifying the characteristics of the Kernel
 - ▶ Including Kernel size, Stride, Padding



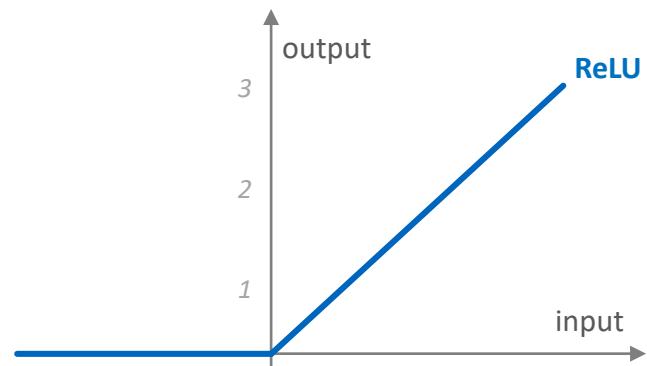
Input Image
(e.g. [48× 48×3])



CNNs Layers

Convolution Layer (Kernel)

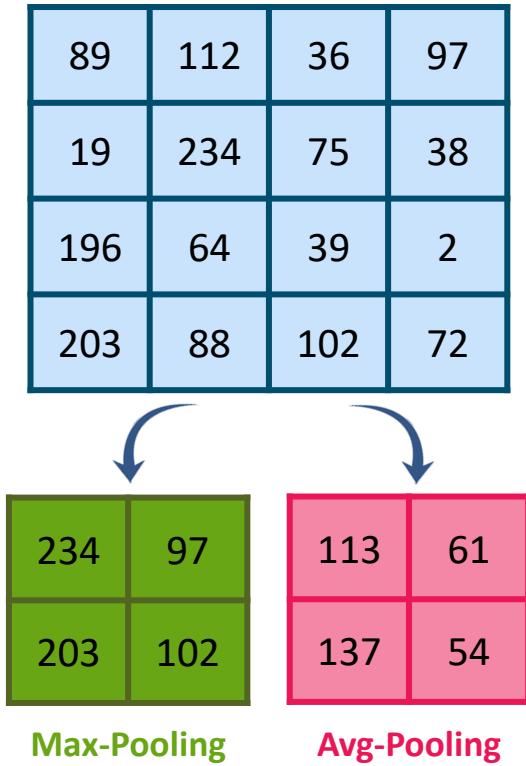
- ▶ Then, adding a supplementary step to the convolution operation
 - ▶ **Activation Function:** often **ReLU** (replacing negative values by zero)
 - ▶ **Note:** The Rectified Linear Unit is not a separate component of the convolutional neural networks' process



CNNs Layers

Pooling Layer (Down sampling)

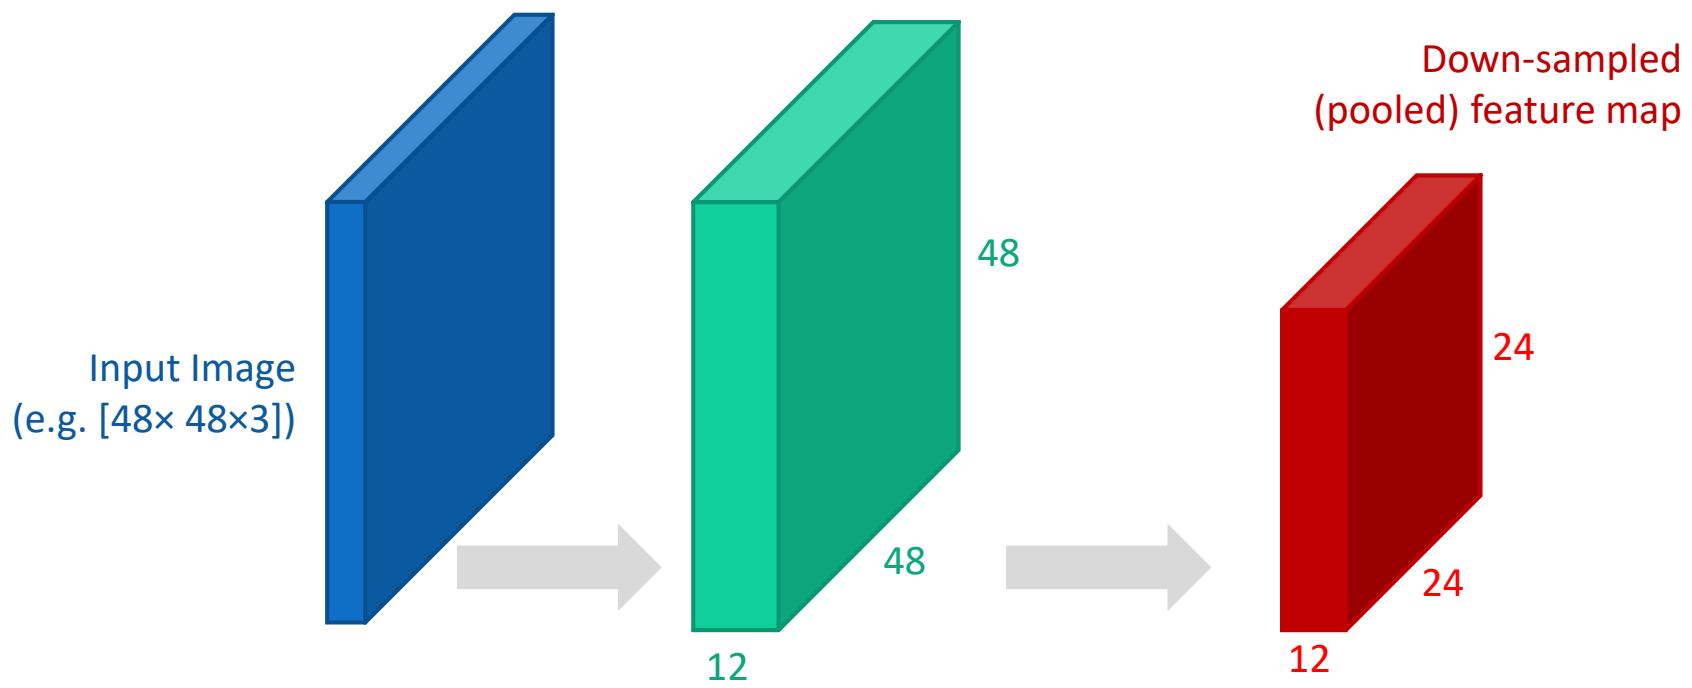
- ▶ **Output:** Pooled feature maps
- ▶ With a Convolution layer, forms a layer in a CNN
- ▶ Reducing the spatial size of the Convolved Feature
- ▶ Decreasing the dimensionality and computation
- ▶ Two-types of Pooling:
 - ▶ **Max-Pooling:** returns the maximum value
 - ▶ Recommended due to de-noising advantage
 - ▶ **Average-Pooling:** returns the average of all the values



CNNs Layers

Pooling Layer (Down sampling)

Kernel with
12 filters



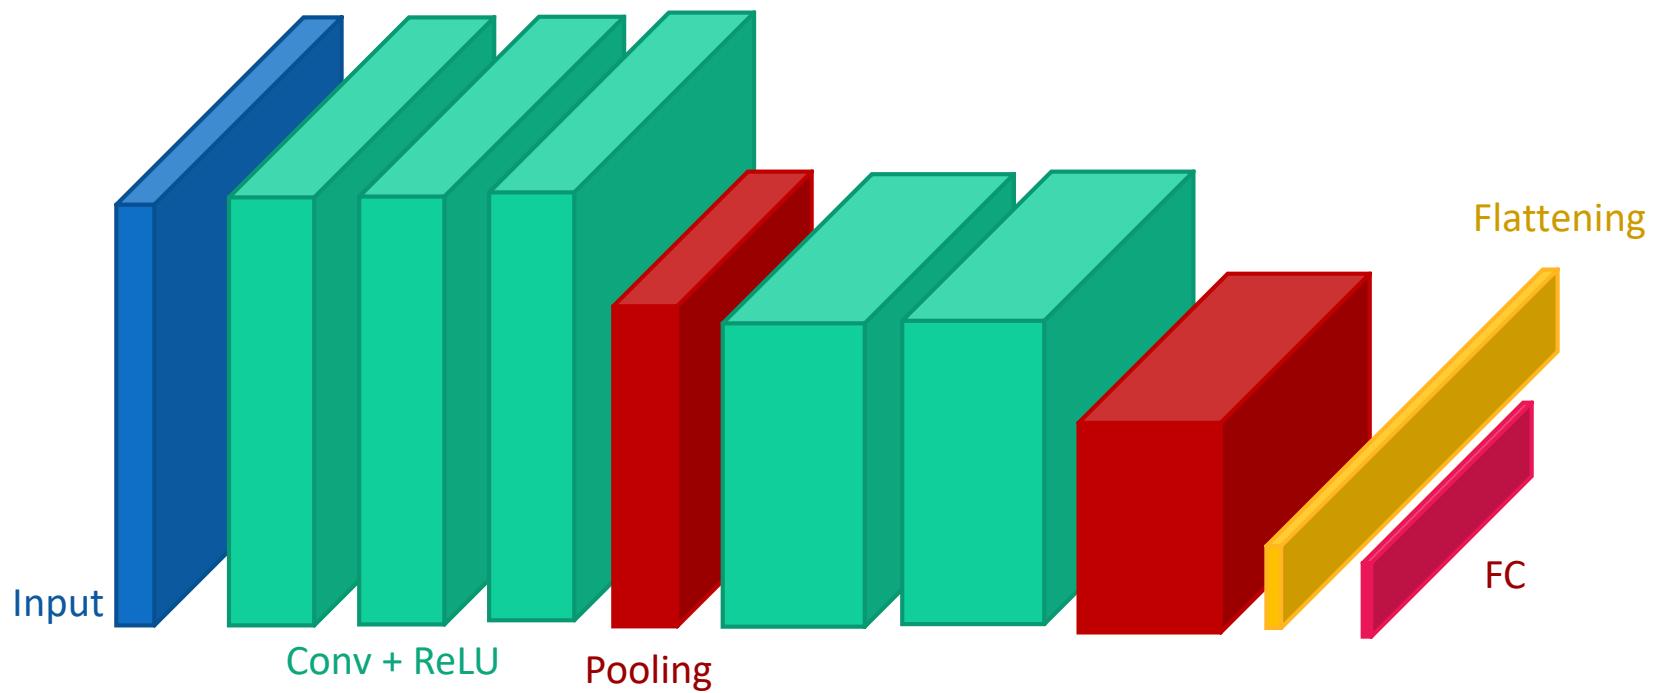
CNNs Layers

Fully Connected (FC) Layer

- ▶ **Output:** Mapped class scores
- ▶ Optimizing and computing class scores
 - ▶ **Note:** the input of this layer should be flattened
 - ▶ **Note:** the output can be normalized using SoftMax algorithm
- ▶ Each neuron is connected to all nodes in the previous step
- ▶ Generally found at the final layers of CNNs
 - ▶ At some points, they can be used along with ReLU layers
 - ▶ The last FC layer is the output layer

CNNs Layers

Let's put them all together!



CNNs Layers

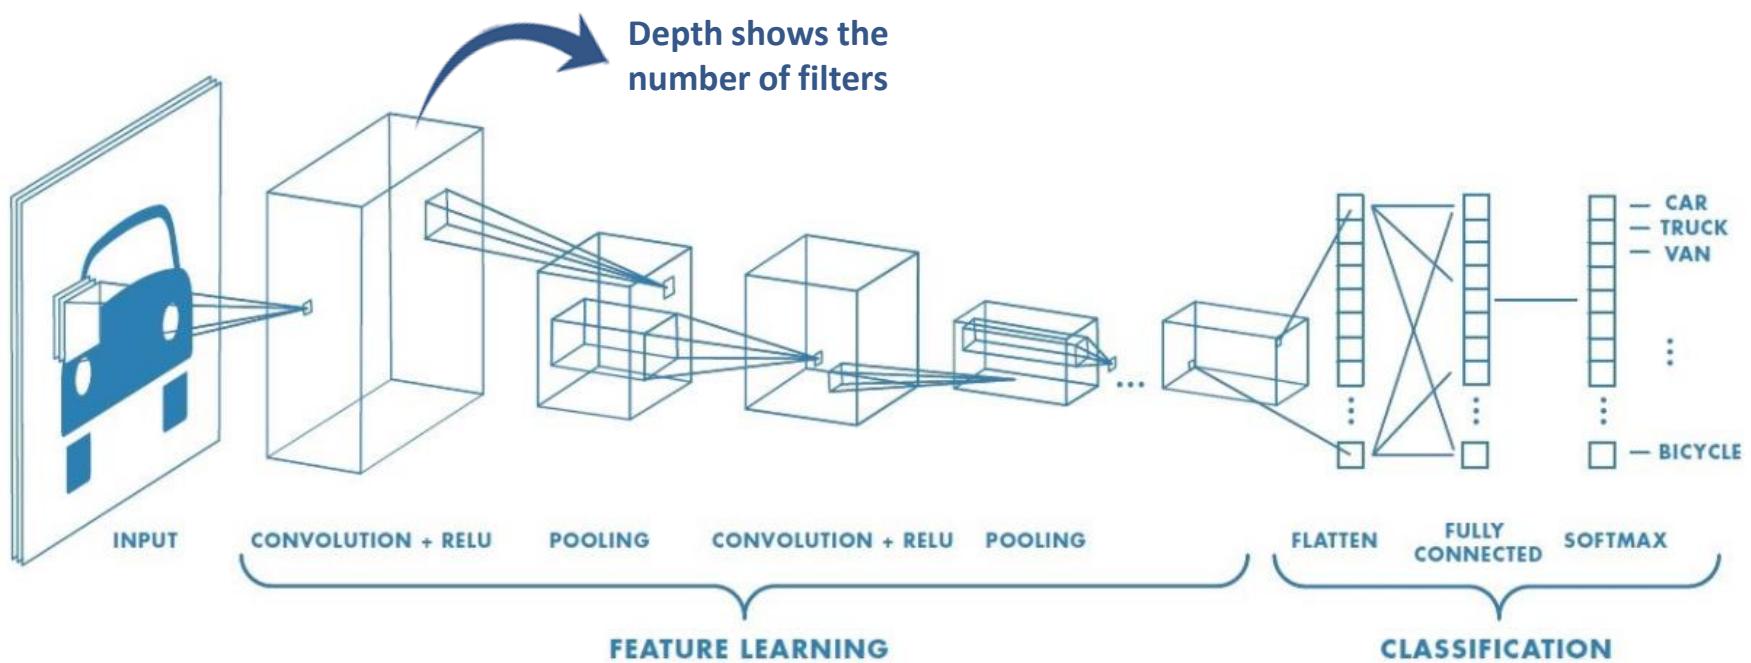
*Inp → ((Conv → ReLU) * n → Pool?) * m → (FC → ReLU) * k → FC*



The magic
pattern!

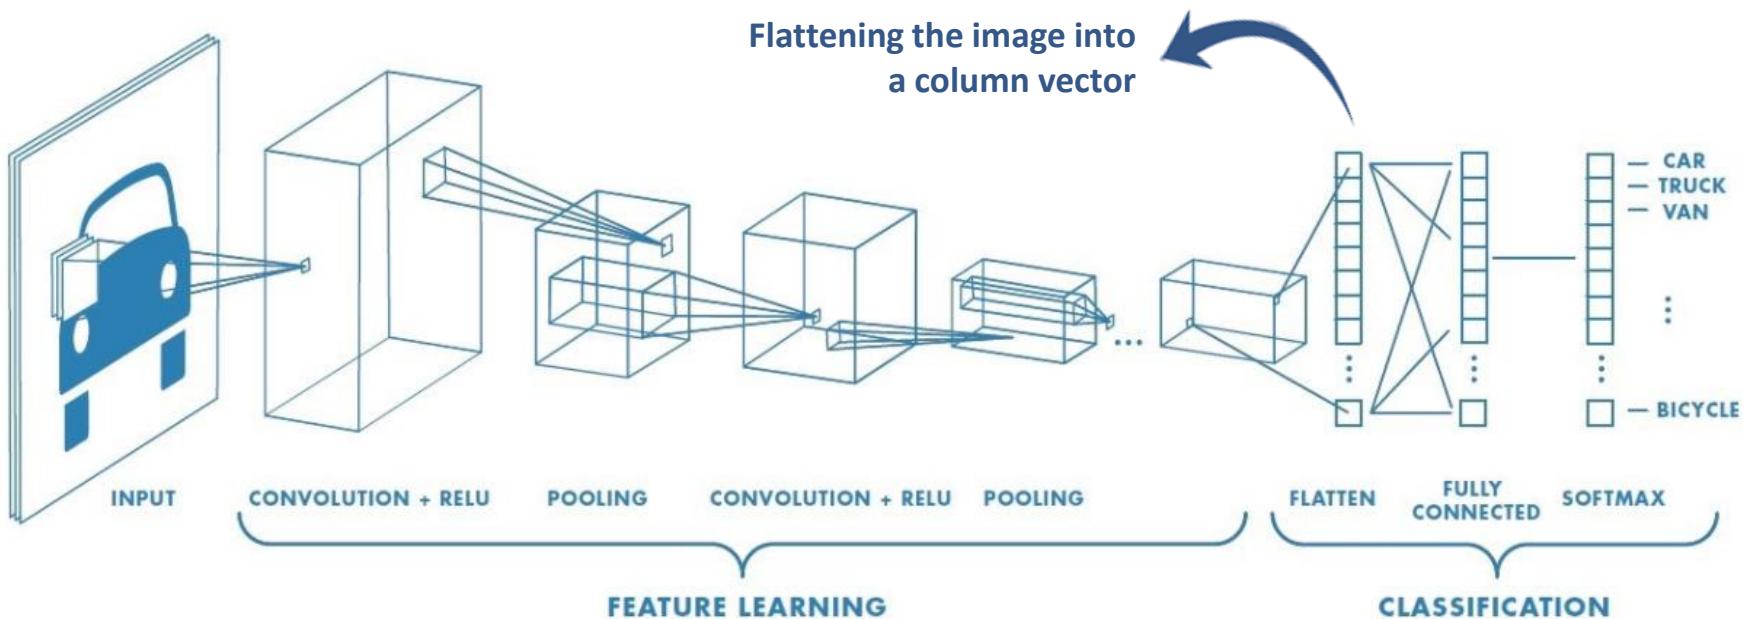
CNNs Layers

A CNN for image classification



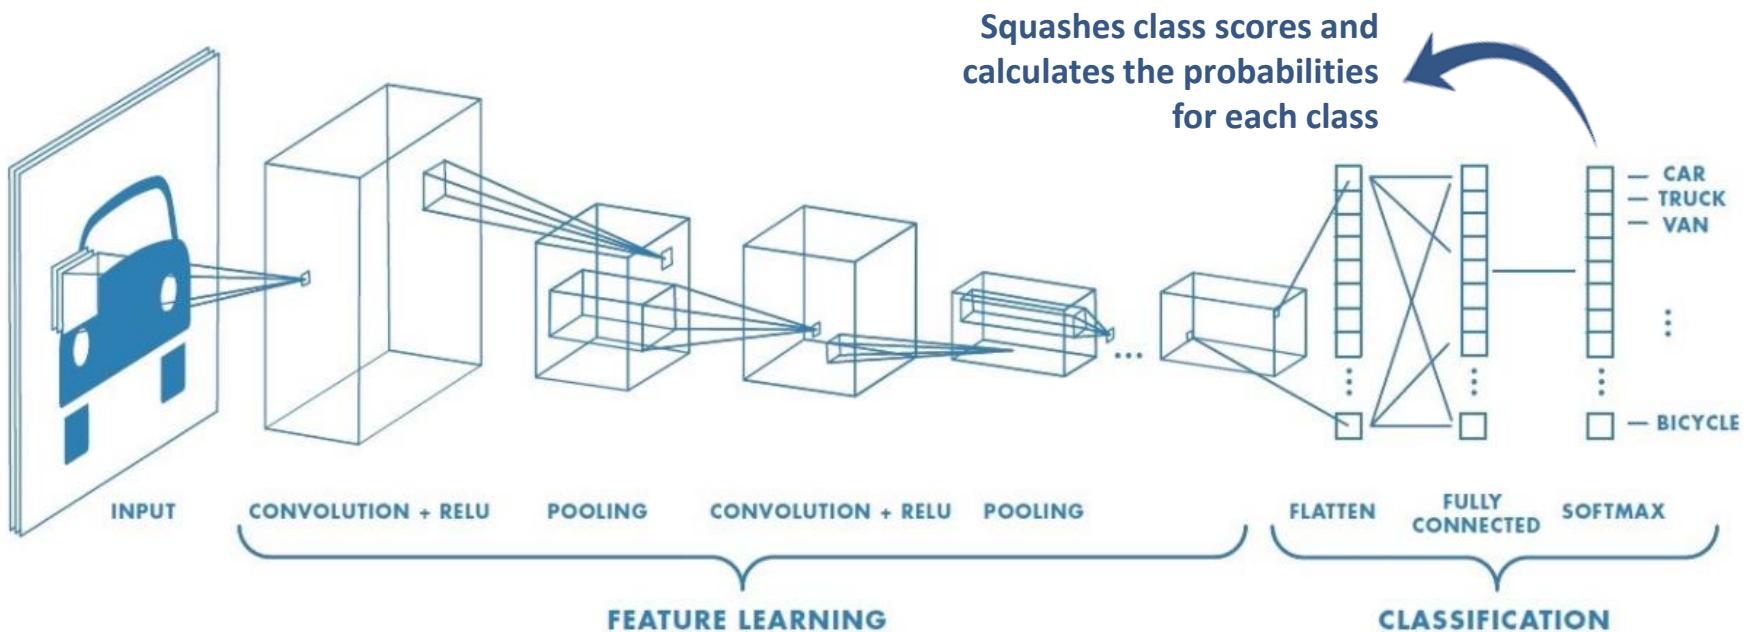
CNNs Layers

A CNN for image classification



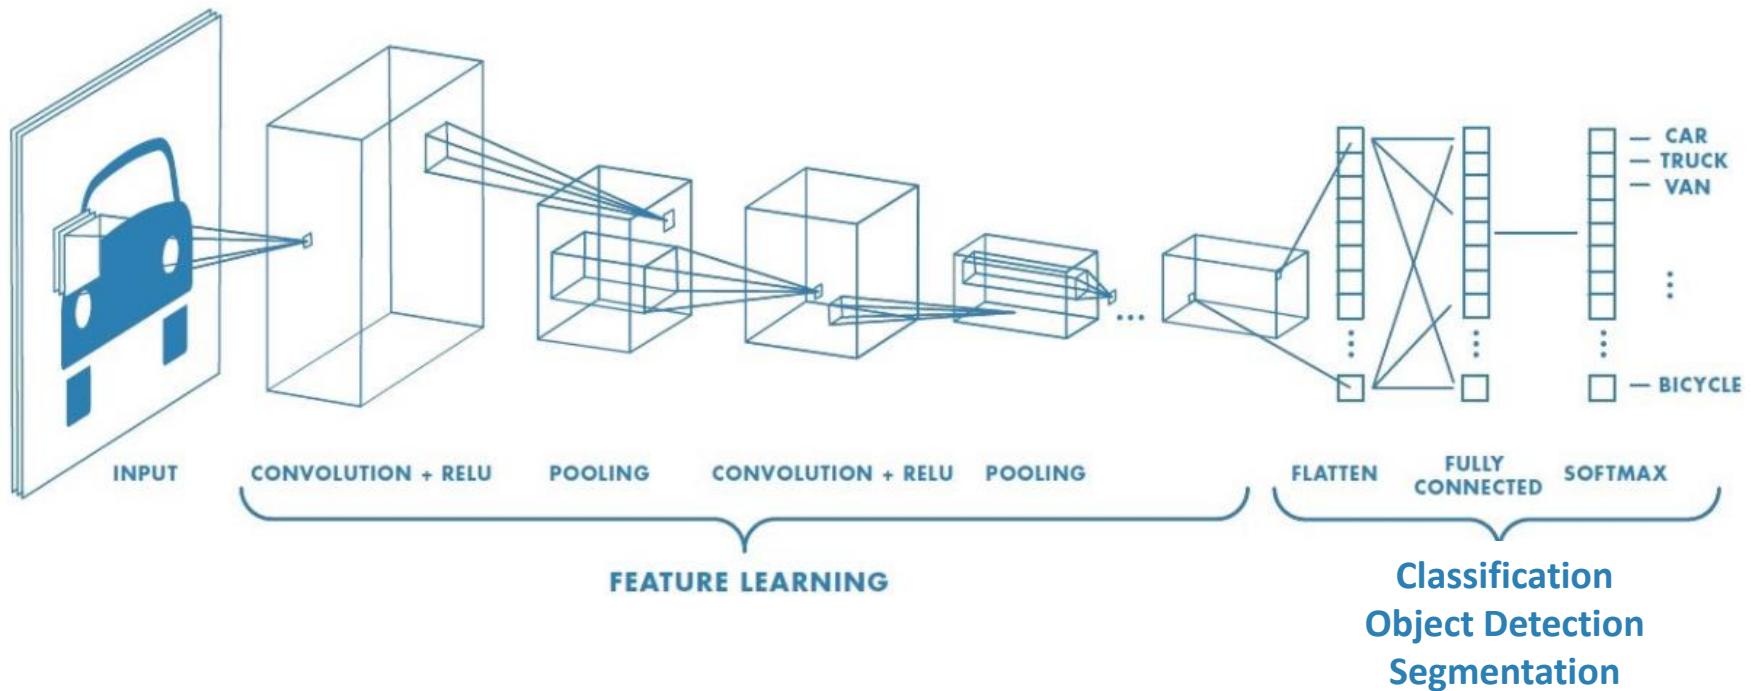
CNNs Layers

A CNN for image classification



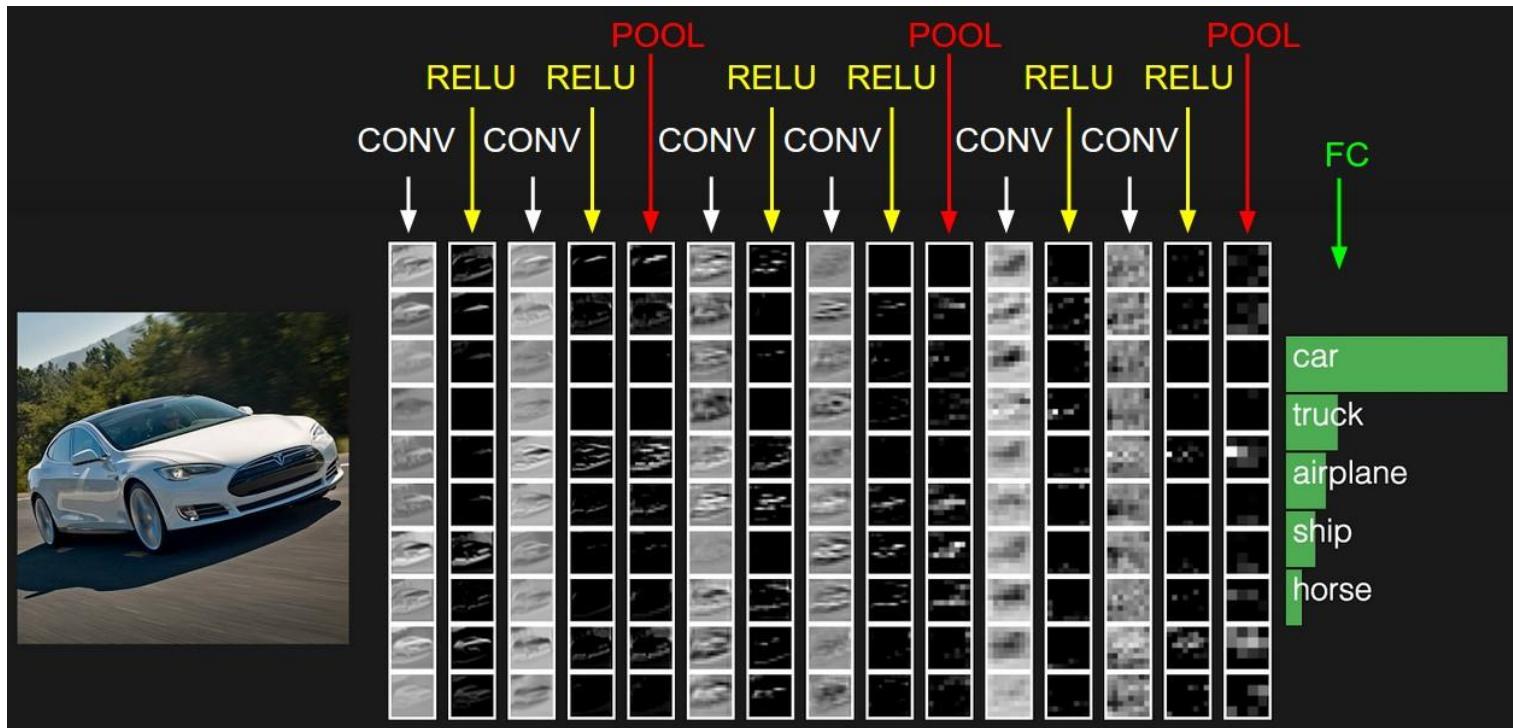
CNNs Layers

A CNN for image classification



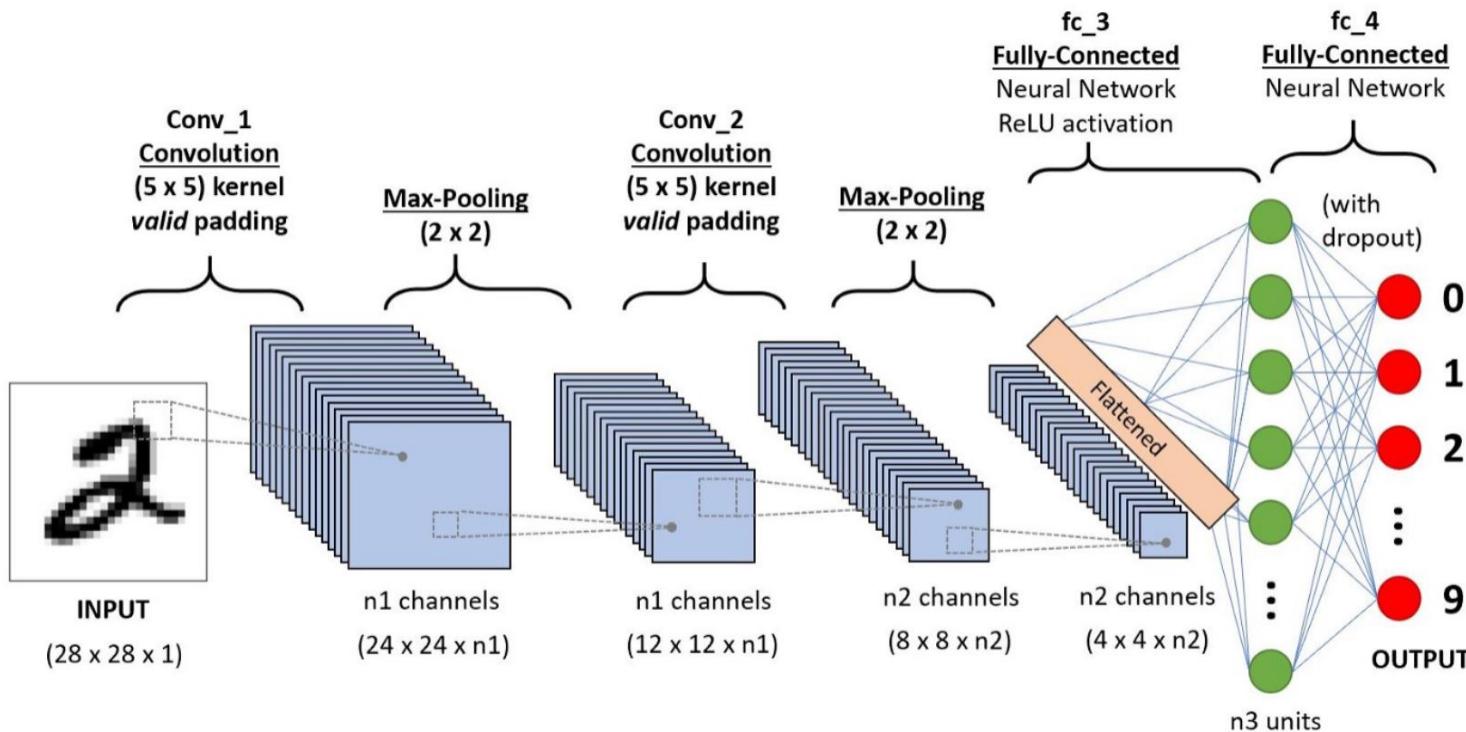
CNNs Layers

A CNN for image classification – Vehicle Classification



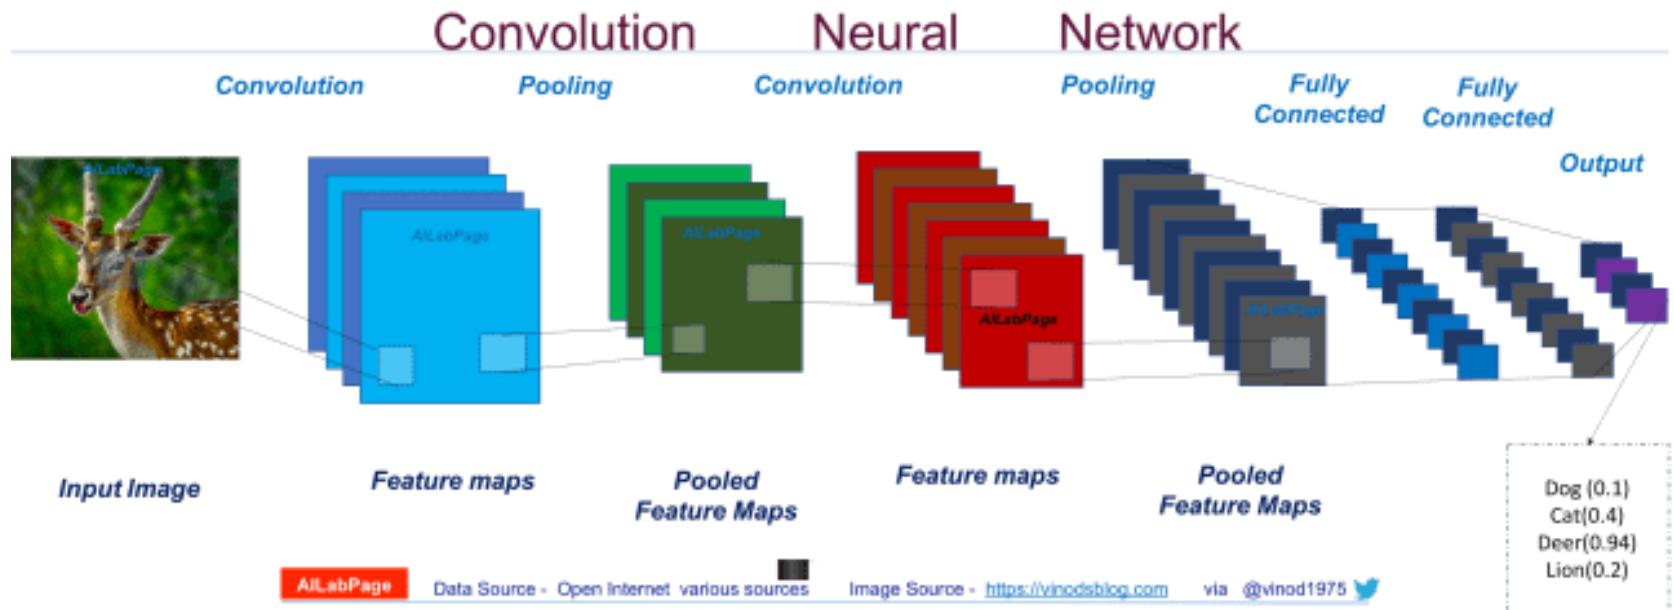
CNNs Layers

A CNN for image classification – **Handwritten digits classification**



CNNs Layers

A CNN for image classification – [Image classification](#)



References

- ▶ <http://introtodeeplearning.com/>
- ▶ <https://towardsdatascience.com/image-feature-extraction-traditional-and-deep-learning-techniques-ccc059195d04>
- ▶ <https://www.encyclopedia.com/computing/news-wires-white-papers-and-books/digital-images>
- ▶ <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>
- ▶ <https://machinelearningmastery.com/classification-versus-regression-in-machine-learning>
- ▶ <https://mathworld.wolfram.com/Convolution.html>
- ▶ <https://aishack.in/tutorials/image-convolution-examples/>
- ▶ <https://cs231n.github.io/convolutional-networks/>

Questions?

