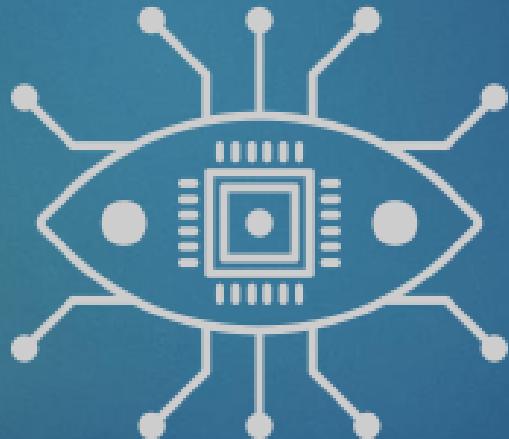


# Deep Learning from Scratch

## Session #7: Convolutional Neural Networks



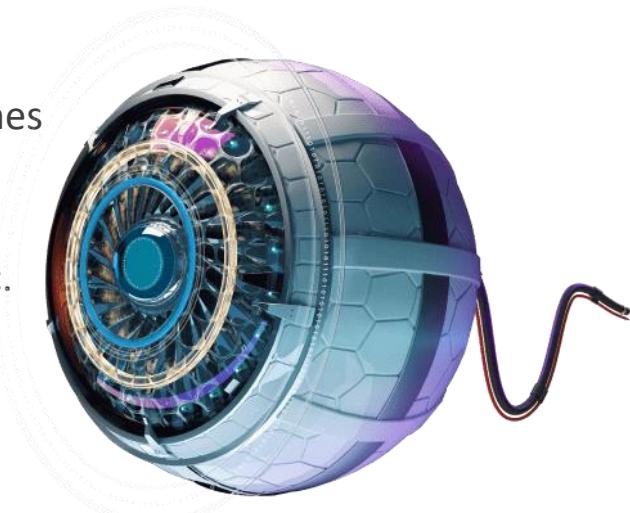
by: Ali Tourani – Summer 2021

# Agenda

- ▶ Computer Vision
- ▶ Feature Extraction
- ▶ Classification vs. Regression
- ▶ Convolution
- ▶ Convolutional Neural Networks
- ▶ CNNs Layers

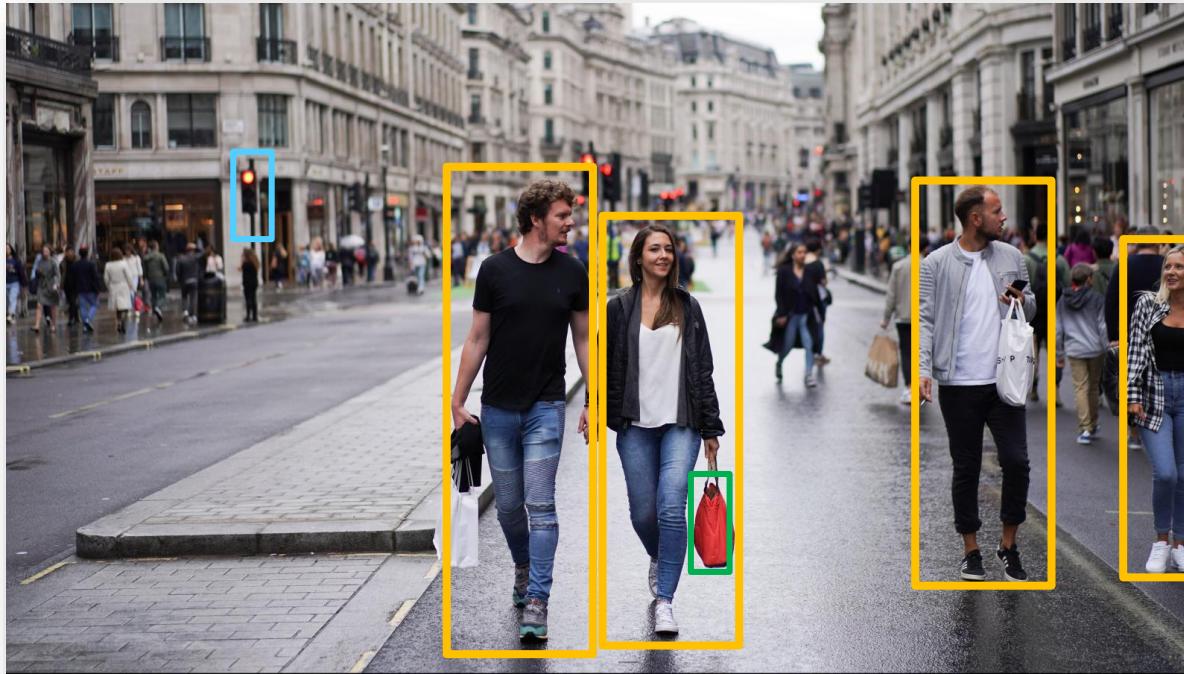
# Computer Vision

- ▶ A very popular field of AI and an **interdisciplinary scientific field**
- ▶ Enabling computers to understand from digital images or videos
  - ▶ Taking action based on the information obtained from **visual inputs**
- ▶ Data are captured from **cameras** or generated in computers
- ▶ Modelling **image processing** using **ML techniques**
  - ▶ Applying ML to recognize patterns from images/frames
- ▶ What actions to do?
  - ▶ Distinguishing between objects, classifying them, etc.



# Computer Vision

By utilizing Computer Vision algorithms we can **identify objects**



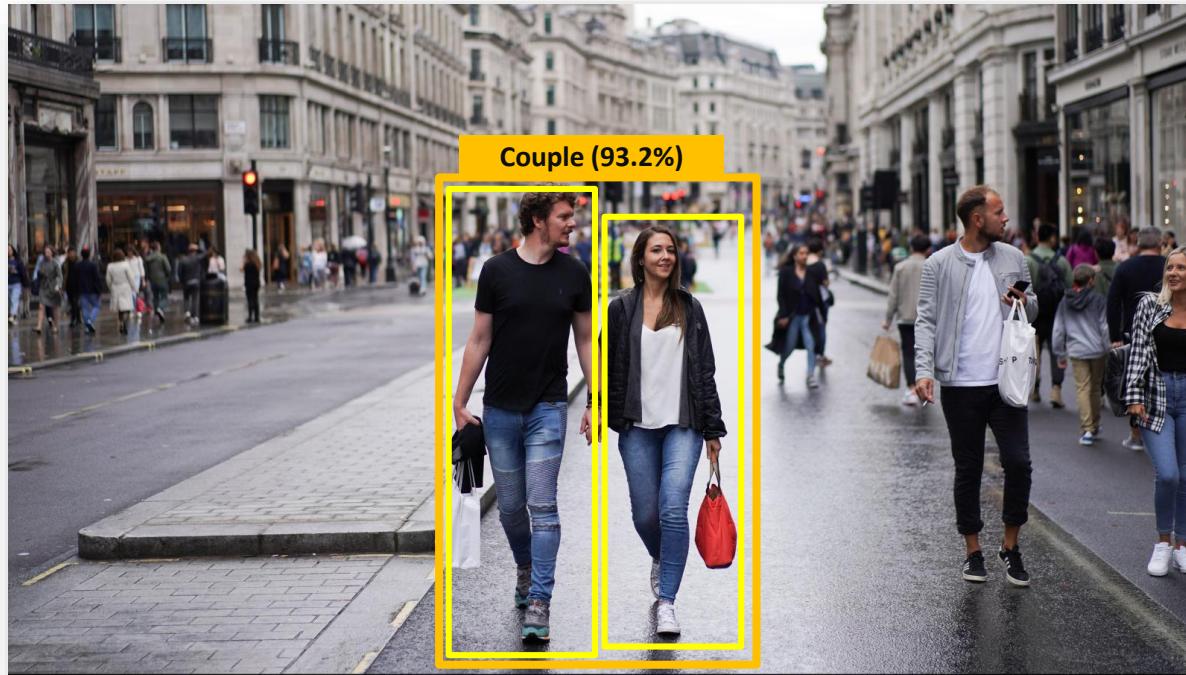
# Computer Vision

By utilizing Computer Vision algorithms we can **predict future actions/events**



# Computer Vision

By utilizing Computer Vision algorithms we can **extract info and discover patterns**



# Computer Vision



# Computer Vision

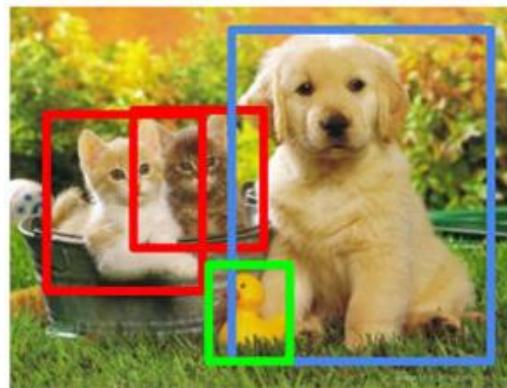
## Image Classification



# Computer Vision

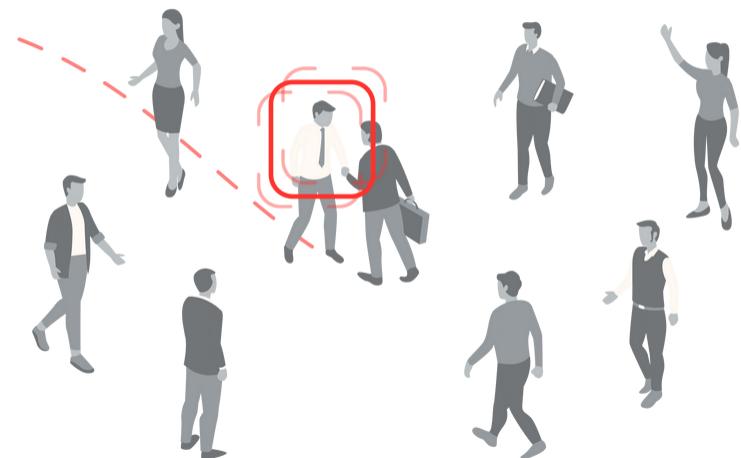
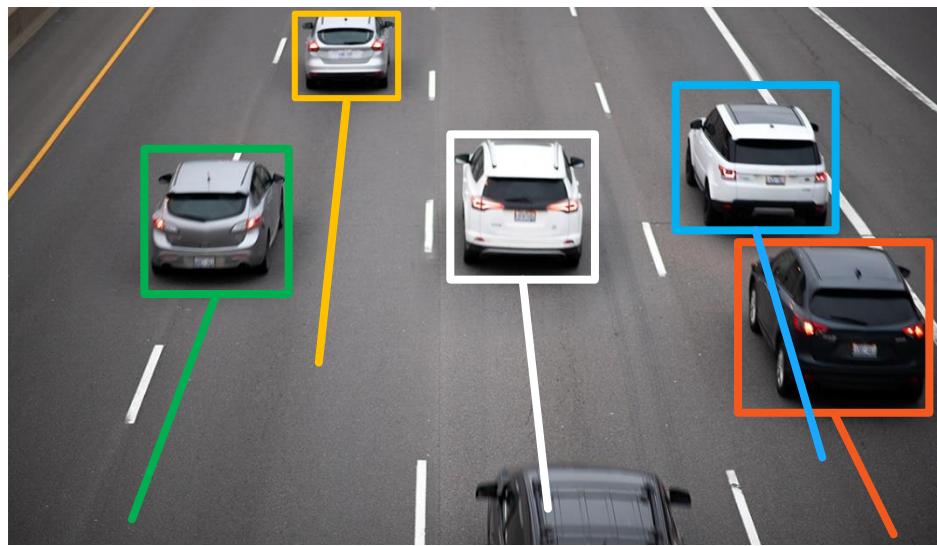
## Object Detection

Color	Class
•	Person
•	Car
•	Cat
•	Dog
•	Duck



# Computer Vision

## Object Tracking



# Computer Vision

## Image Retrieval

Query



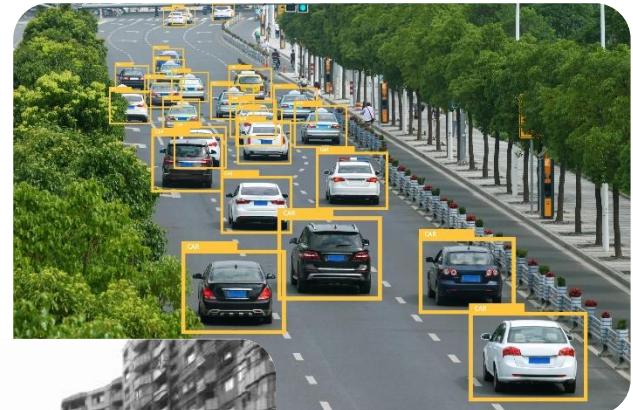
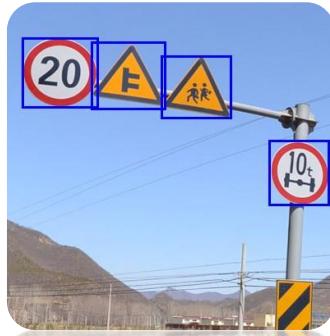
Answer



# Computer Vision

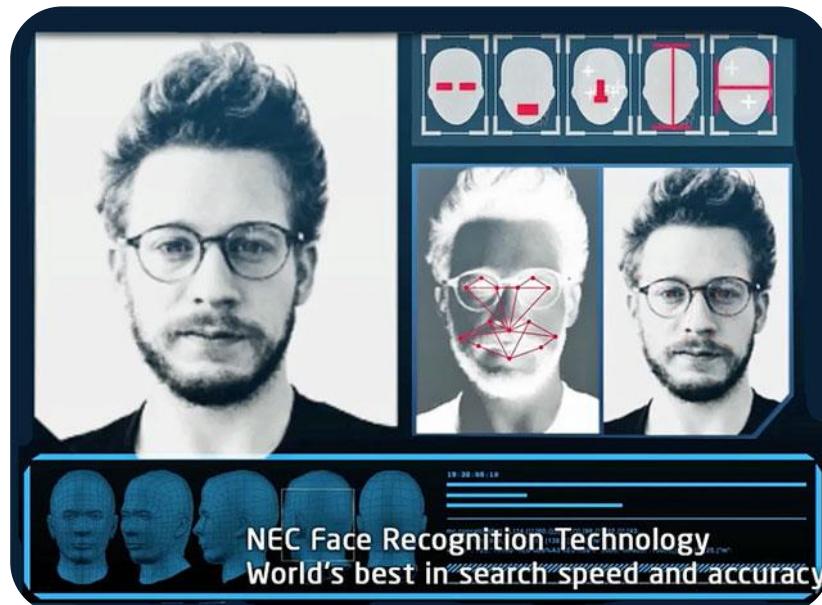
## Applications of CV – Autonomous Vehicles

- ▶ Vehicle Detection
- ▶ Pedestrian Detection
- ▶ Traffic Sign Recognition
- ▶ Lane Markers Detection



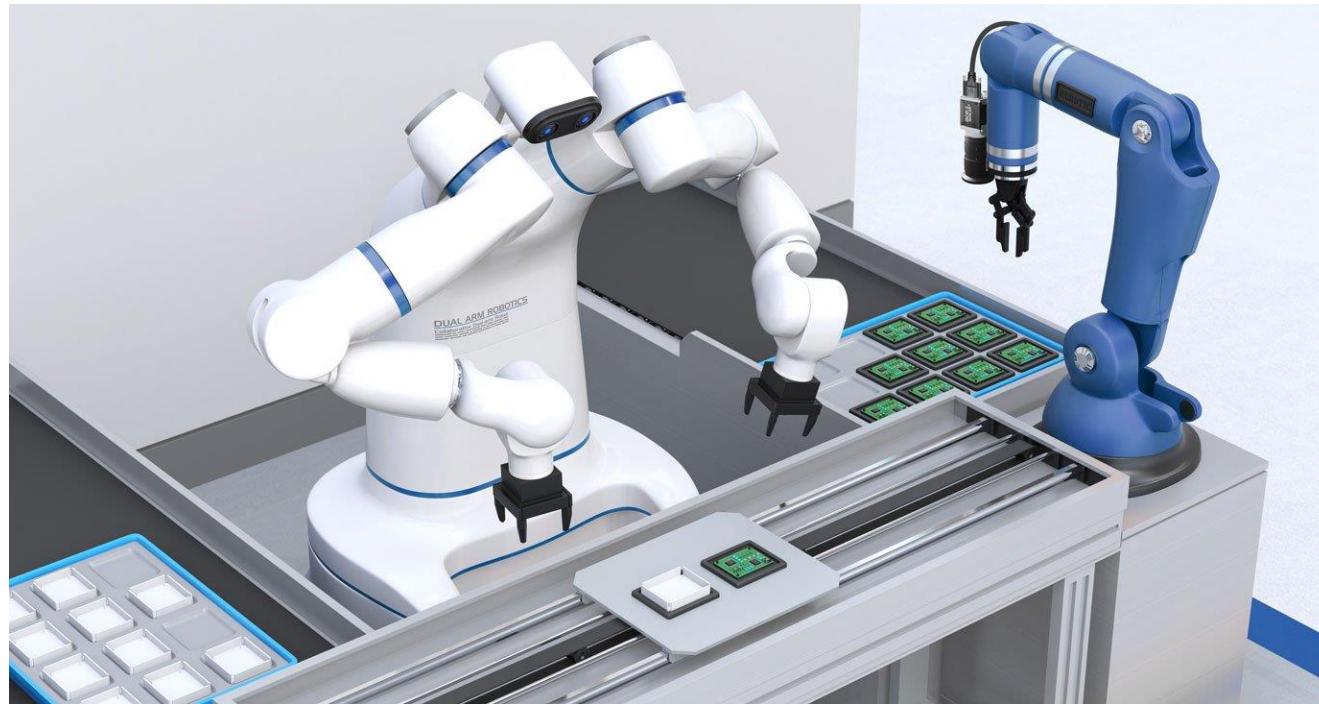
# Computer Vision

## Applications of CV – Face Recognition



# Computer Vision

## Applications of CV – Robotics



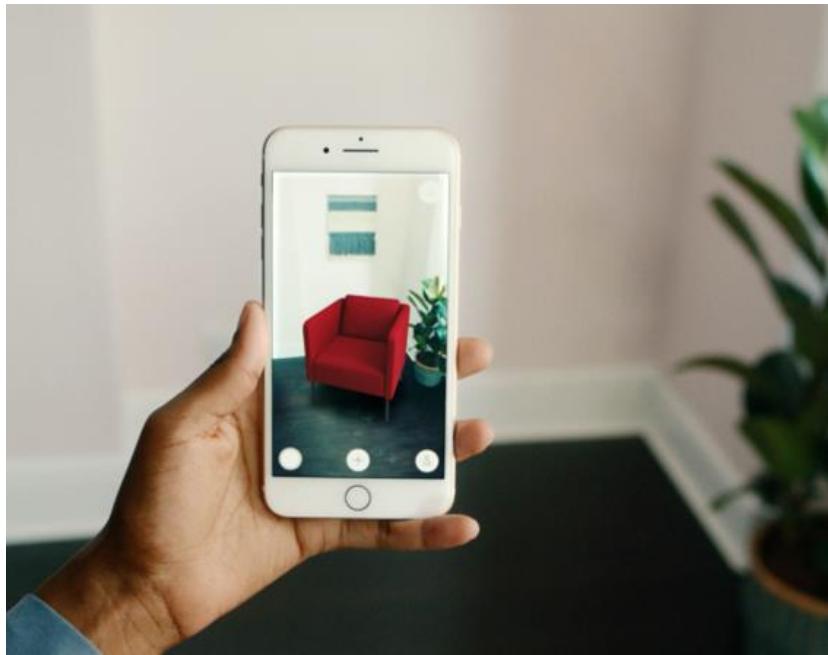
# Computer Vision

## Applications of CV – Google Translate



# Computer Vision

## Applications of CV – Mobile Computing



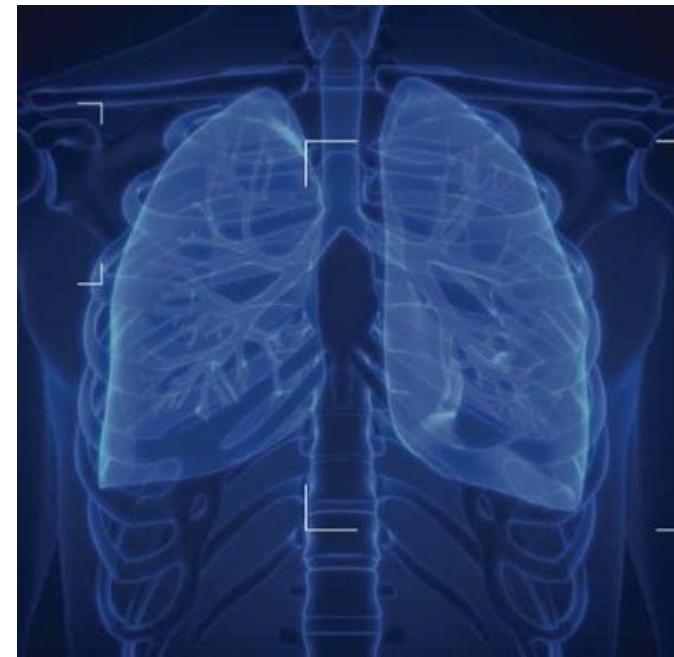
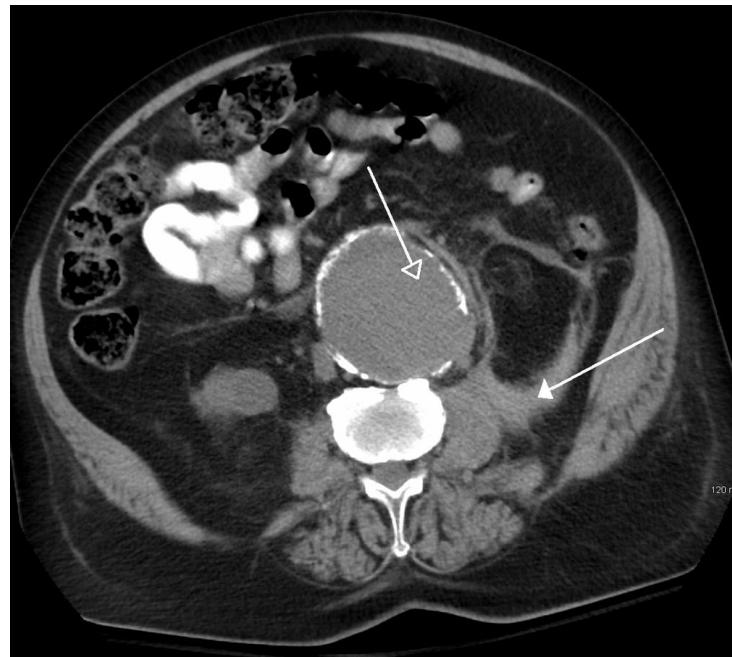
# Computer Vision

## Applications of CV – Abandoned Object Detection



# Computer Vision

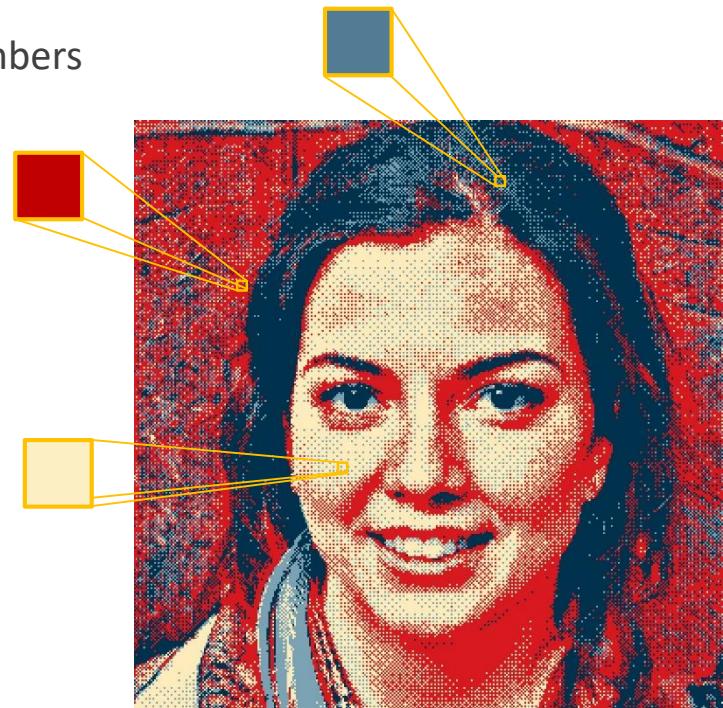
## Applications of CV – Medicine



# Feature Extraction

How do computers understand digital images?

- ▶ Digital images
  - ▶ Representations of real image as a set of numbers
  - ▶ Divided into small areas, AKA pixels
  - ▶ Pixels hold a small set of numbers for:
    - ▶ Color
    - ▶ Brightness
    - ▶ Intensity of light



# Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



Binary



Grayscale

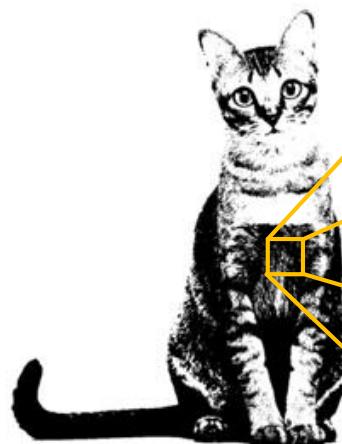


Color

# Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



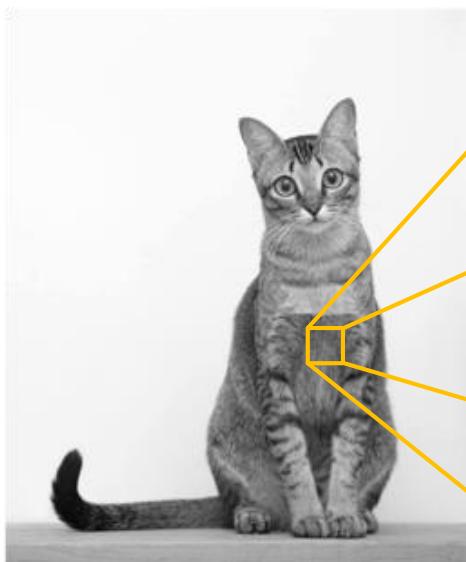
Binary image

1	1	1	1	1	0	1	1	0	0
1	1	1	0	0	0	0	0	1	1
1	1	1	1	0	0	0	0	1	1
1	1	1	0	0	0	0	0	1	1
1	1	0	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	1	1
1	1	1	1	1	0	0	0	0	0
1	1	1	0	0	0	0	0	0	0

# Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



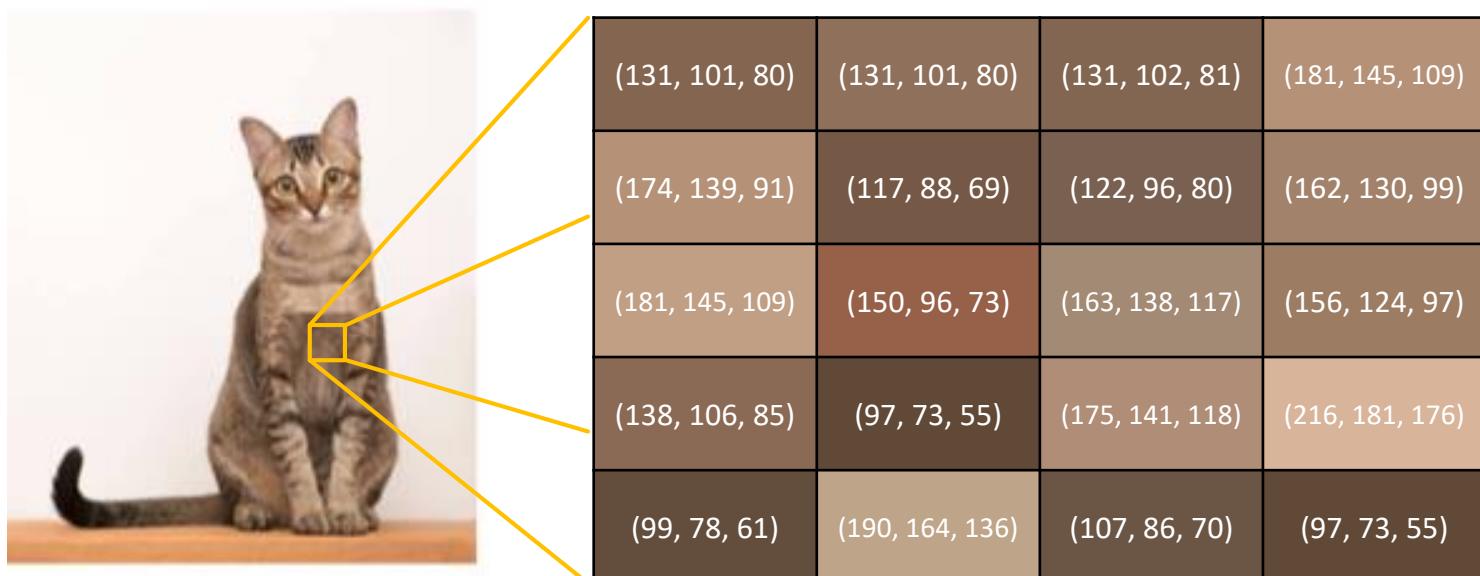
Grayscale image

121	99	78	69	129	194	93	182	81	42
78	92	188	96	95	189	204	219	94	67
105	182	79	94	79	199	219	241	201	95
185	178	98	65	132	14	18	192	182	90
98	182	201	92	94	182	194	182	71	95
84	95	81	95	102	201	29	95	185	206
96	47	173	82	96	195	181	148	103	97
98	57	142	95	84	197	61	184	73	93

# Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images

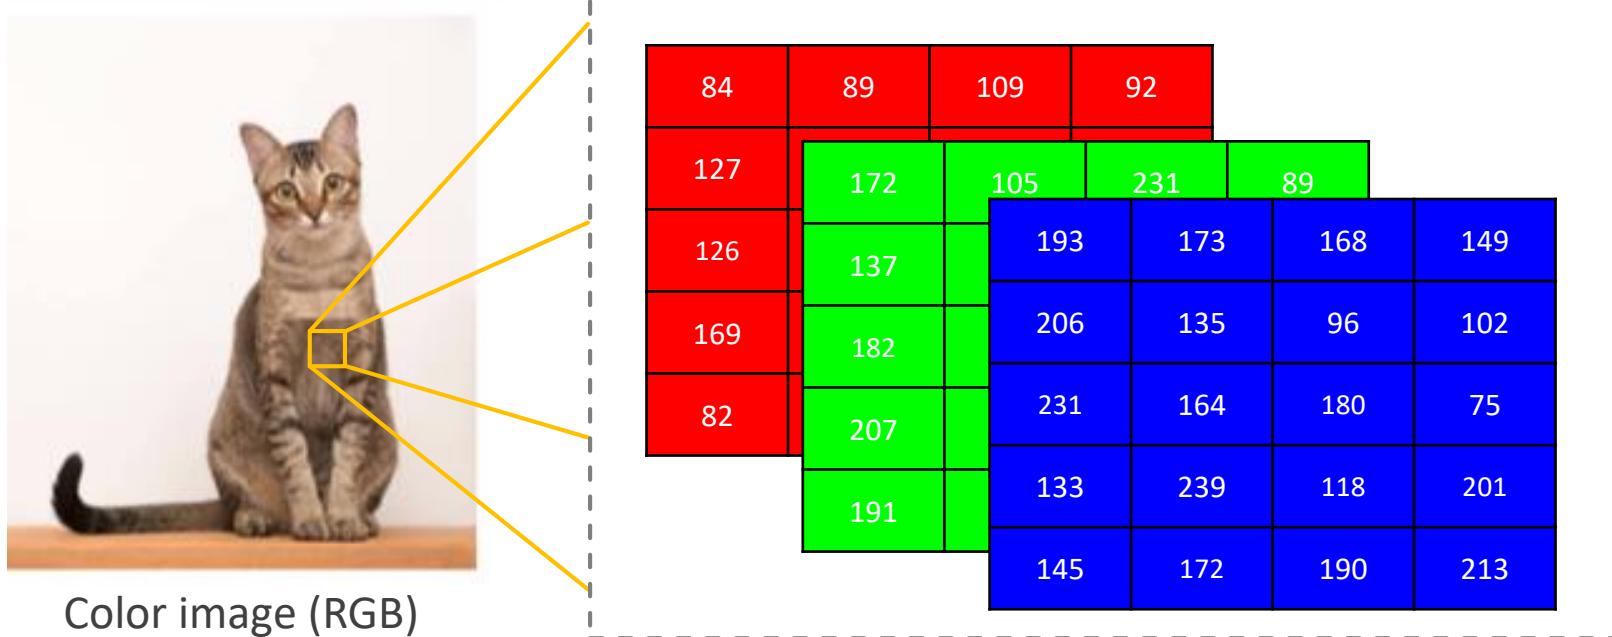


Color image (RGB)

# Feature Extraction

How do computers understand digital images?

- ▶ Types of Digital Images



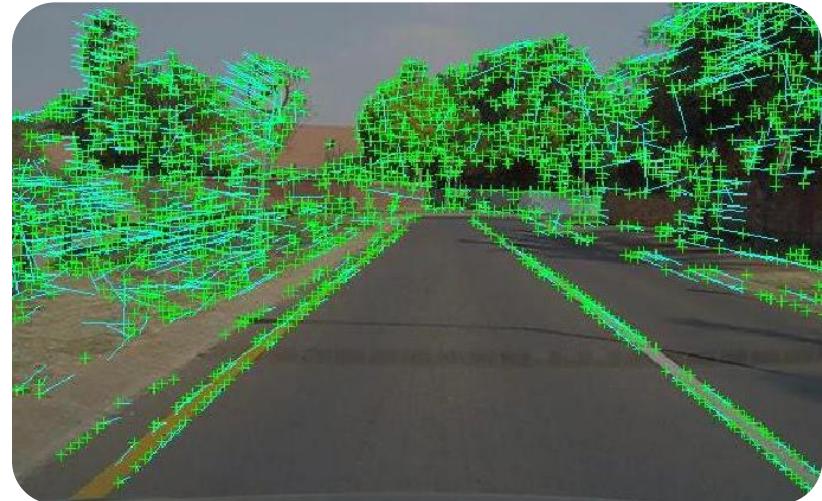
# Feature Extraction

**What are features of an image?**

- ▶ Visual parts or patterns in an image
- ▶ Useful to identify it and extract information from it

**Where to find the features?**

- ▶ Corners
- ▶ Edges
- ▶ Regions of Interest (RoI) points
- ▶ Etc.



# Feature Extraction

Where to find these features?



# Feature Extraction

Where to find these features?

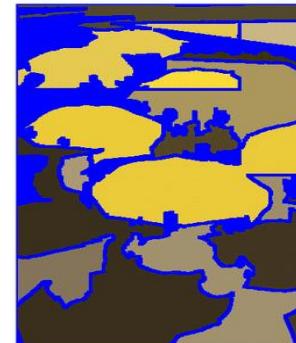


# Feature Extraction

Where to find these features?



*Regions*



# Feature Extraction

Looking for features using the pixel representation

- ▶ Low-level features
  - ▶ Capturing local appearance/texture statistics of objects
  - ▶ Highly effective in classification tasks
- ▶ High-level features
  - ▶ Semantic concepts that can be interpreted in a scene
  - ▶ The probability of observing an object in an image

## Hierarchy of Features

### Low-level features

- Algorithms: SIFT, SURF, etc.
- Interest points, edges, etc.



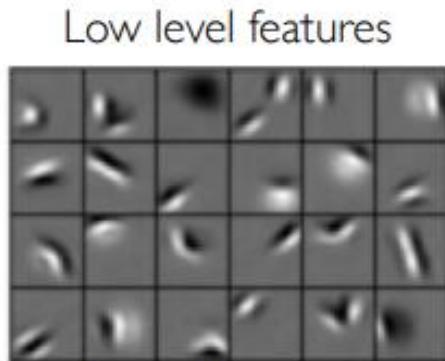
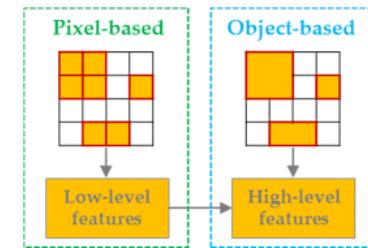
### High-level features

- Mainly used in CV
- Birds, leaves

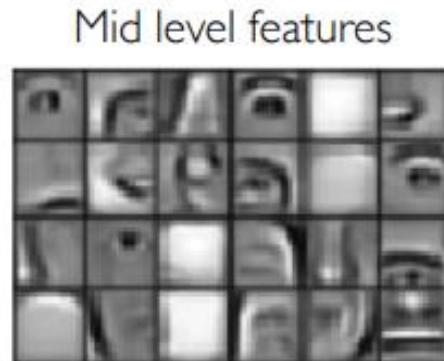
# Feature Extraction

Looking for features using the pixel representation

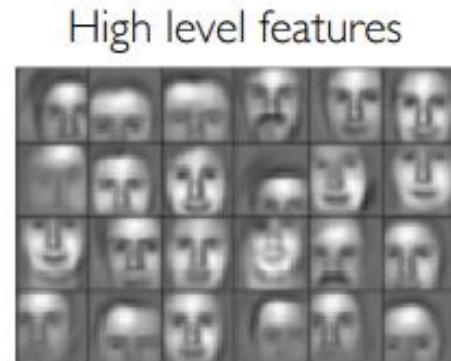
- ▶ Low-level and High-level features
  - ▶ Recall: Session#1 - Basics



Low level features  
Edges, dark spots



Mid level features  
Eyes, ears, nose



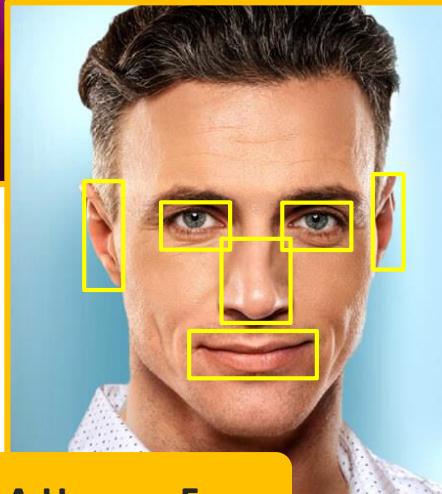
High level features  
Facial structure

# Feature Extraction

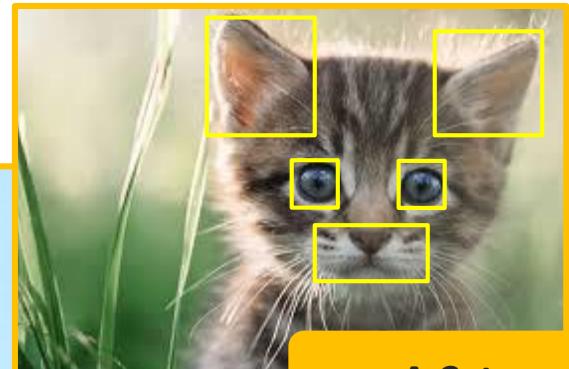
What are the high-level features that classify these images?



A Car



A Human Face



A Cat

# Feature Extraction

What are the high-level features that classify these images?



Playing football

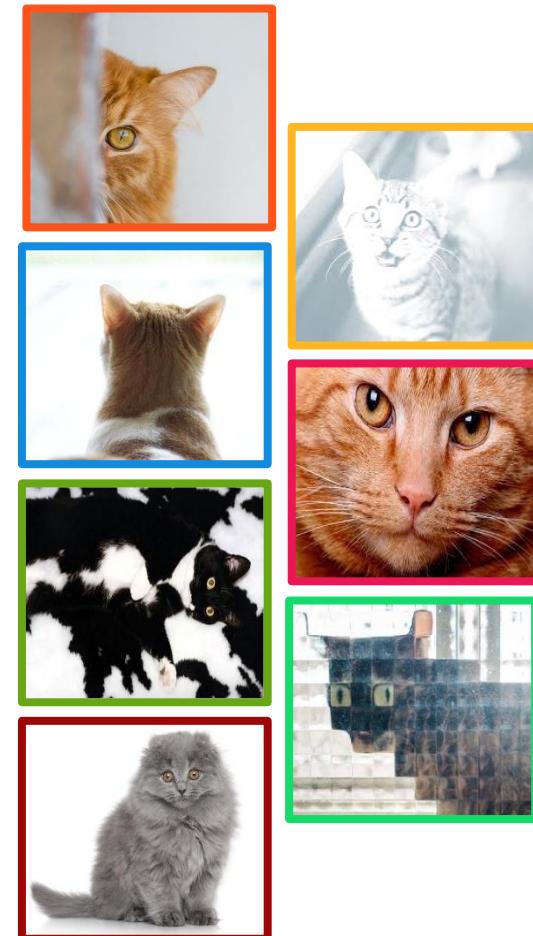


Playing basketball

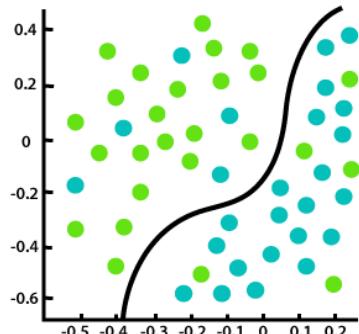
# Feature Extraction



What are the challenges?



# Classification vs. Regression



## Classification

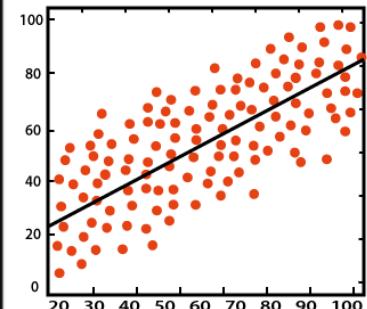
- Predicting a label
- Outputs are discrete
- Outputs = labels
- Input values can be continuous

**Use case:** forecasting the weather (classes: sunny, rainy, cloudy, snowy)

## Regression

- Predicting a quantity
- Outputs are continuous (e.g. an integer or floating point value)
- Outputs can be converted into labels to present a classification problem

**Use case:** predicting the value of a car two years later



# Classification vs. Regression

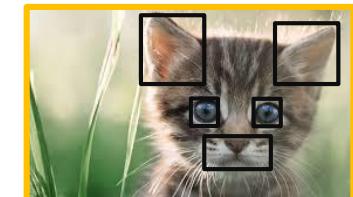
## Classification in Computer Vision



Input image

Pixel representation							
121	99	78	69	129	194	93	
78	92	188	96	95	189	204	
105	182	79	94	79	199	219	
185	178	98	65	132	14	18	
98	182	201	92	94	182	194	
84	95	81	95	102	201	29	
96	47	173	82	96	195	181	
98	57	142	95	84	197	61	

....



Unique features of a cat

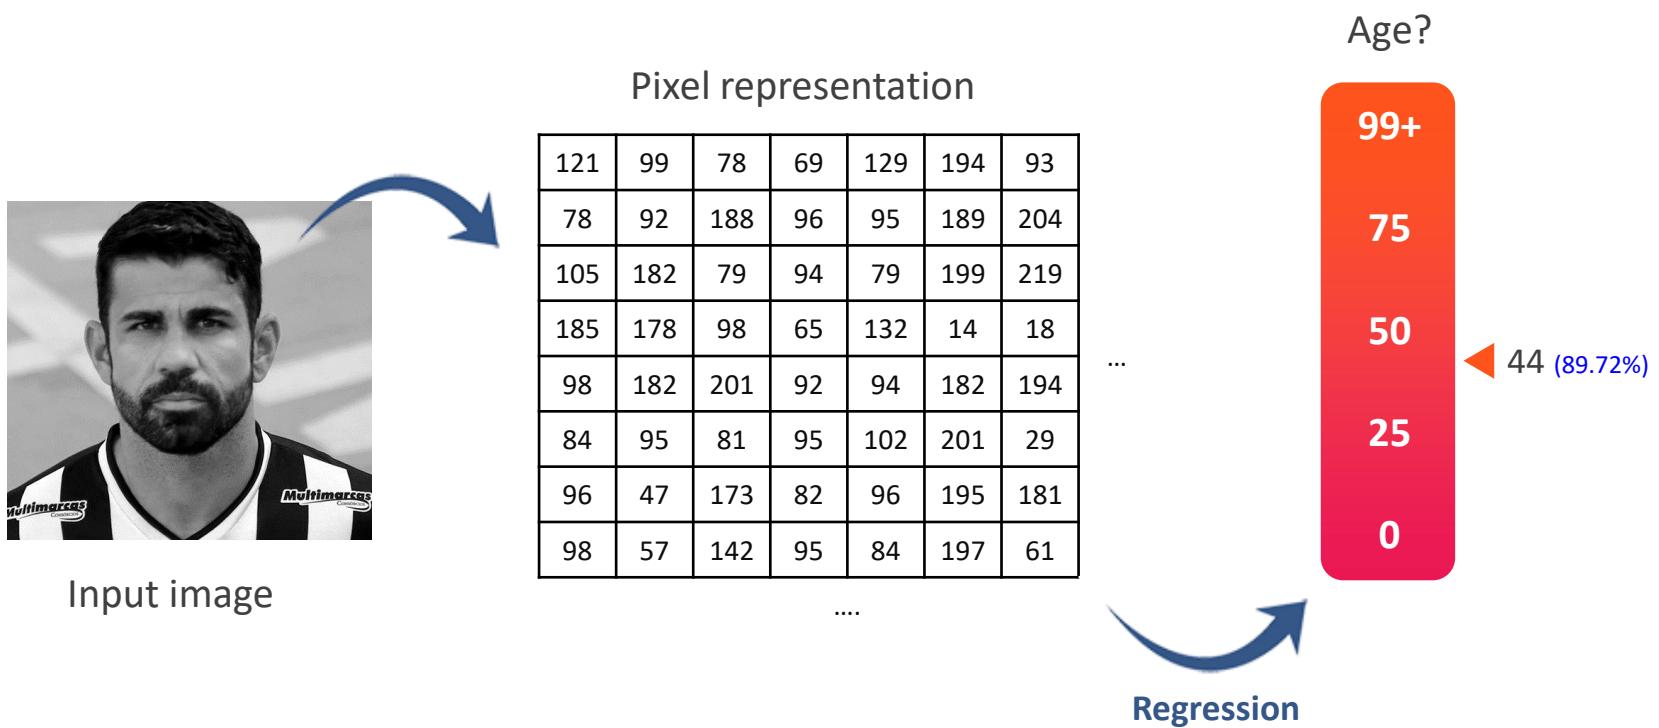
Class	Probability	
Cat	0.85	✓
Dog	0.09	✗
Mouse	0.06	✗



Classification

# Classification vs. Regression

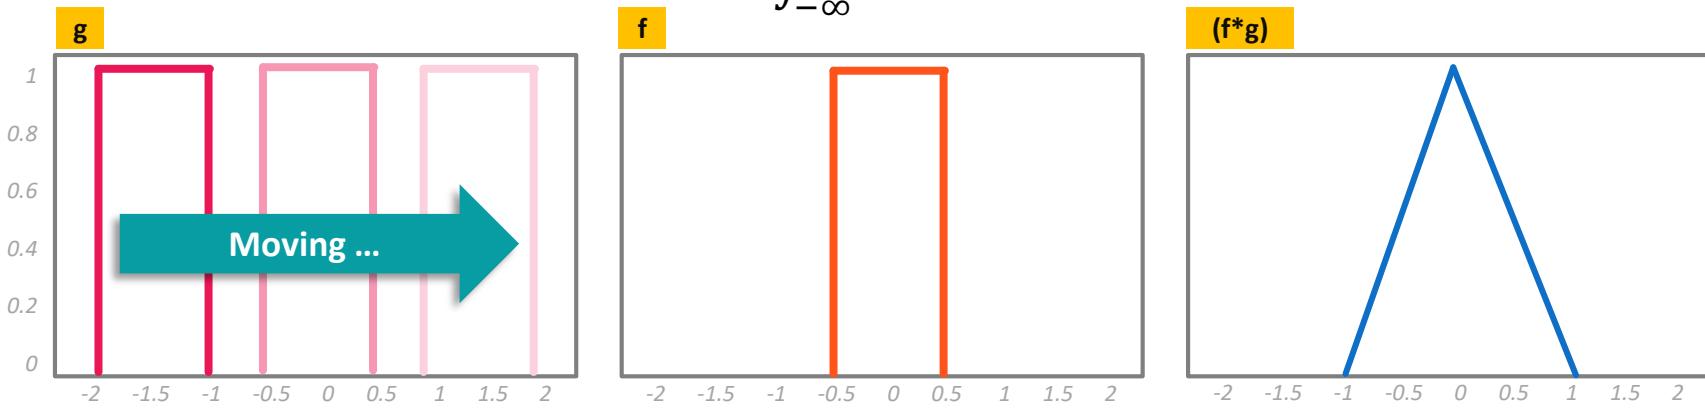
## Regression in Computer Vision



# Convolution

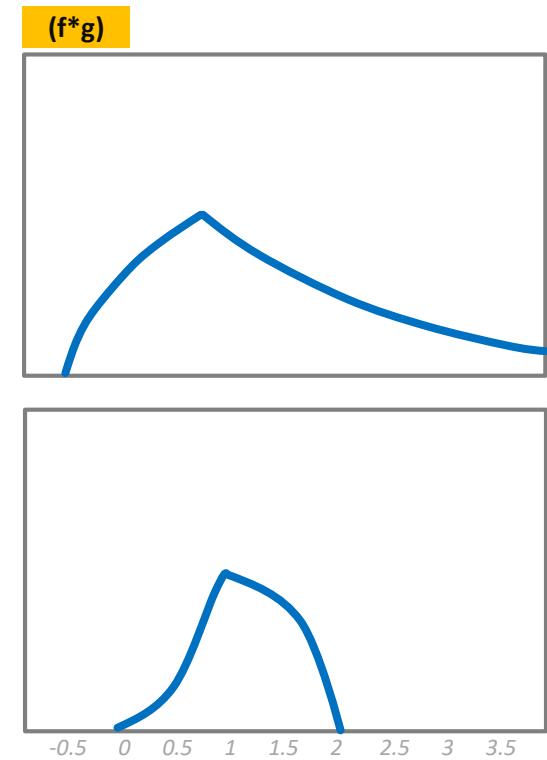
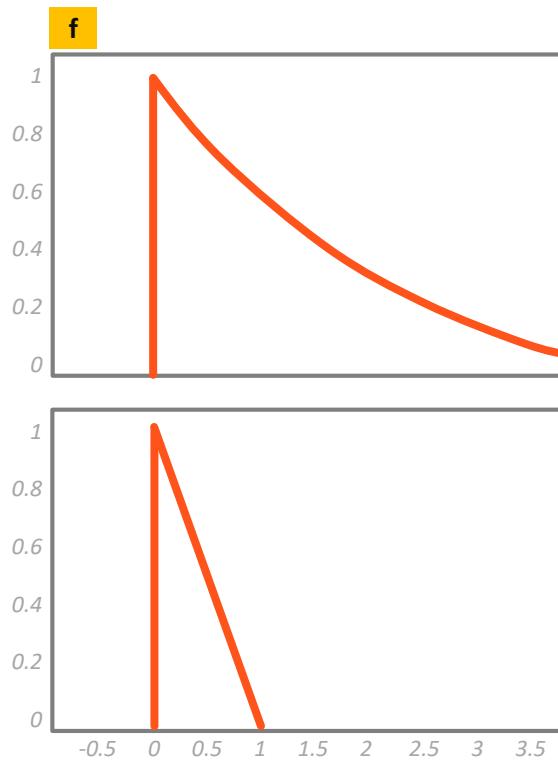
- ▶ A mathematical operation on **two functions ( $f$  and  $g$ )** to produce a **third one**
- ▶ **Result:** Identifying how  $f$  can modify the shape of  $g$
- ▶ **Definition:** an integral that expresses the amount of overlap of  $f$  as it is shifted over  $g$

$$[f * g](t) \equiv \int_{-\infty}^{+\infty} f(\tau)g(t - \tau). d\tau$$



# Convolution

$$[f * g](t) \equiv \int_{-\infty}^{+\infty} f(\tau)g(t - \tau). d\tau$$



# Convolution

What about applying convolution on images?

I(0,0)	I(1,0)	I(2,0)	I(3,0)	I(4,0)	I(5,0)	I(6,0)
I(0,1)	I(1,1)	I(2,1)	I(3,1)	I(4,1)	I(5,1)	I(6,1)
I(0,2)	I(1,2)	I(2,2)	I(3,2)	I(4,2)	I(5,2)	I(6,2)
I(0,3)	I(1,3)	I(2,3)	I(3,3)	I(4,3)	I(5,3)	I(6,3)
I(0,4)	I(1,4)	I(2,4)	I(3,4)	I(4,4)	I(5,4)	I(6,4)
I(0,5)	I(1,5)	I(2,5)	I(3,5)	I(4,5)	I(5,5)	I(6,5)
I(0,6)	I(1,6)	I(2,6)	I(3,6)	I(4,6)	I(5,6)	I(6,6)

Input image



H(0,0)	H(1,0)	H(2,0)
H(0,1)	H(1,1)	H(2,1)
H(0,2)	H(1,2)	H(2,2)

Filter/Kernel/Mask

O(0,0)				

Output image  
(feature map)

# Convolution

What about applying convolution on images?

1	0	1	0	1
1	1	1	0	0
1	0	1	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} ? \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image  
(feature map)

# Convolution

What about applying convolution on images?

1 <small><math>\times 1</math></small>	0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	0	1
1 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1 <small><math>\times 1</math></small>	0	0
1 <small><math>\times 1</math></small>	0 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 \\ \vdots \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

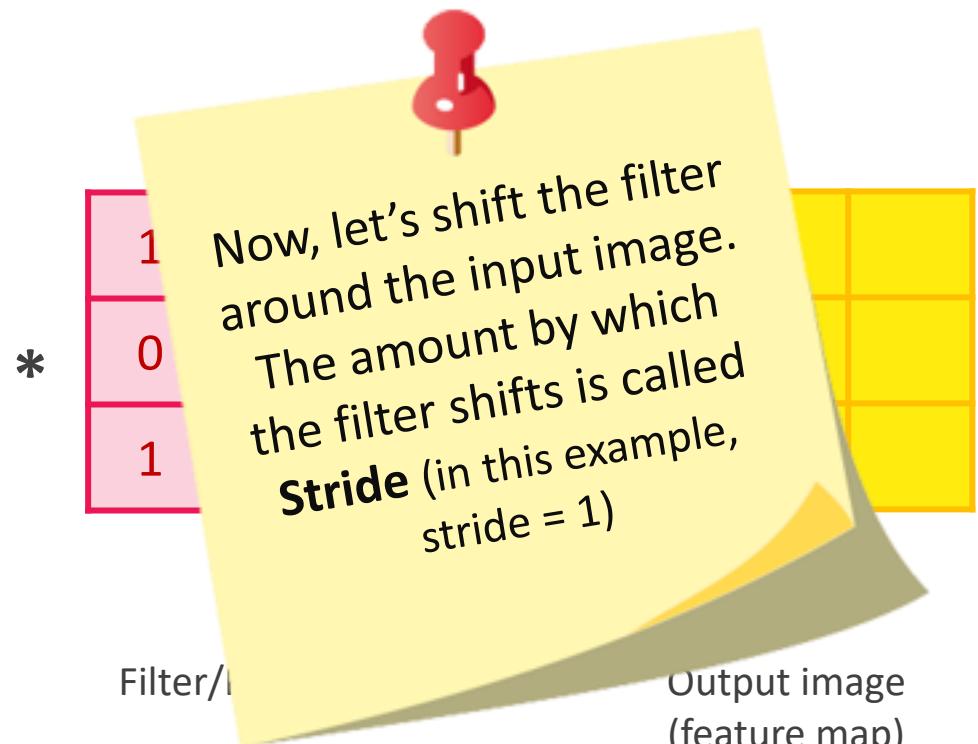
Output image  
(feature map)

# Convolution

What about applying convolution on images?

1 $\times 1$	0 $\times 0$	1 $\times 1$	0	1
1 $\times 0$	1 $\times 1$	1 $\times 1$	0	0
1 $\times 1$	0 $\times 0$	1 $\times 1$	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)



# Convolution

What about applying convolution on images?

1	0 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	0 <small><math>\times 1</math></small>	1
1	1 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	0 <small><math>\times 1</math></small>	0
1	0 <small><math>\times 1</math></small>	1 <small><math>\times 0</math></small>	1 <small><math>\times 1</math></small>	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 & 2 \\ & \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image  
(feature map)

# Convolution

What about applying convolution on images?

1	0	1 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>
1	1	1 <sub>x0</sub>	0 <sub>x1</sub>	0 <sub>x1</sub>
1	0	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 & 2 & 3 \\ & & \\ & & \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image  
(feature map)

# Convolution

What about applying convolution on images?

1	0	1	0	1
1	1	1	0	0
1	0	1	1	0
0	0	1	0	1
0	1	1	1	0

Input image (binary)

$$\begin{matrix} & \begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{matrix} & = & \begin{matrix} 6 & 2 & 3 \\ 4 & 3 & 4 \\ 4 & 4 & 3 \end{matrix} \end{matrix}$$

Filter/Kernel/Mask

Output image  
(feature map)

# Convolution

What about applying convolution on images?

1 × 1	1 × 0	1 × 1	0	0
0 × 0	1 × 1	1 × 0	1	0
0 × 1	0 × 0	1 × 1	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved feature

1	1 × 1	1 × 0	0 × 1	0
0	1 × 0	1 × 1	1 × 0	0
0	0 × 1	1 × 0	1 × 1	1
0	0	1	1	0
0	1	1	0	0

Image

4	3	

Convolved feature

1	1	1 × 1	0 × 0	0 × 1
0	1	1 × 0	1 × 1	0 × 0
0	0	1 × 1	1 × 0	1 × 1
0	0	1	1	0
0	1	1	0	0

Image

4	3	4

Convolved feature

1	1	1	0	0
0 × 1	1 × 0	1 × 1	1	0
0 × 0	0 × 1	1 × 0	1	1
0 × 1	0 × 0	1 × 1	1	0
0	1	1	0	0

Image

4	3	4
2		

Convolved feature

1	1	1	0	0
0	1 × 1	1 × 0	1 × 1	0
0	0 × 0	1 × 1	1 × 0	1
0	0 × 1	1 × 0	1 × 1	0
0	1	1	0	0

Image

4	3	4
2	4	

Convolved feature

1	1	1	0	0
0	1	1 × 1	1 × 0	0 × 1
0	0	1 × 0	1 × 1	1 × 0
0	0	1 × 1	1 × 0	0 × 1
0	1	1	0	0

Image

4	3	4
2	4	3

Convolved feature

# Convolution

What about applying convolution on images?



Input image

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9



Simple box blur

-1	-1	-1
2	2	2
-1	-1	-1



Horizontal lines

# Convolution

What about applying convolution on images?



Input image

$$\begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline -1 & 8 & -1 \\ \hline -1 & -1 & -1 \\ \hline \end{array}$$



Edge detection

$$\begin{array}{|c|c|c|} \hline 0 & -1 & 0 \\ \hline -1 & 4 & -1 \\ \hline 0 & -1 & 0 \\ \hline \end{array}$$

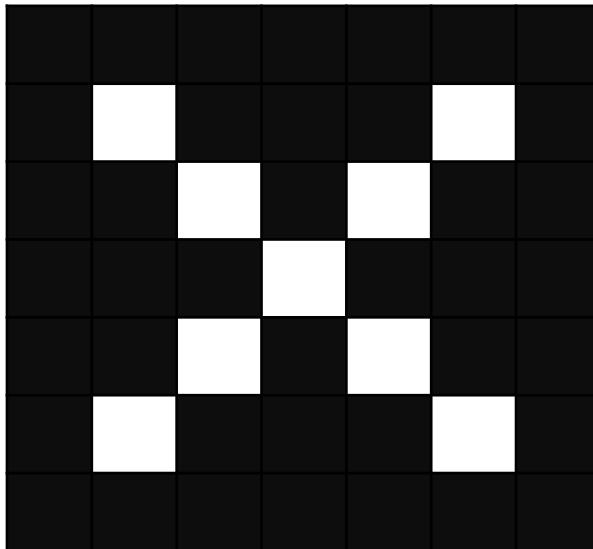


The Laplacian operator

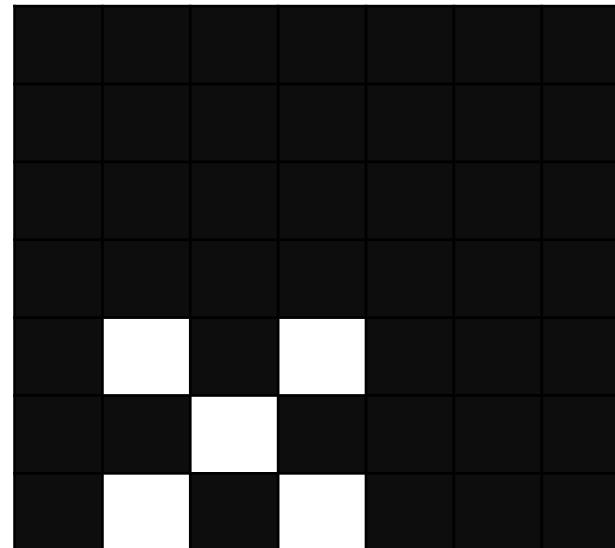
# Convolution

## Why is it important to use Convolution?

- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



The letter “X” in a binary image

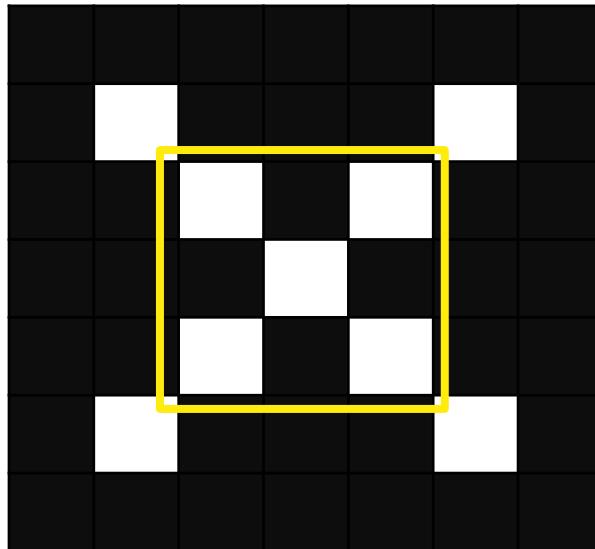


The scaled version

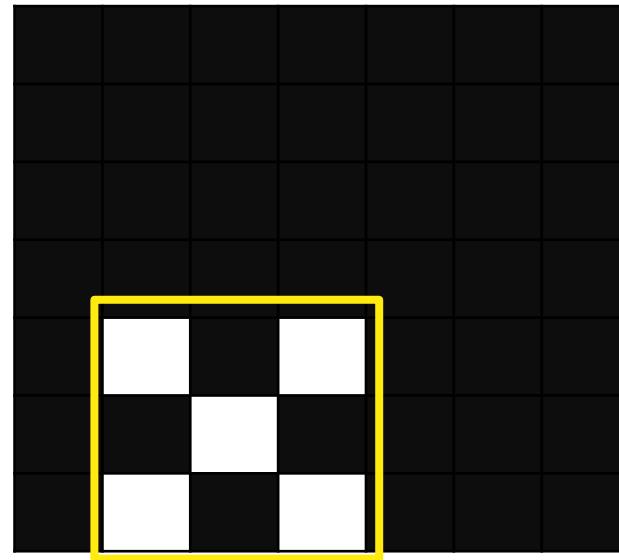
# Convolution

## Why is it important to use Convolution?

- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



The letter “X” in a binary image



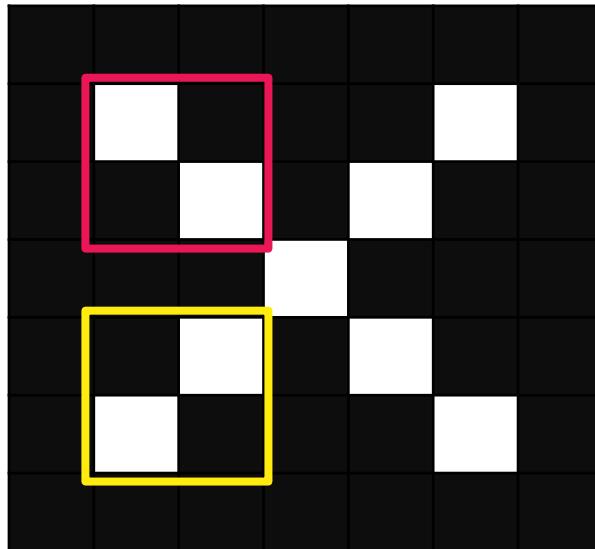
The scaled version

The same features!

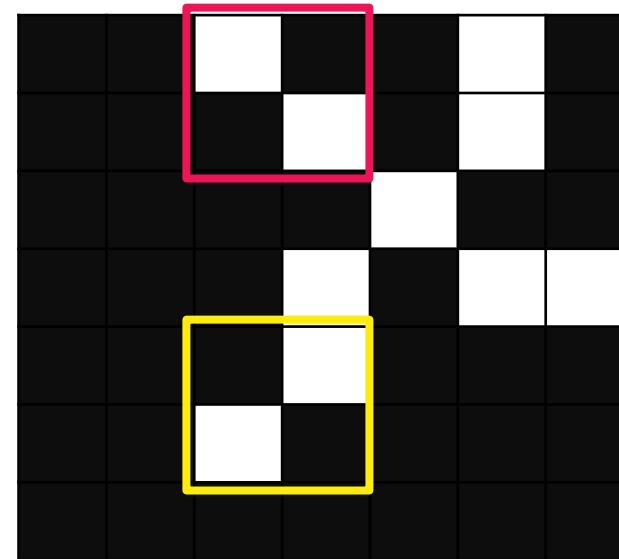
# Convolution

## Why is it important to use Convolution?

- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



The letter “X” in a binary image



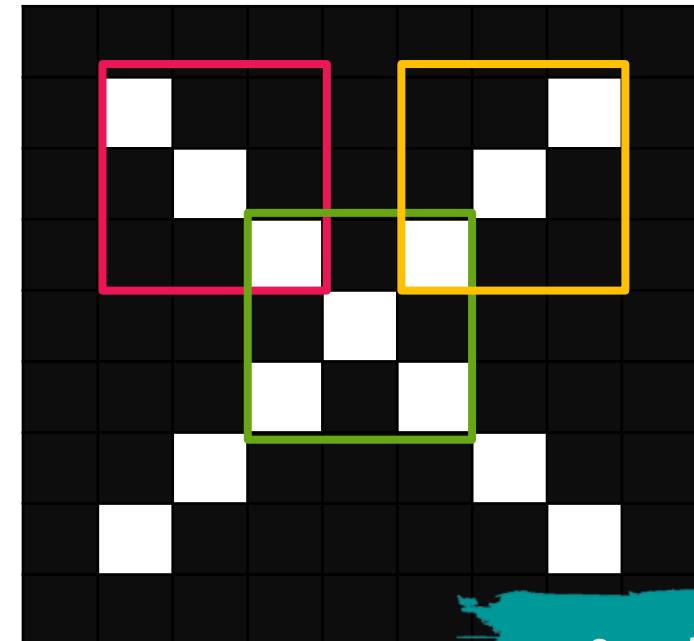
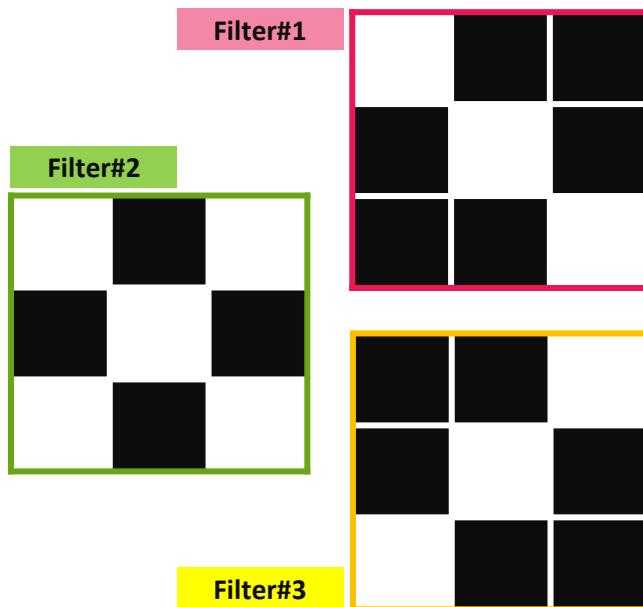
The scaled version

The same features!

# Convolution

Why is it important to use Convolution?

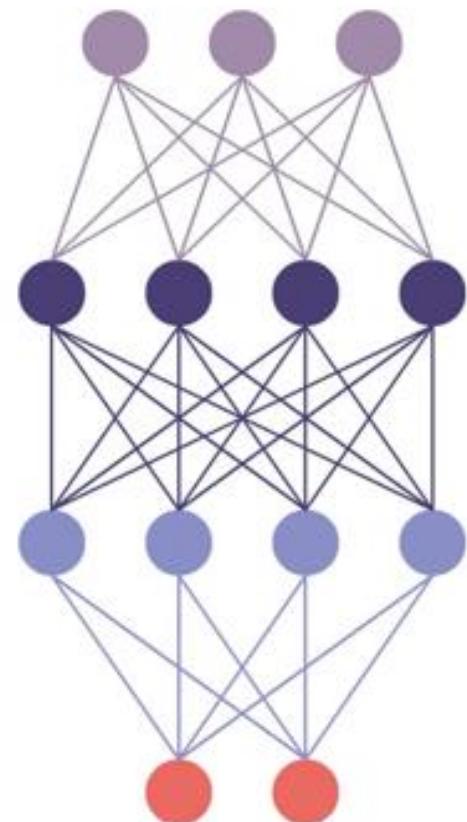
- ▶ A simple case study: classifying the letter “X” even if it is rotated/shifted



Sample filters

# Convolutional Neural Networks

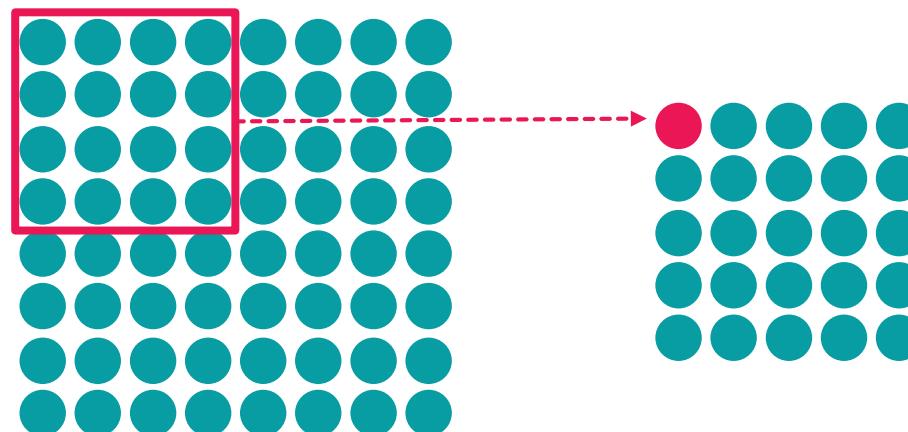
- ▶ Fully-connected Neural Networks (Dense NNs)
  - ▶ Recall: [Session#1 - Basics](#)
- ▶ Can we use them for processing mage inputs?
  - ▶ Regular NNs do not **scale well** to full images
  - ▶ We need many **calculations** and parameters!
  - ▶ We will lose the **spatial information**
    - ▶ **Reason:** flattening the 2D image into 1D vector
    - ▶ **Example:** the nearness of pixels to others



# Convolutional Neural Networks

## How to preserve the spatial structure?

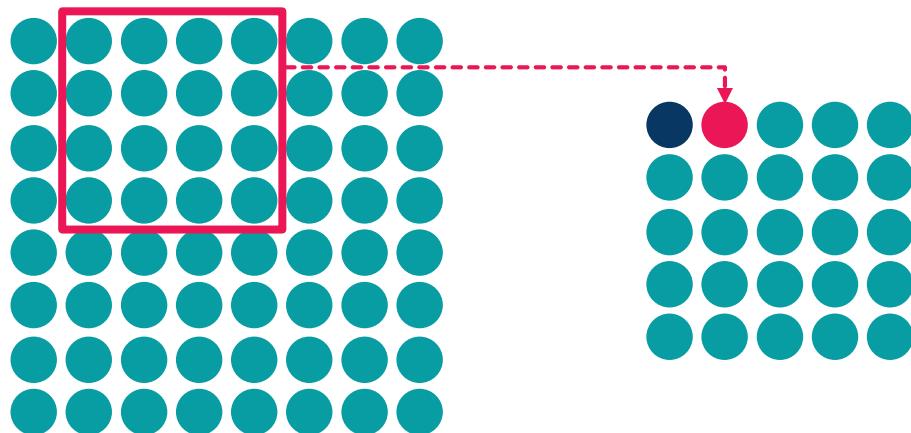
- ▶ Connecting **slices of the input**, instead of connecting all input values to neurons in the hidden layer
  - ▶ Each patch (slice) to a single neuron in the next layer
- ▶ This approach preserves the spatial information + visual features



# Convolutional Neural Networks

## How to preserve the spatial structure?

- ▶ Connecting **slices of the input**, instead of connecting all input values to neurons in the hidden layer
  - ▶ Each patch (slice) to a single neuron in the next layer
- ▶ This approach preserves the spatial information + visual features

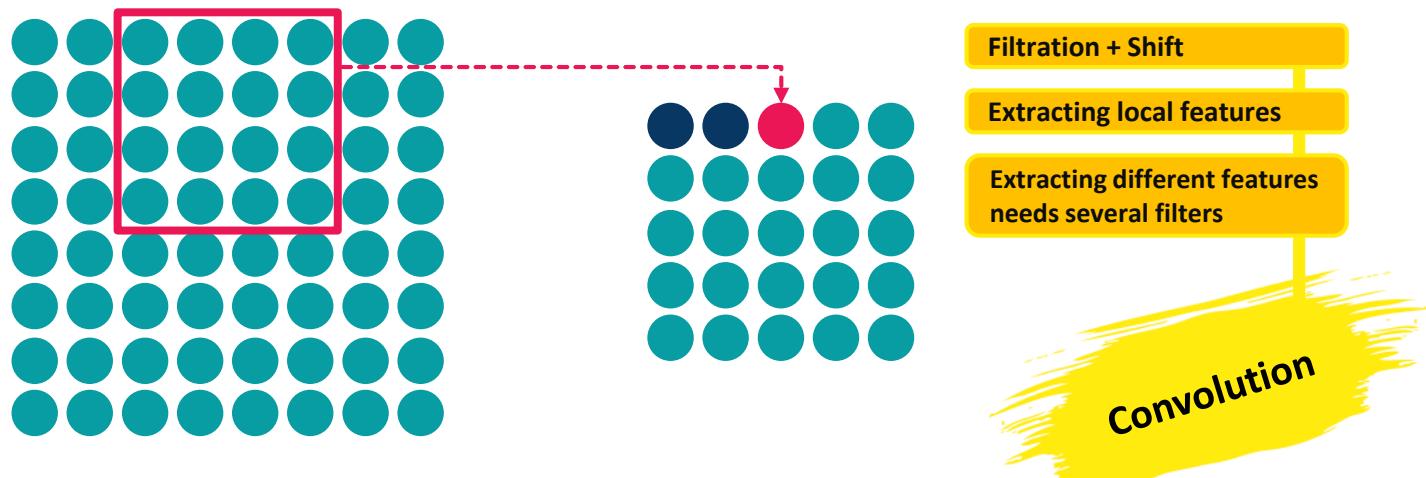


Simply, a sliding window can do the rest!

# Convolutional Neural Networks

## How to preserve the spatial structure?

- ▶ Connecting **slices of the input**, instead of connecting all input values to neurons in the hidden layer
  - ▶ Each patch (slice) to a single neuron in the next layer
- ▶ This approach preserves the spatial information + visual features

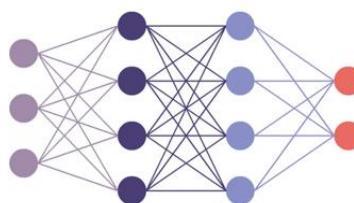


# Convolutional Neural Networks

- ▶ CNNs (*ConvNets*) utilize the mathematical foundation of convolution for **computer vision** tasks
- ▶ Inspired by the organization of the Visual Cortex
- ▶ Much lower **pre-processing** comparing to other classification algorithms
  - ▶ CNNs can automatically learn the filters and characteristics
- ▶ The overall process in a CNN:



Input image



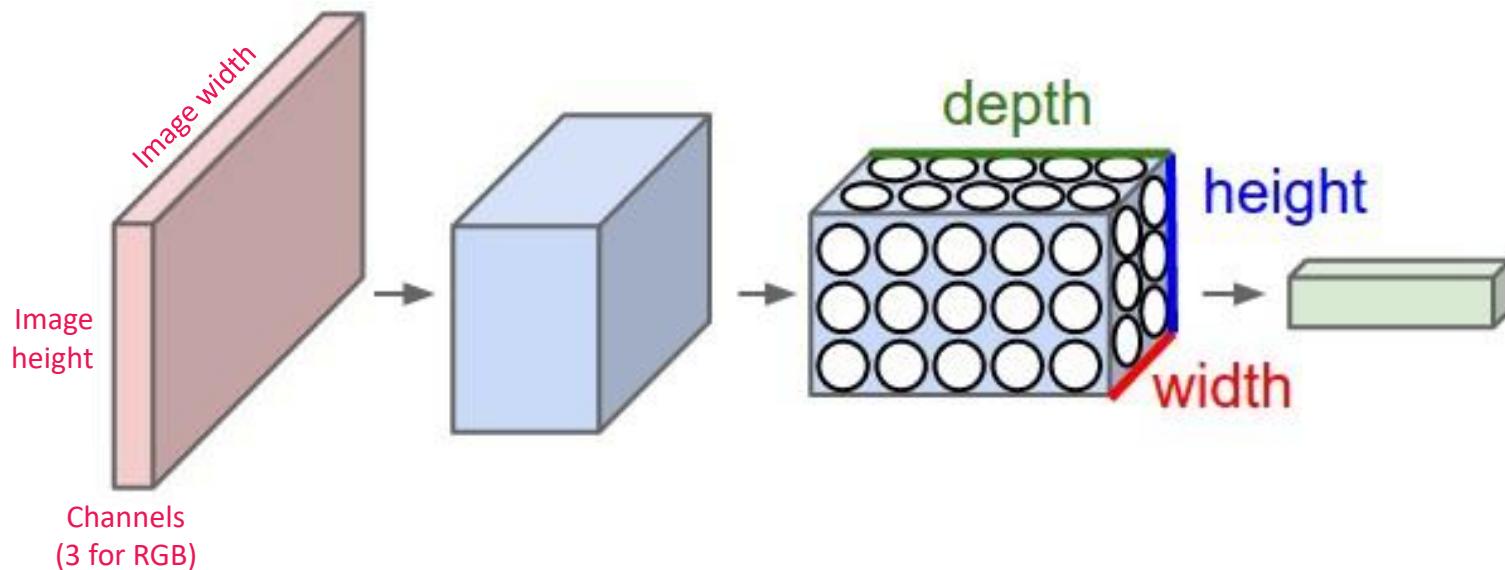
Learnable weights  
and biases



Differentiation of  
objects in the image

# Convolutional Neural Networks

- ▶ In contrast with **regular NNs**, the layers of a CNN have neurons arranged in **three dimensions**: **width, height, depth**
  - ▶ **Note:** “depth” is the third dimension of an activation volume
  - ▶ **Note:** each layer converts a 3D input to a 3D output of neurons’ AFs



# Convolutional Neural Networks

## Roles of CNNs

- ▶ Assuming that the inputs are images → more efficient network
- ▶ Reducing images into a form that is **easier for processing**
- ▶ Preserving the **critical features** of images for prediction
- ▶ Reducing the **complexity** of the image classification task

## Popular Architectures of CNNs **(Click to see the papers)**

AlexNet

GoogLeNet

VGGNet

LeNet

ResNet

# Convolutional Neural Networks

README.md



## Convolutional Neural Networks (CNNs)

CNNs can easily handle sequential data processing.



### Codes

#	File	Description
0	<a href="#">CNN layers</a>	Introduction to Keras CNN layers

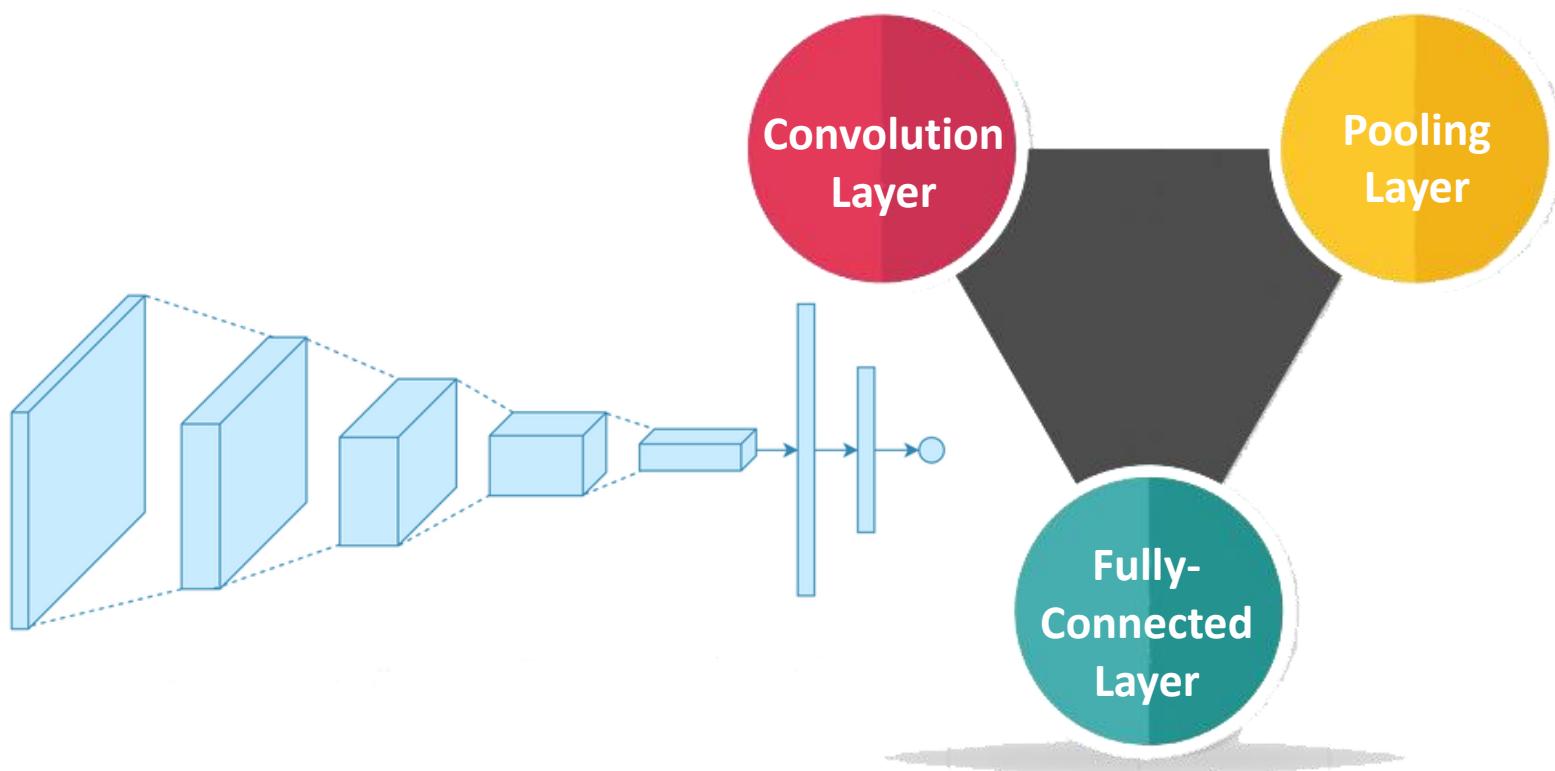
### ✖ Use cases (to be added)

- Abandoned Object Detection
- Text Translation
- Face Recognition
- Vehicle Detection
- Pedestrian Detection

Full code on GitHub



# CNNs Layers



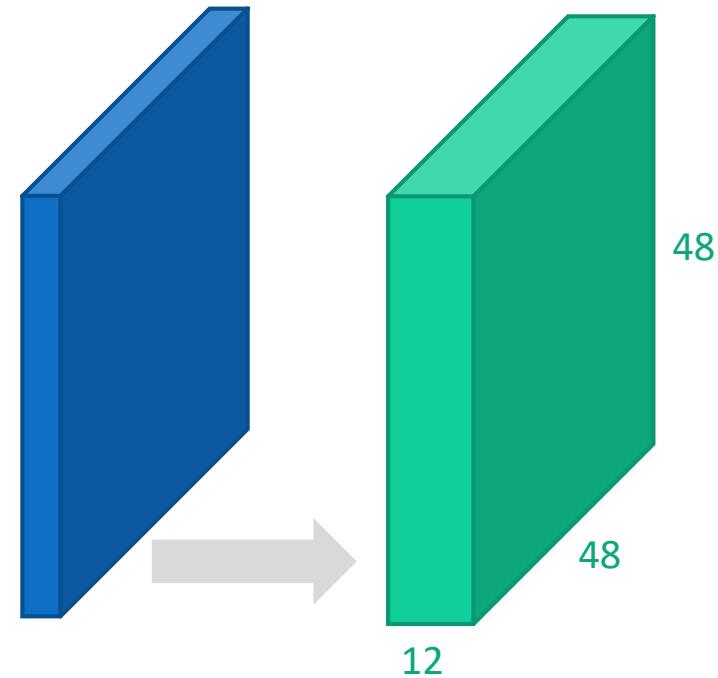
# CNNs Layers

## Convolution Layer (Kernel)

- ▶ **Output:** Feature maps
- ▶ Extracting high-level features (e.g., edges)
- ▶ Specifying the characteristics of the Kernel
  - ▶ Including Kernel size, Stride, Padding



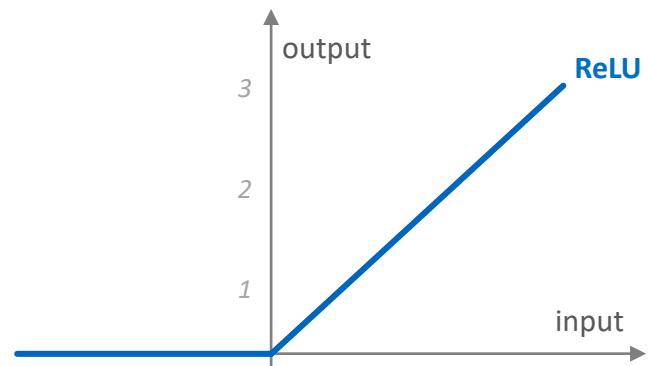
Input Image  
(e.g. [48× 48×3])



# CNNs Layers

## Convolution Layer (Kernel)

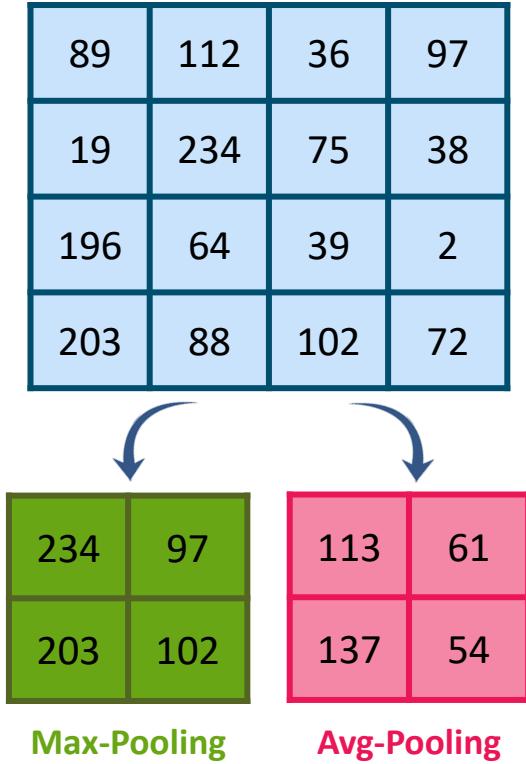
- ▶ Then, adding a supplementary step to the convolution operation
  - ▶ **Activation Function:** often **ReLU** (replacing negative values by zero)
  - ▶ **Note:** The Rectified Linear Unit is not a separate component of the convolutional neural networks' process



# CNNs Layers

## Pooling Layer (Down sampling)

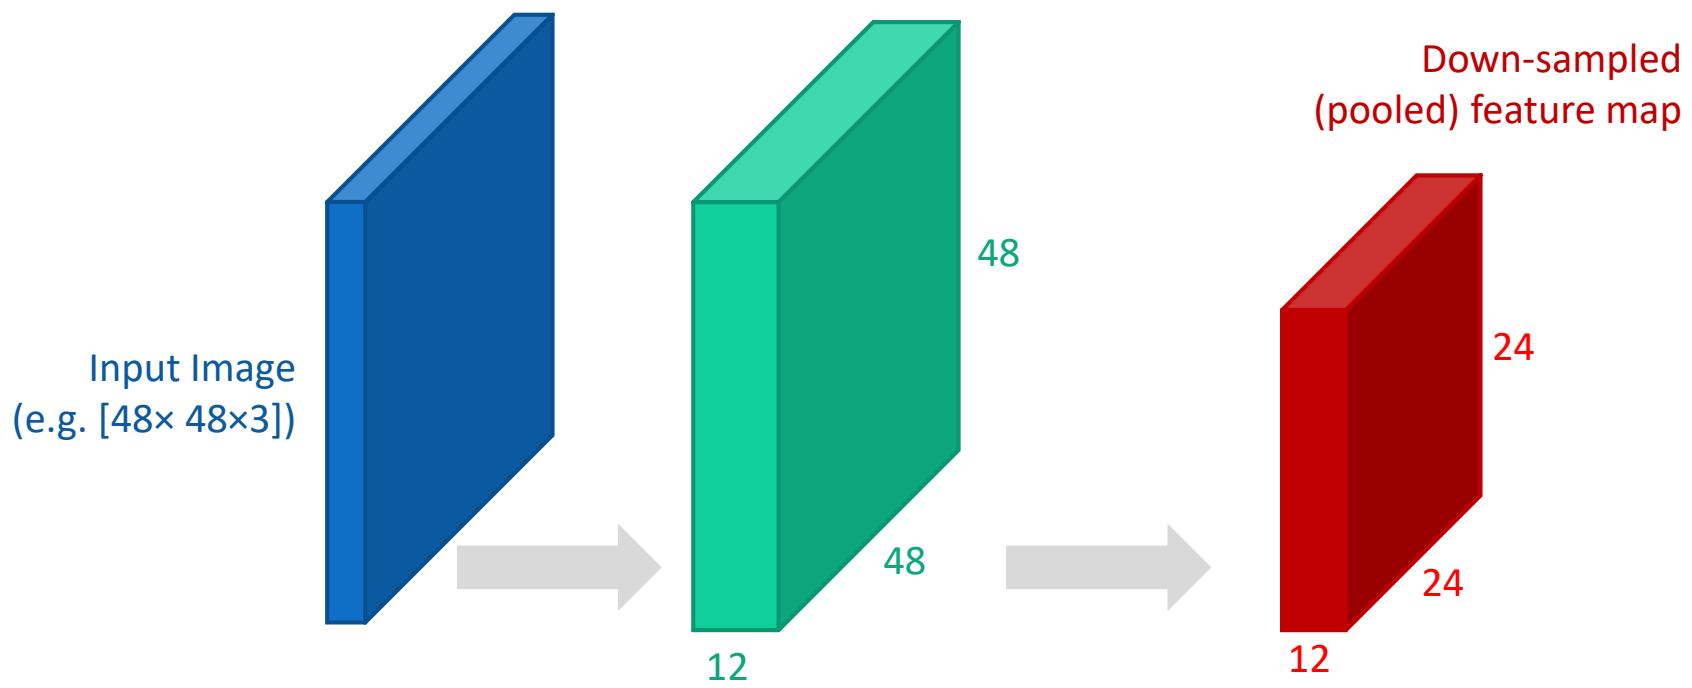
- ▶ **Output:** Pooled feature maps
- ▶ With a Convolution layer, forms a layer in a CNN
- ▶ Reducing the spatial size of the Convolved Feature
- ▶ Decreasing the dimensionality and computation
- ▶ Two-types of Pooling:
  - ▶ **Max-Pooling:** returns the maximum value
    - ▶ Recommended due to de-noising advantage
  - ▶ **Average-Pooling:** returns the average of all the values



# CNNs Layers

## Pooling Layer (Down sampling)

Kernel with  
12 filters



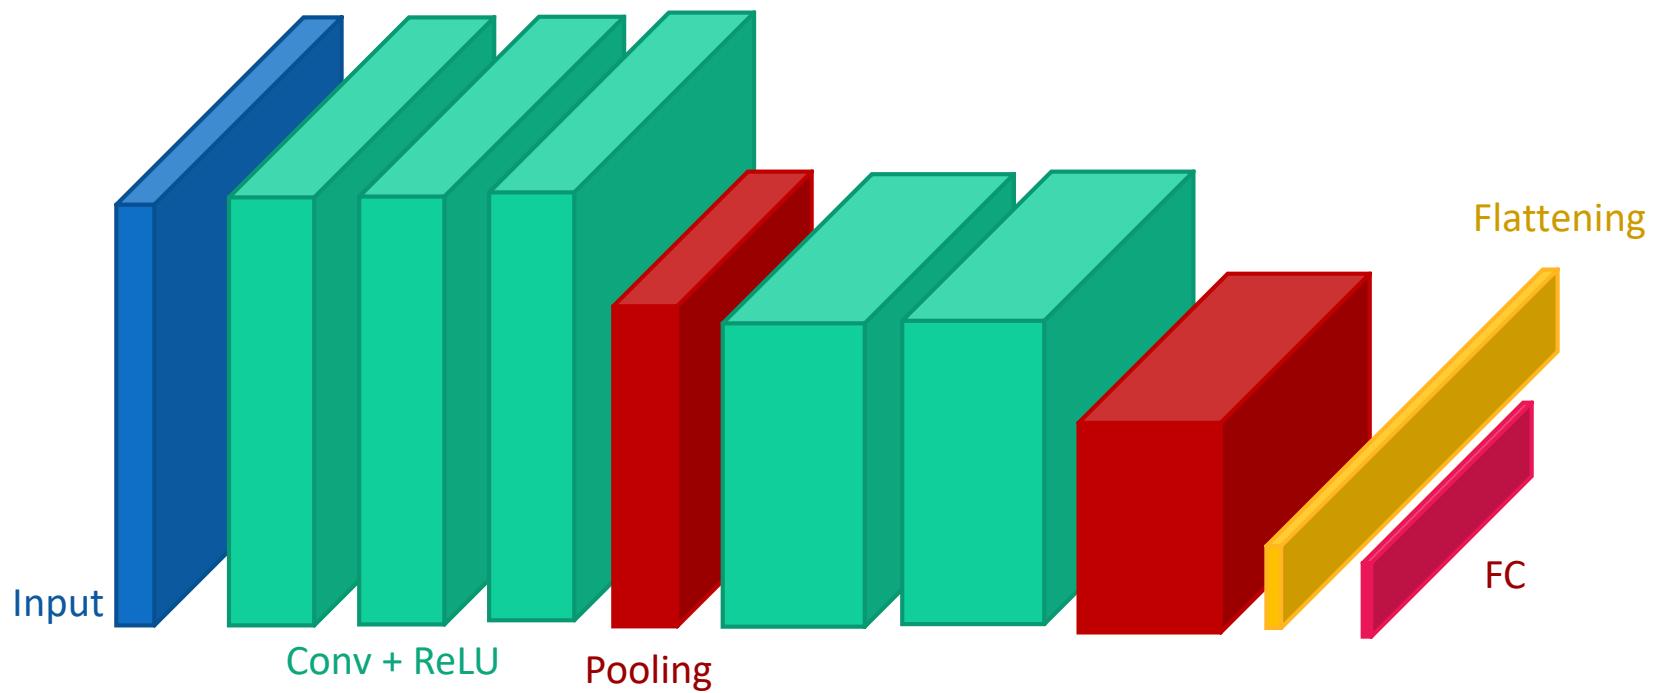
# CNNs Layers

## Fully Connected (FC) Layer

- ▶ **Output:** Mapped class scores
- ▶ Optimizing and computing class scores
  - ▶ **Note:** the input of this layer should be flattened
  - ▶ **Note:** the output can be normalized using SoftMax algorithm
- ▶ Each neuron is connected to all nodes in the previous step
- ▶ Generally found at the final layers of CNNs
  - ▶ At some points, they can be used along with ReLU layers
  - ▶ The last FC layer is the output layer

# CNNs Layers

Let's put them all together!



# CNNs Layers

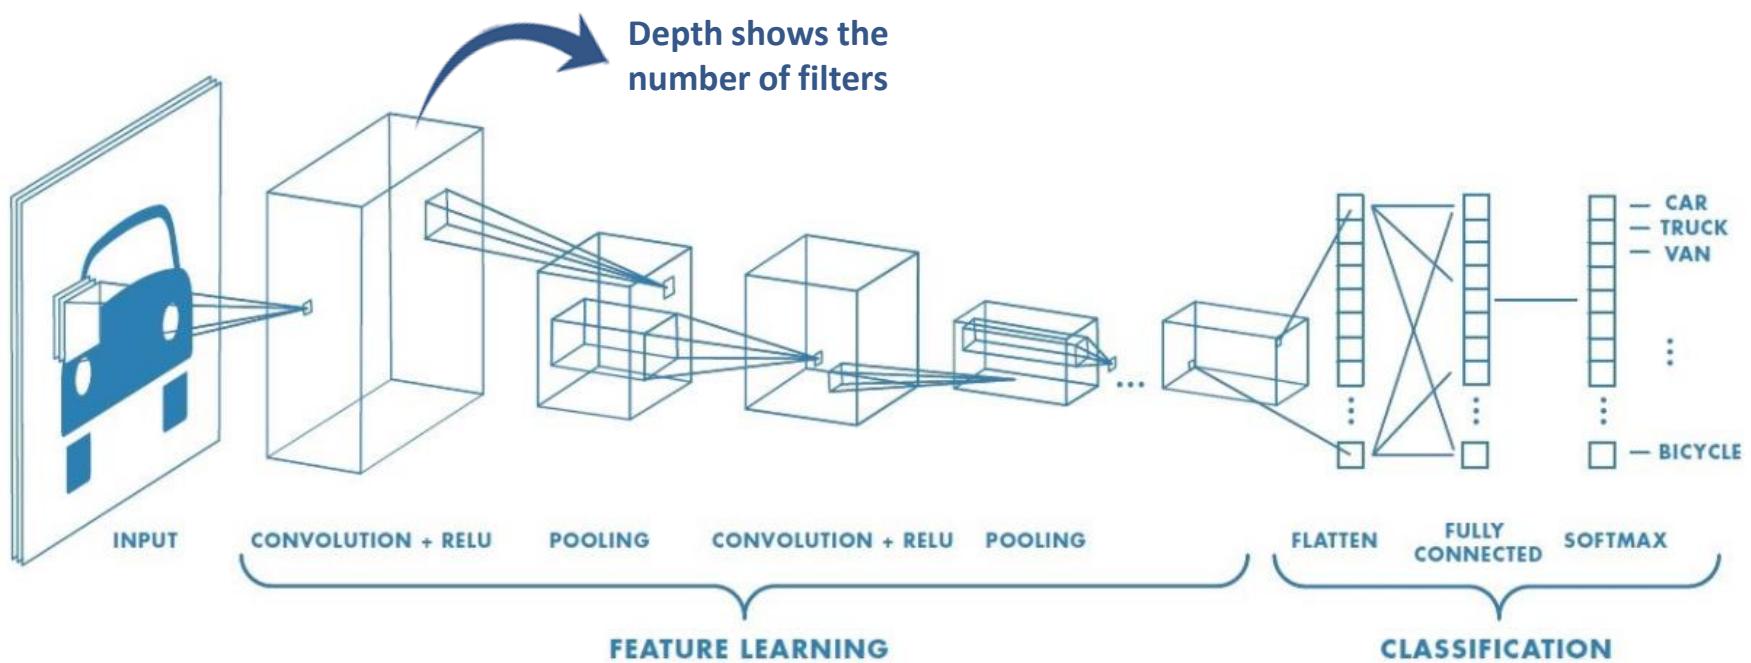
*Inp → ((Conv → ReLU) \* n → Pool?) \* m → (FC → ReLU) \* k → FC*



The magic  
pattern!

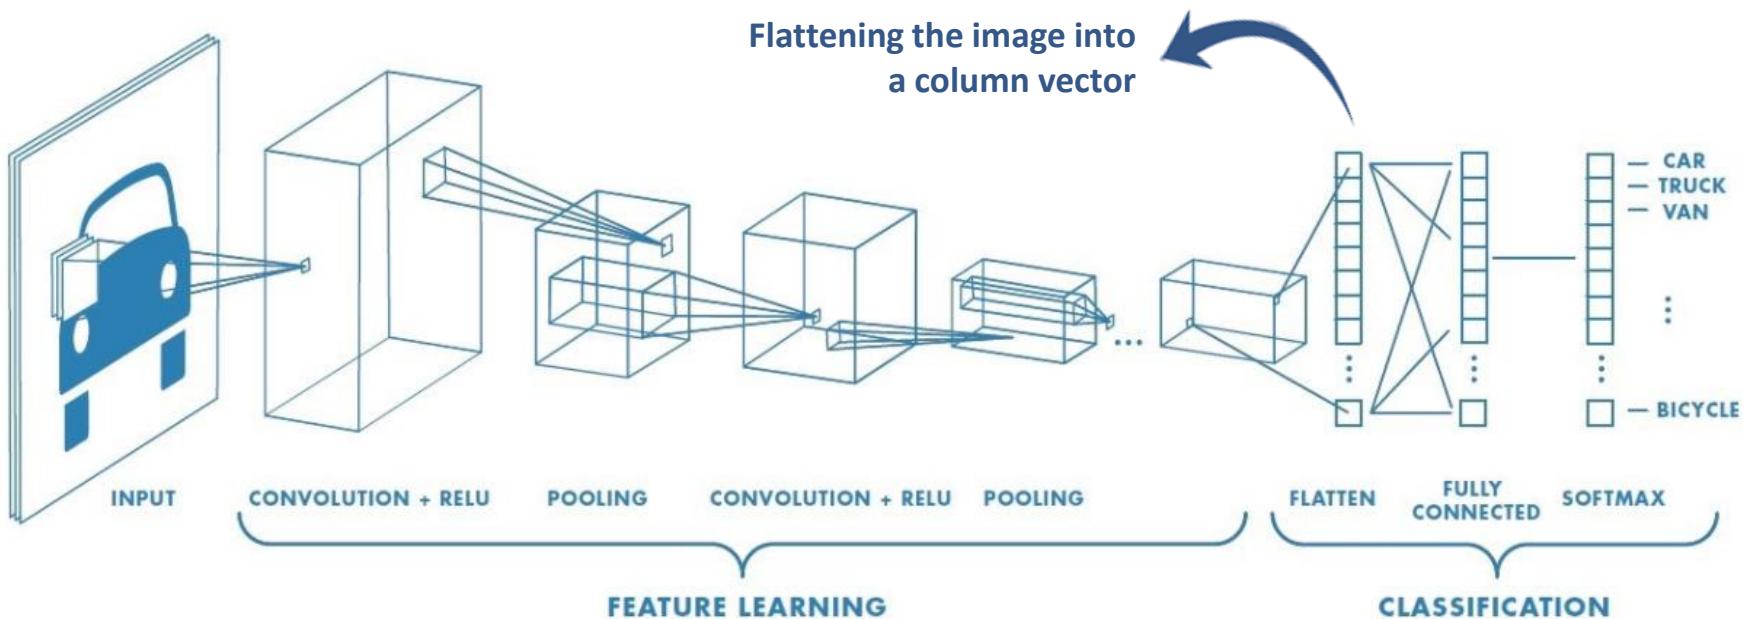
# CNNs Layers

A CNN for image classification



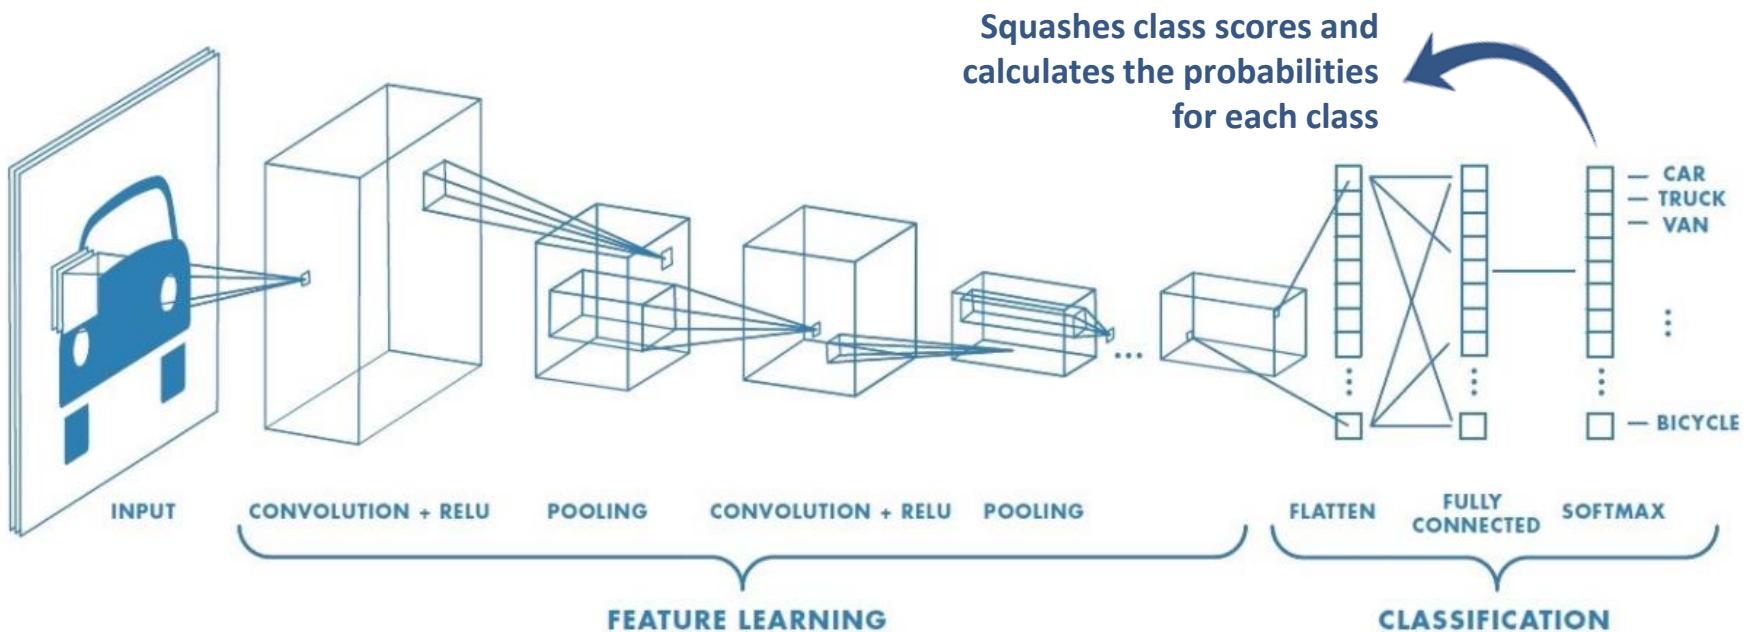
# CNNs Layers

A CNN for image classification



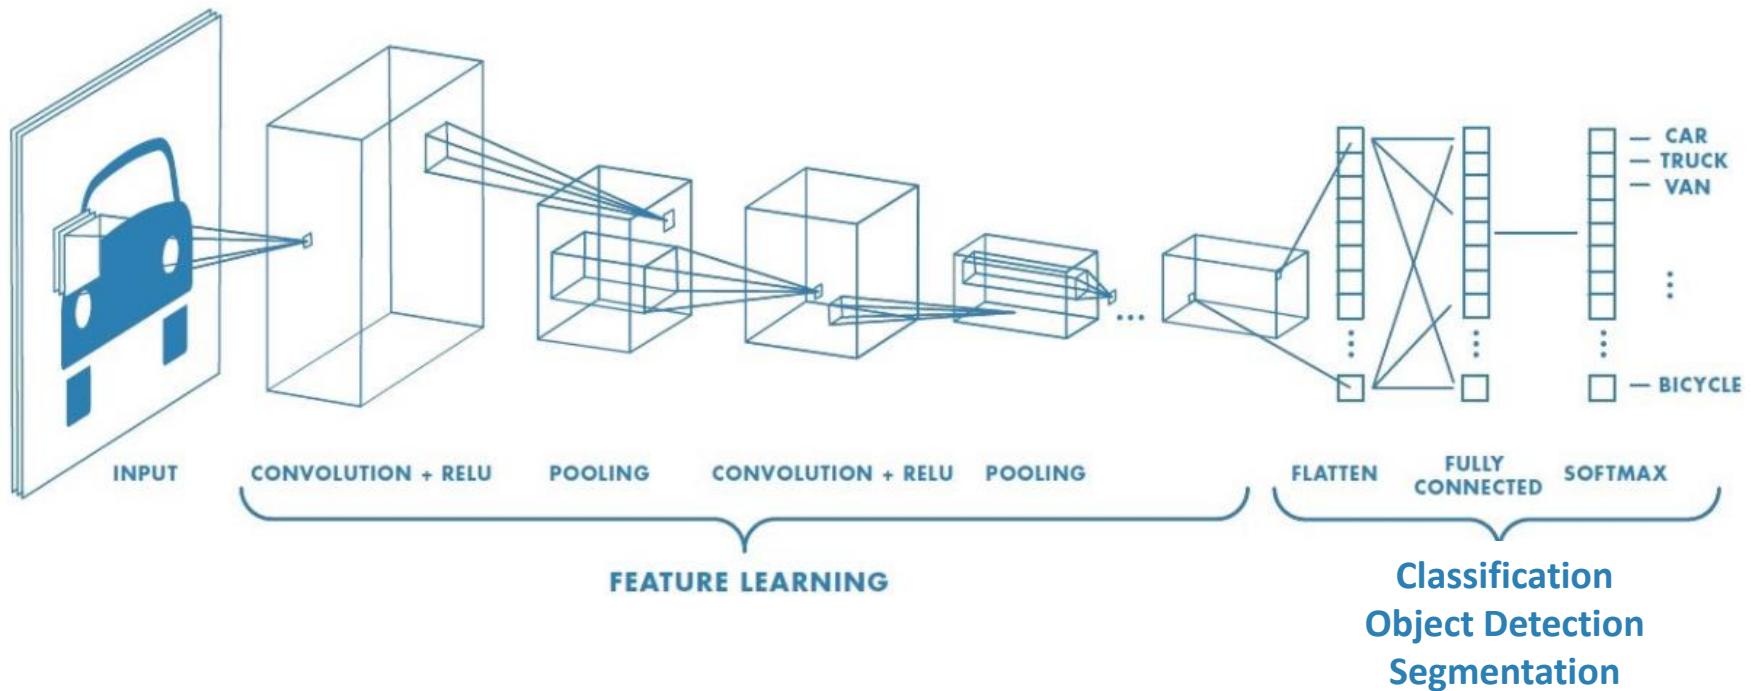
# CNNs Layers

A CNN for image classification



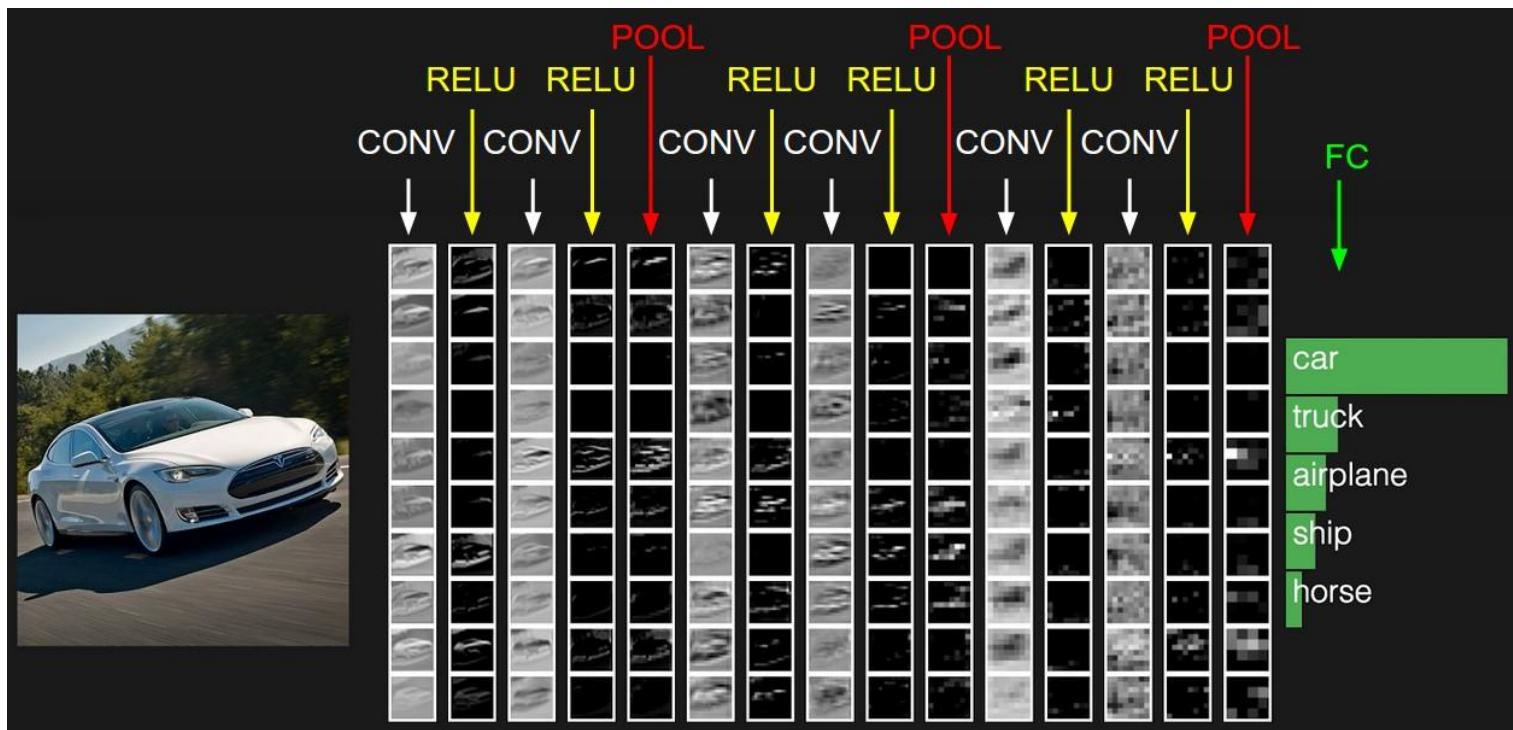
# CNNs Layers

A CNN for image classification



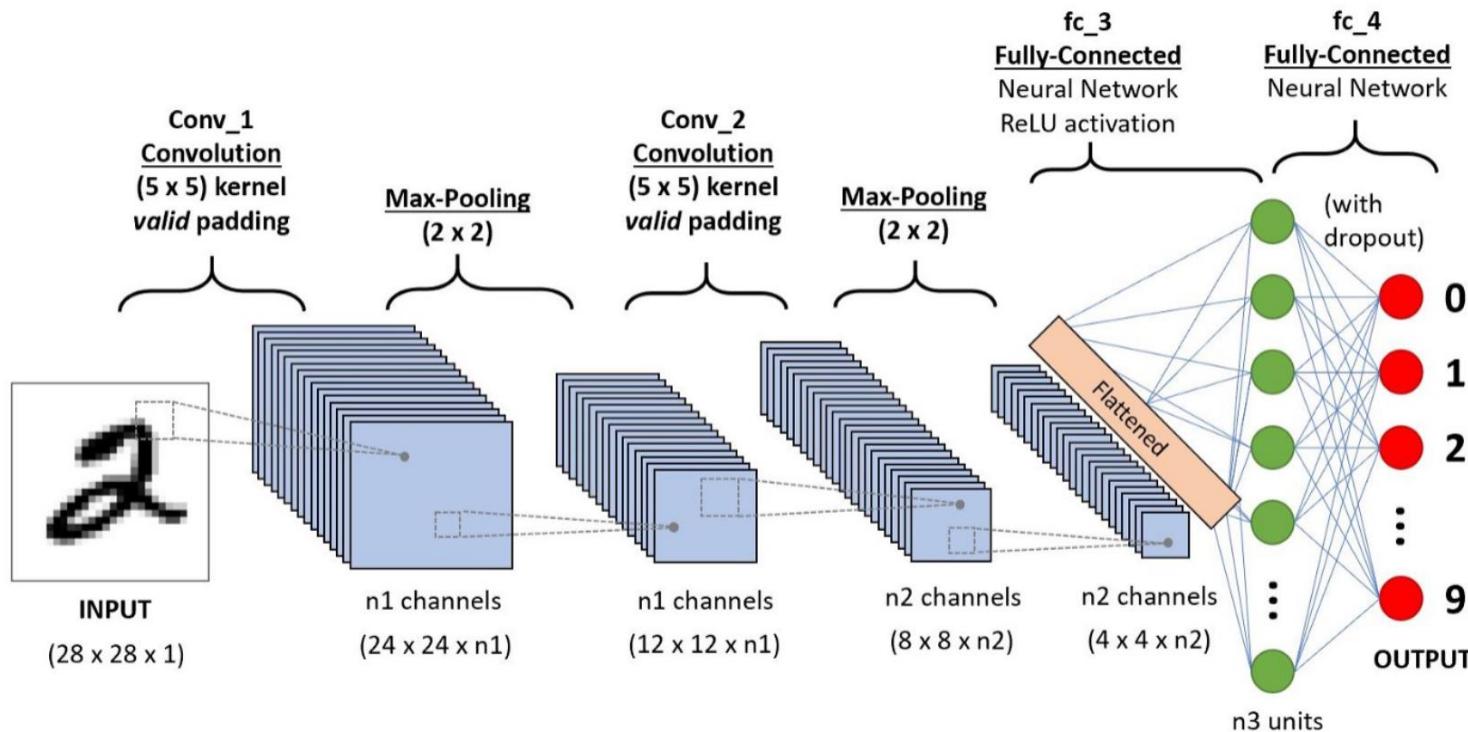
# CNNs Layers

A CNN for image classification – **Vehicle Classification**



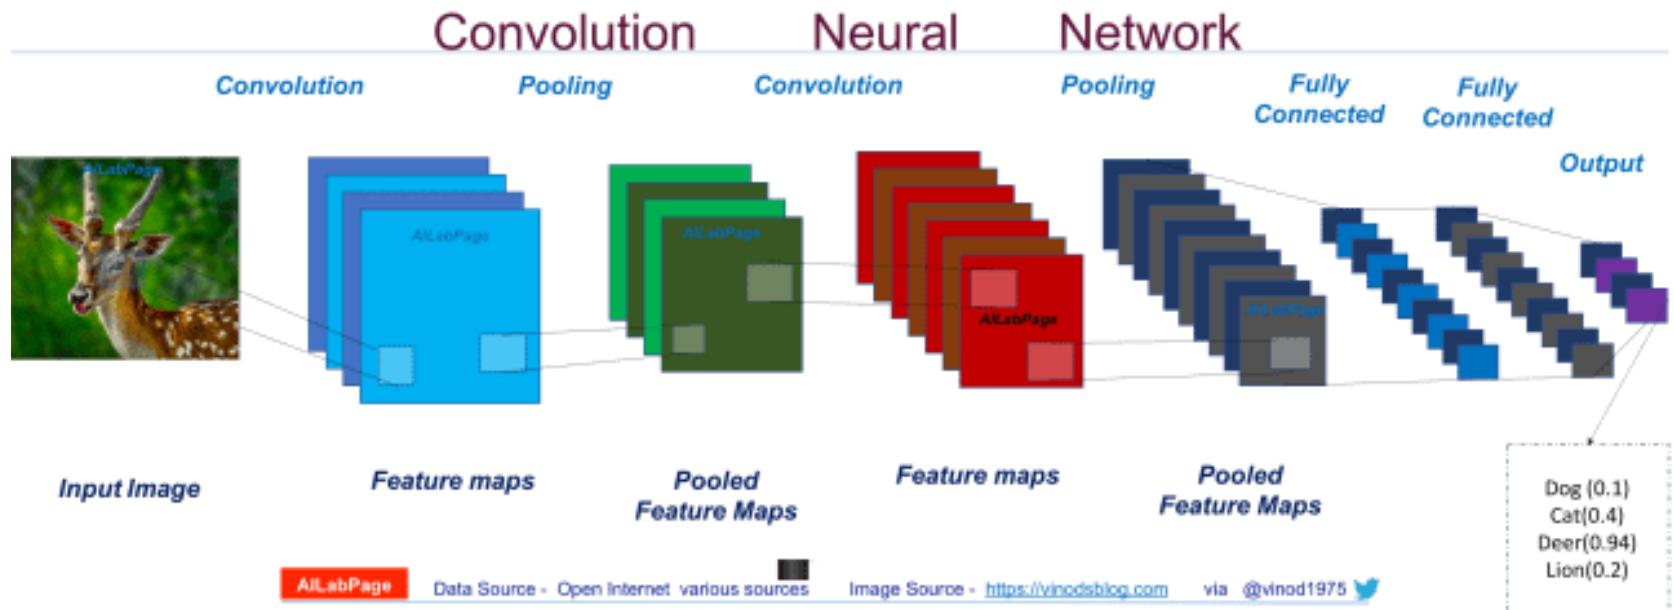
# CNNs Layers

A CNN for image classification – **Handwritten digits classification**



# CNNs Layers

A CNN for image classification – [Image classification](#)



# References

- ▶ <http://introtodeeplearning.com/>
- ▶ <https://towardsdatascience.com/image-feature-extraction-traditional-and-deep-learning-techniques-ccc059195d04>
- ▶ <https://www.encyclopedia.com/computing/news-wires-white-papers-and-books/digital-images>
- ▶ <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>
- ▶ <https://machinelearningmastery.com/classification-versus-regression-in-machine-learning>
- ▶ <https://mathworld.wolfram.com/Convolution.html>
- ▶ <https://aishack.in/tutorials/image-convolution-examples/>
- ▶ <https://cs231n.github.io/convolutional-networks/>

# Questions?

