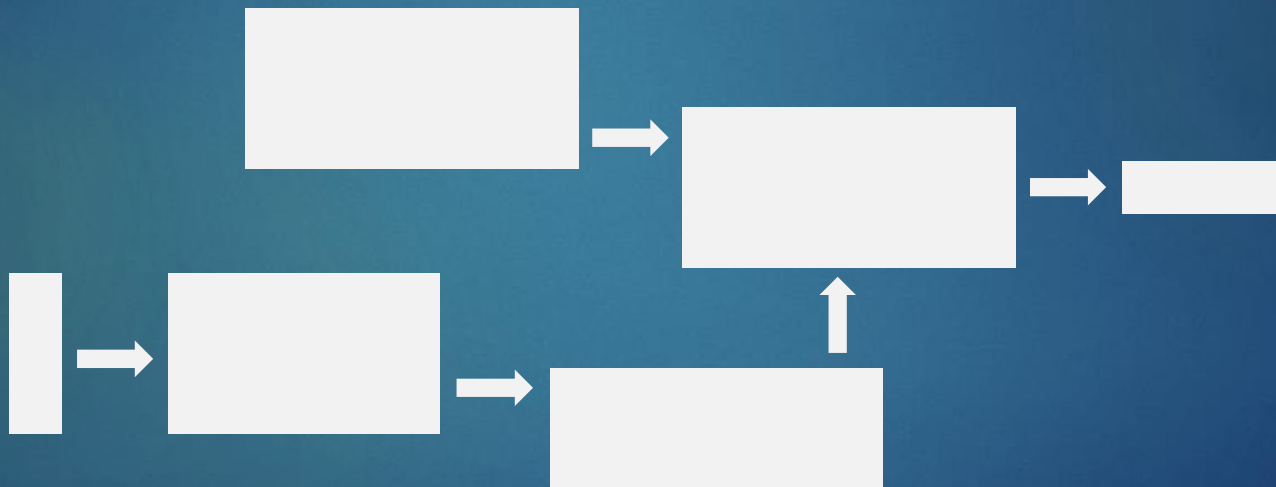


Deep Learning from Scratch

Session #8: Generative Adversarial Networks



by: **Ali Tourani – Summer 2021**

Agenda

- ▶ Supervised vs. Unsupervised Learning
- ▶ Generative Modeling
- ▶ Latent Variable Models
- ▶ Autoencoders
- ▶ Variational Autoencoders
- ▶ Generative Adversarial Networks

Supervised vs. Unsupervised Learning

Supervised Learning

- ▶ Learning which takes place in the presence of a teacher!
- ▶ The algorithm learns from **labeled** training data
 - ▶ Refers to the observed labels (classes) when faces unforeseen data
- ▶ **Challenges?**
 - ▶ The tough process of building models
 - ▶ Needs technical data science expertise
 - ▶ Labeling many data is a disaster!

Regression

Classification



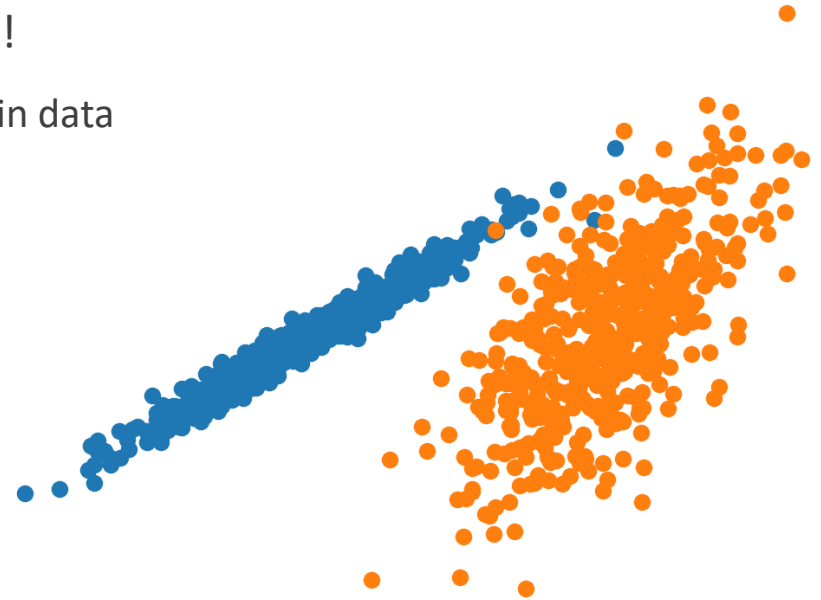
Supervised vs. Unsupervised Learning

Unsupervised Learning

- ▶ Let's allow the model to work on its own and **discover** information!
- ▶ We do not need **labeled data** anymore!
 - ▶ Can find all kinds of unknown patterns in data
- ▶ **Challenges?**
 - ▶ More complex processing tasks
 - ▶ Easier data acquisition
 - ▶ Less limitations in building models

Association

Clustering



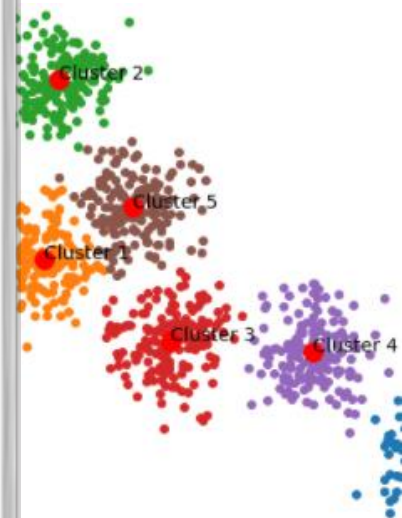
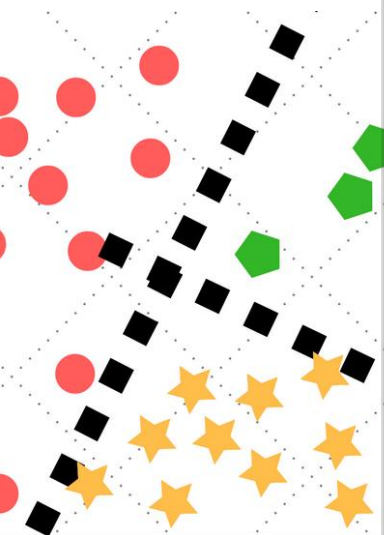
Supervised vs. Unsupervised Learning

Supervised Learning

- Access to the pairs of *(data, label)*
- Trying to find a way to *map data to label*
- Use cases in regression, classification, object detection, etc.

Unsupervised Learning

- Access to only *data* and no *label*
- Trying to learn the *underlying structures* of data
- Use cases in clustering, feature reduction, etc.



Supervised vs. Unsupervised Learning

But, why is unsupervised learning **important**?

- ▶ In many cases, we do not know exactly what we are looking for! We just want to find patterns, no matters how!
- ▶ For instance, in the **Cybersecurity applications**

Supervised approach

Attacks may miss, because the machine has not seen it before



Unsupervised approach

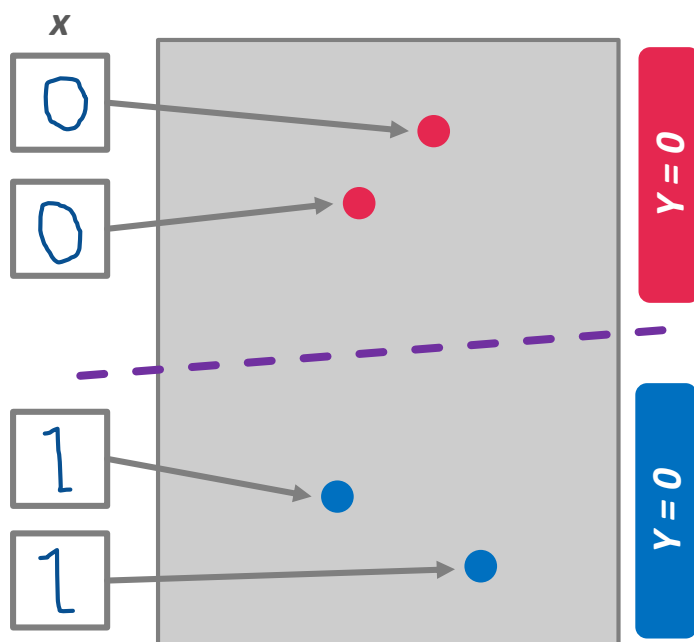
The machine tries to detect any abnormal actions

Generative Modeling

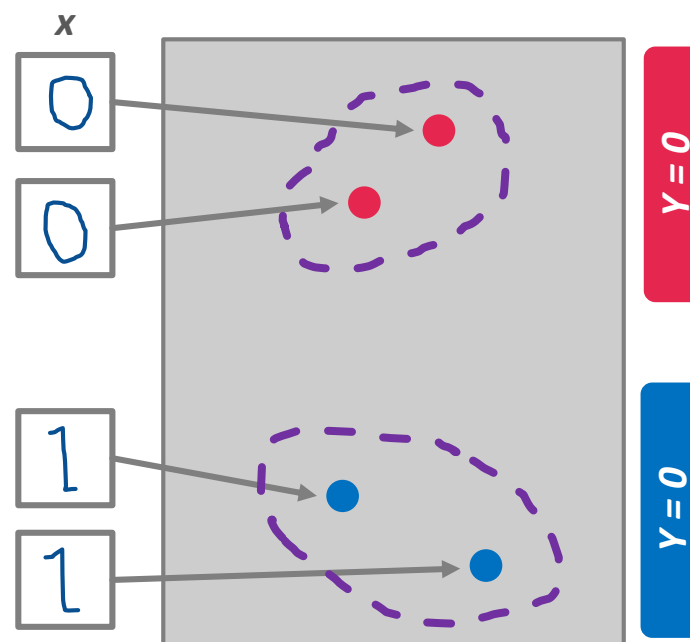
- ▶ Can we build models that can generate new samples of data?
 - ▶ **Answer:** Yes! By using **Generative Models**.
- ▶ Generative Modeling (GM):
 - ▶ Contrasts with **discriminative modeling**
 - ▶ Is the use of **AI and statistics** to produce representations of observed data
 - ▶ **Goal: Representing the distribution of the training data**
 - ▶ Is used in **unsupervised** approaches to help machines predict any probabilities
 - ▶ Utilizes **reduced understandings** of data for modeling
 - ▶ **Use case:** Models that predict the next word in a sequence

Generative Modeling

Generative vs. Discriminative modeling



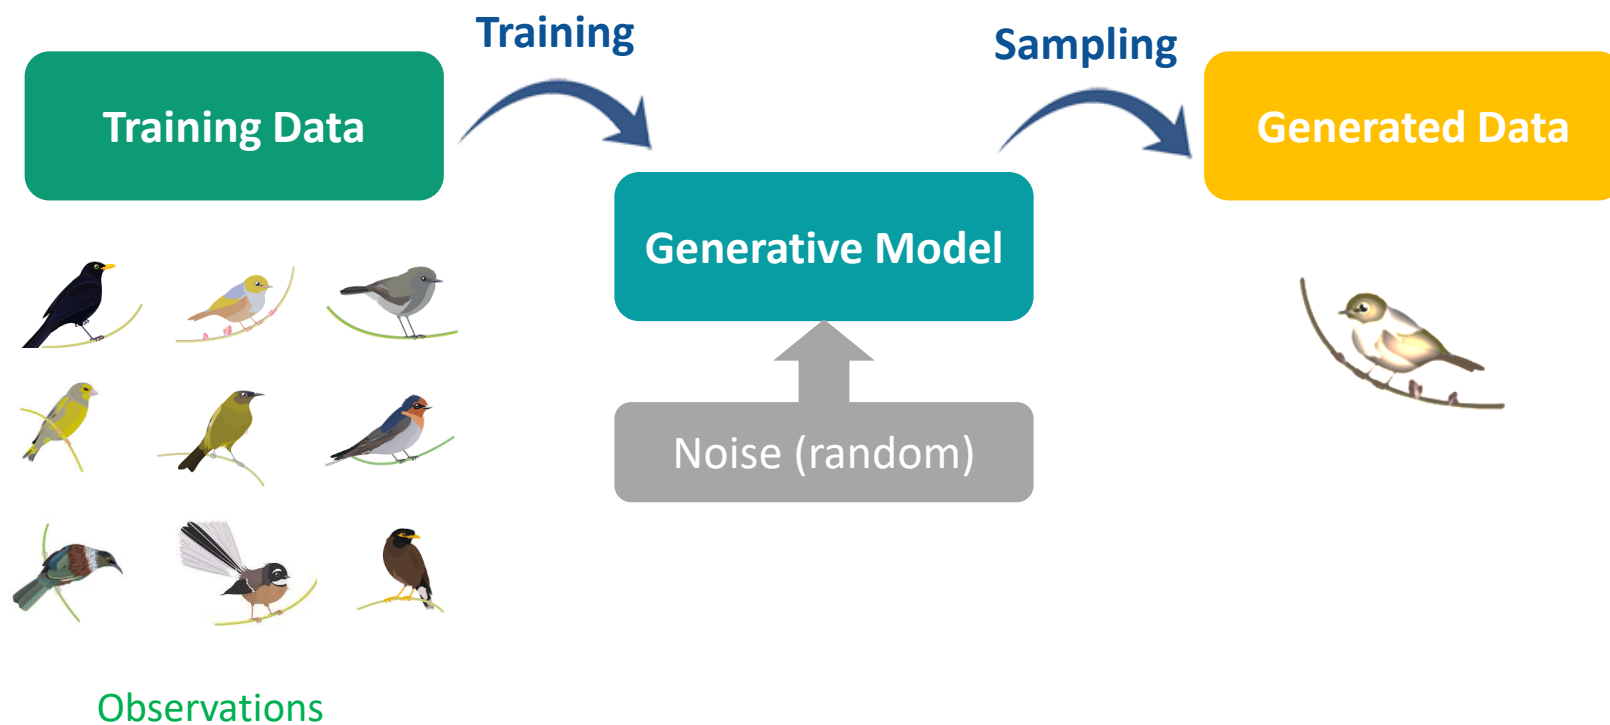
Discriminative Model



Generative Model

Trying to distinguish 1's and 0's by generating samples that fall close to their real values

Generative Modeling



Generative Modeling

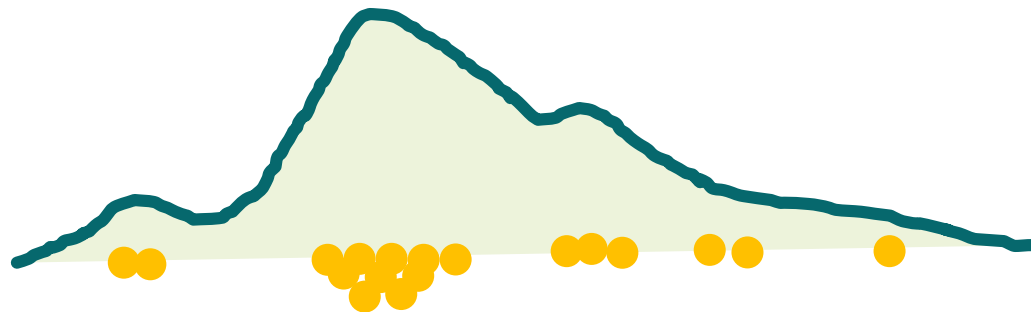
⊕ I had such a



Generative Modeling

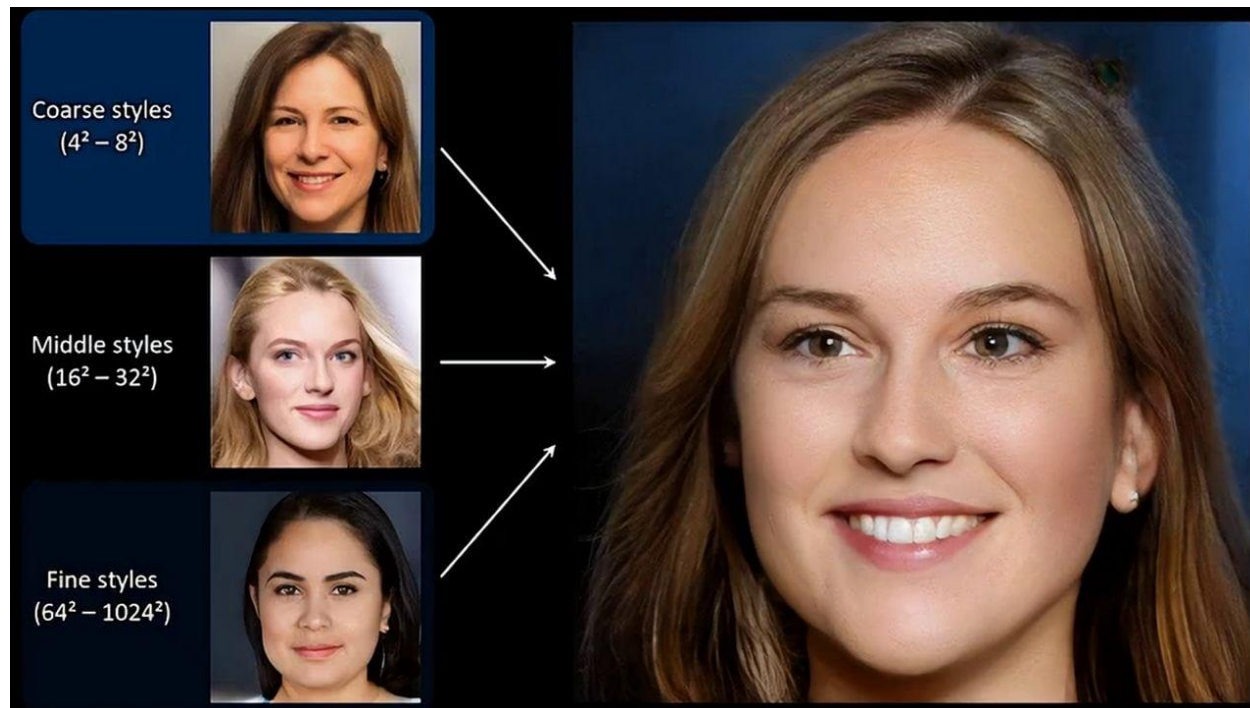
How to achieve this?

- ▶ Using the **features** extracted from training data
- ▶ The model should try to generate new sets of features
- ▶ Generated samples should follow the same (or similar) rules that created the training instances
- ▶ Building a model that **mimics the probabilistic distribution** in the training dataset



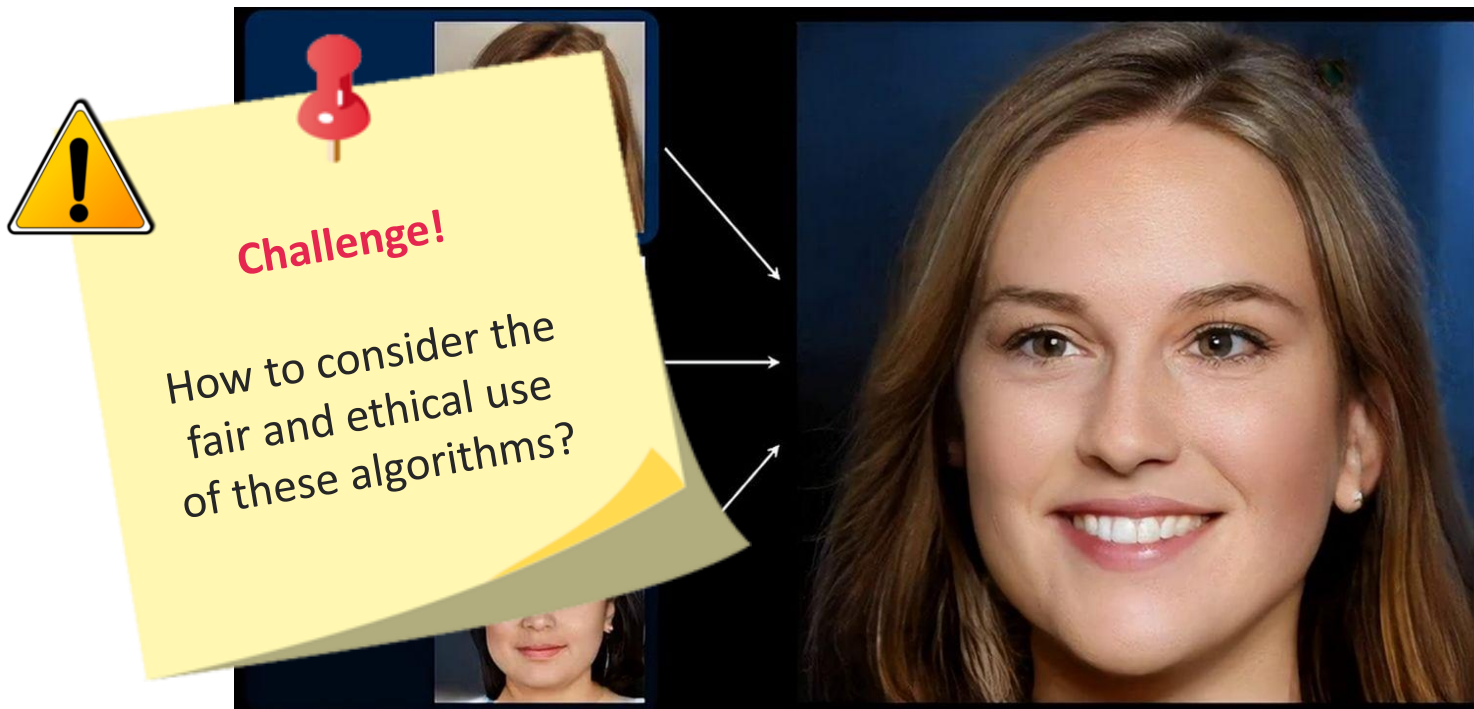
Generative Modeling

GMs can **uncover underlying features** in a dataset to create new items



Generative Modeling

GMs can **uncover underlying features** in a dataset to create new items



Generative Modeling

Outlier/anomaly detection using GMs

- ▶ Can also be used for outlier (anomaly) detection

Outliers are extreme values that deviate from other observations on data



Generative Modeling

Outlier/anomaly detection using GMs

- ▶ The goal is to detect rare or unseen items
- ▶ How to do that?
 - ▶ **Approach#1:** Leveraging generative models to detect outliers
 - ▶ **Approach#2:** Using outliers during training to improve accuracy

What we
have
trained



Some outliers
with different
features

Latent Variable Models

What are Latent Variables?

- ▶ Variables that **cannot be observed**
- ▶ They are detectable by their effects on observable data
- ▶ They can be modeled to be observed
- ▶ In Machine Learning:

A Latent Variable Model (LVM) is a probability distribution over x (**observed at learning time in the dataset**) and y (**unseen data**) variables

$$p(x, y)$$

Latent Variable Models

Why are LVMs important?


- ▶ Some data might be naturally unobserved (outliers)
- ▶ They enable us to leverage our prior knowledge when defining a model
- ▶ Using LVMs, we can learn the **Explanatory Factors** from observed data

We are trying to describe the animal using its **shadow**, while its physical appearances is not **clear** in this image




Latent Variable Models

LVMs and different types of ANNs



Autoencoders (AEs)
Variational AEs (VAEs)




Generative Adversarial
Networks (GANs)

Autoencoders

- ▶ A DNN used for **Feature (representation) Learning** in unlabeled data
 - ▶ Discovering the **representations** needed for feature detection
 - ▶ Learning a **lower-dimensional** feature representation
- ▶ Attempting to **regenerate** the input from its representation (encoding)
 - ▶ Ignoring the noise while training

Contains two modules:

- ▶ **Encoder:** maps raw data into vectors of LVs
- ▶ **Decoder:** reconstructs the observation




Autoencoders

The architecture of an autoencoder

Encoder

Decoder



Autoencoders



Important Notes on AEs


- ▶ The goal in an AE is to leverage NNs for **representation learning**
 - ▶ Having a low dimensional latent space (compression)
- ▶ AEs reconstruct the input **approximately** by keeping only the most relevant parts of data
- ▶ The code is actually a **compressed knowledge representation**
- ▶ Training process in AEs is a bit different!
 - ▶ The model is trained to **use features** to reconstruct the original data
 - ▶ Trying to **minimize the difference** between the input and reconstructed data
 - ▶ There are various approaches for this, like Mean Square Error (MSE)

Autoencoders




Important Notes on AEs


- ▶ In fact, AEs present a form of compression
 1. Compressing the input into much smaller latent space
 2. Building the input back (reconstruction)
- ▶ Lower dimensionality latent space → Poorer reconstruction output



Ground Truth



4D latent space



2D latent space

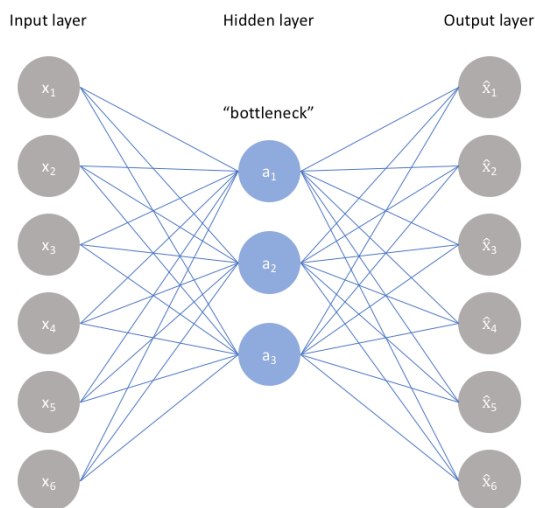
Autoencoders



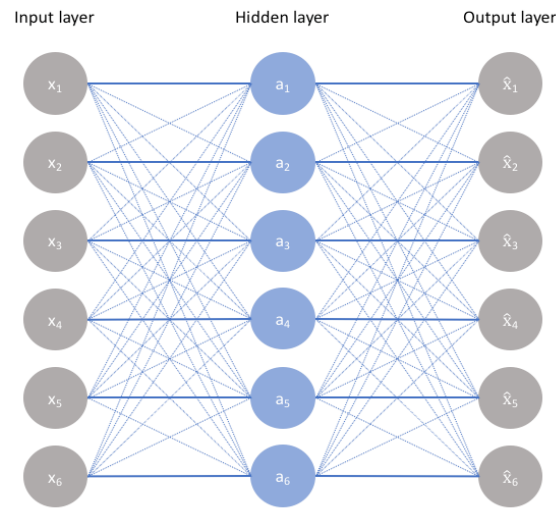
Important Notes on AEs

- ▶ Bottleneck in AEs prevents the network from **memorizing** the input values instead of **learning the compression** of the input data

**With
Bottleneck**




**Without
Bottleneck**



Autoencoders

Applications of AEs – Image Reconstruction




Autoencoders

Applications of AEs – Image Compression



Autoencoders

Applications of AEs – Dimensionality Reduction



Variational Autoencoders

- ▶ Unlike the traditional AEs, VAEs contain a **variational twist** for coding
 - ▶ A latent space with a **mixture of distributions**, instead of a fixed vector
 - ▶ Instead of direct learning of the latent variables, learns **Mean & Variance**
- ▶ A different mathematical formulation
- ▶ **Goal:** generating higher quality representations and samples
- ▶ Loss function in a VAE is a bit more complex:

$$\text{Loss} = \text{reconstruction loss} + \text{regularization loss}$$

The difference between input data
and the reconstructed output


Divergence between the inferred latent
and fixed prior on latent distributions

Variational Autoencoders

The architecture of an autoencoder

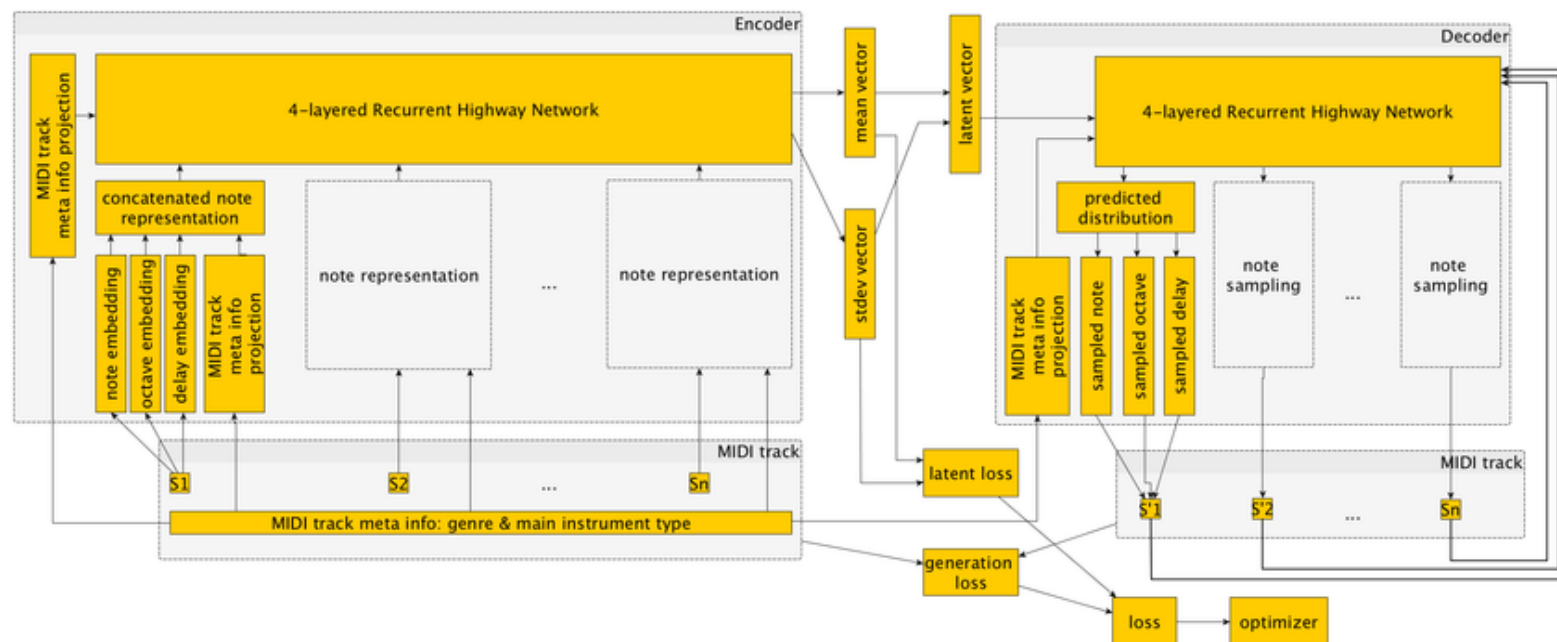
Encoder

Decoder



Variational Autoencoders

Applications of AEs – Music Generation



Variational Autoencoders

Applications of AEs – Image Generation



Variational Autoencoders

README.md

Autoencoders (AEs) and Variational Autoencoders (VAEs)

Autoencoders are deep neural network architectures used for **Feature (representation) Learning** in unlabeled data. The goal in an AE is to leverage NNs for representation learning. Autoencoders reconstruct the input approximately by keeping only the most relevant parts of data.

Encoder

Decoder


Codes

| # | File | Description |
|---|------------------------------------|--|
| 0 | Basic introduction | Introduction to Autoencoder codes in Keras |

Full code on GitHub


Generative Adversarial Networks

- ▶ A deep learning framework to generate new data **similar to the training set**
- ▶ **Goal:** sampling from a complex distribution
 - ▶ Building some approximations of this distribution
- ▶ **Idea:** starting from **random noise** (a simple data), then learning a functional **transformation** that goes from **noise to the data distribution**



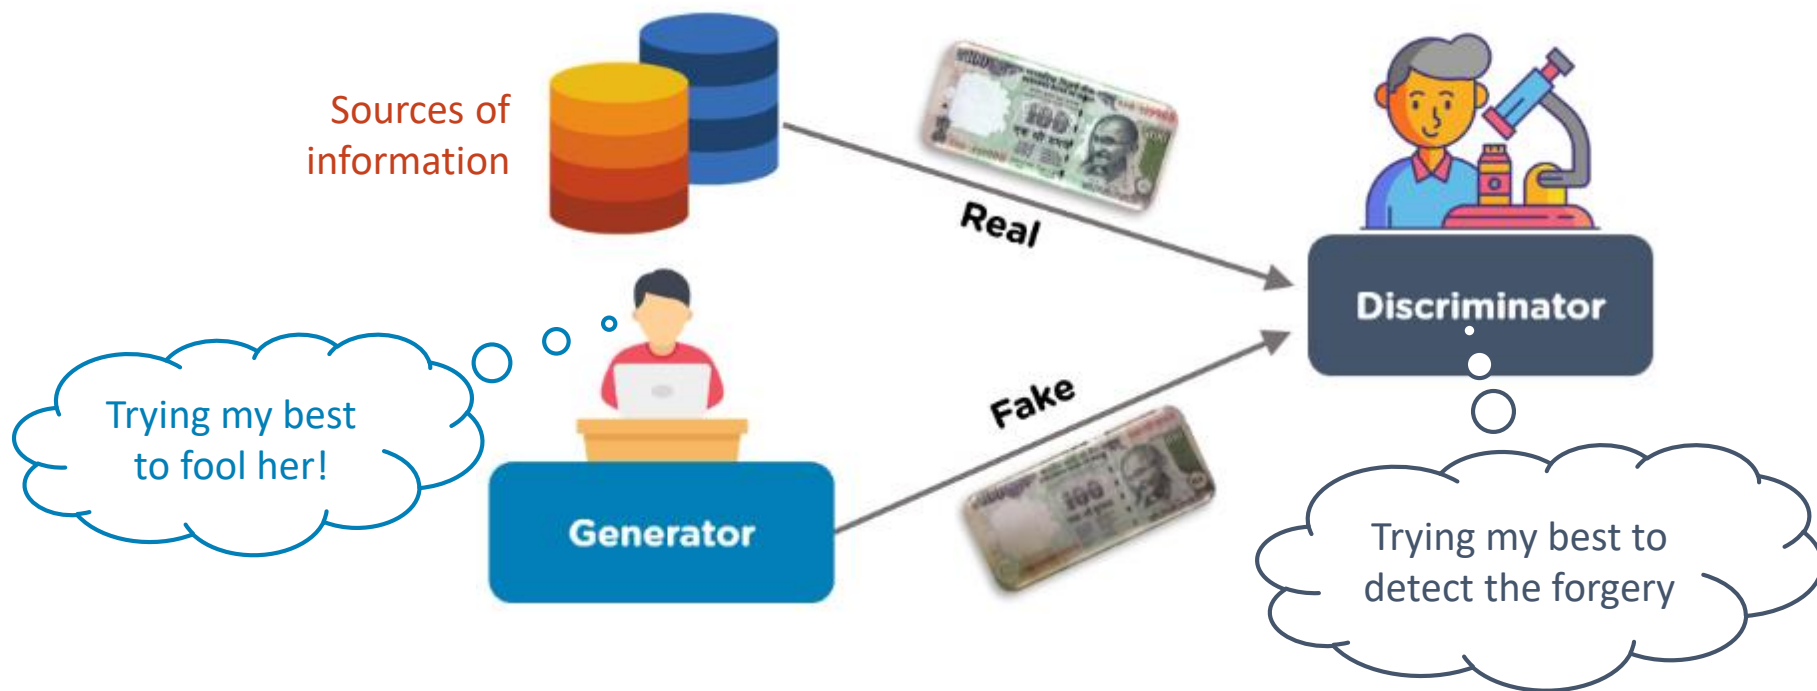
Generative Adversarial Networks

Basic idea



Generative Adversarial Networks

Basic idea



Generative Adversarial Networks

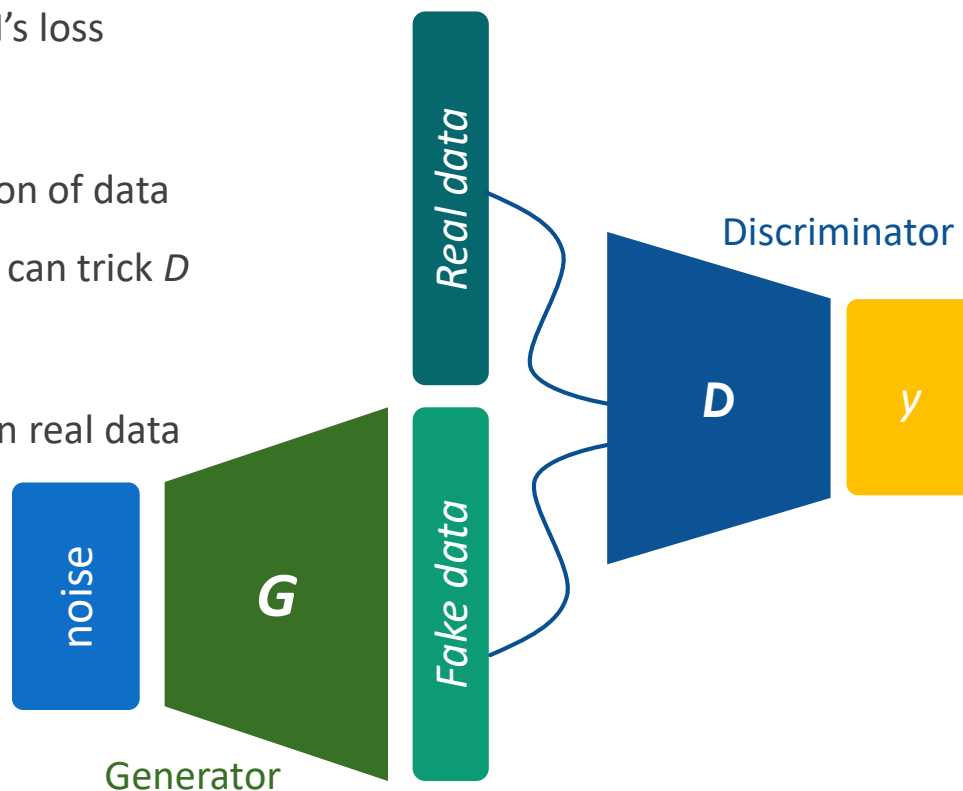
- ▶ Contains two NN that compete with each other in a game
 - ▶ One NNs' gain is another NN's loss

Generator

- ▶ Turning noise into an imitation of data
- ▶ **Goal:** generate samples that can trick D


Discriminator

- ▶ Trying to distinguish between real data and fake (generated) ones




Generative Adversarial Networks

How do GANs work?





Generator: these are birds!



Generative Adversarial Networks


How do GANs work?




Discriminator: I don't think so!

Generative Adversarial Networks

How do GANs work?




The discriminator needs to be trained over time to identify real and unreal items



Discriminator: I don't think so!

Generative Adversarial Networks

How do GANs work?





Generator: these are birds!



Real data

Generative Adversarial Networks


How do GANs work?



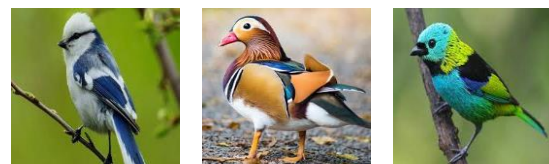
Discriminator: Hmm! Not sure yet!

Generative Adversarial Networks

How do GANs work?





Generator: these are birds!



Real data

Generative Adversarial Networks

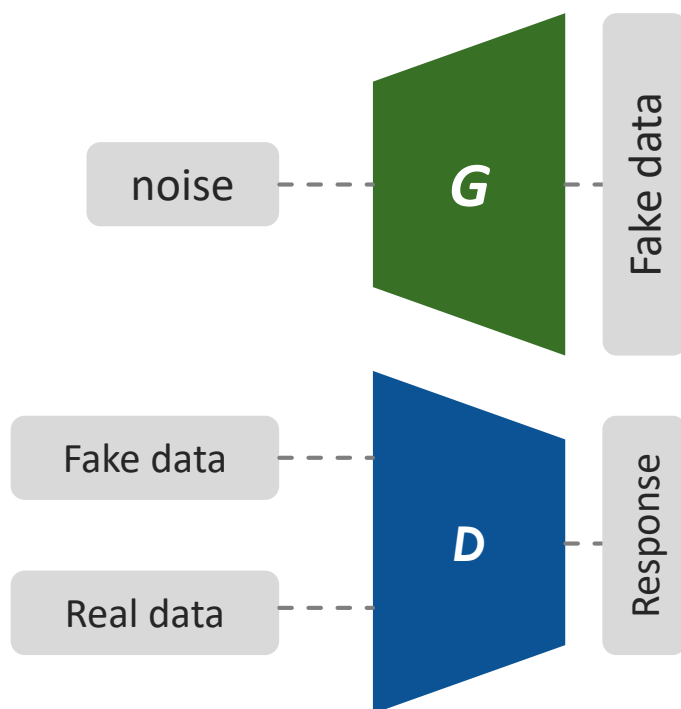
How do GANs work?



Discriminator: Maybe!

Generative Adversarial Networks

How do GANs work?



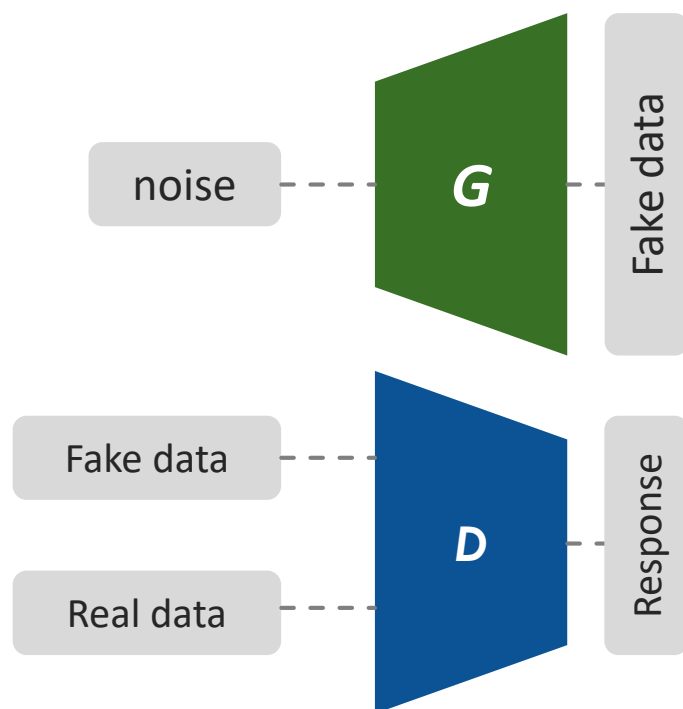
Generator: these are birds!



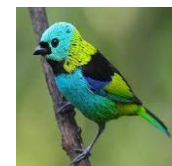
Real data

Generative Adversarial Networks

How do GANs work?



Fake data

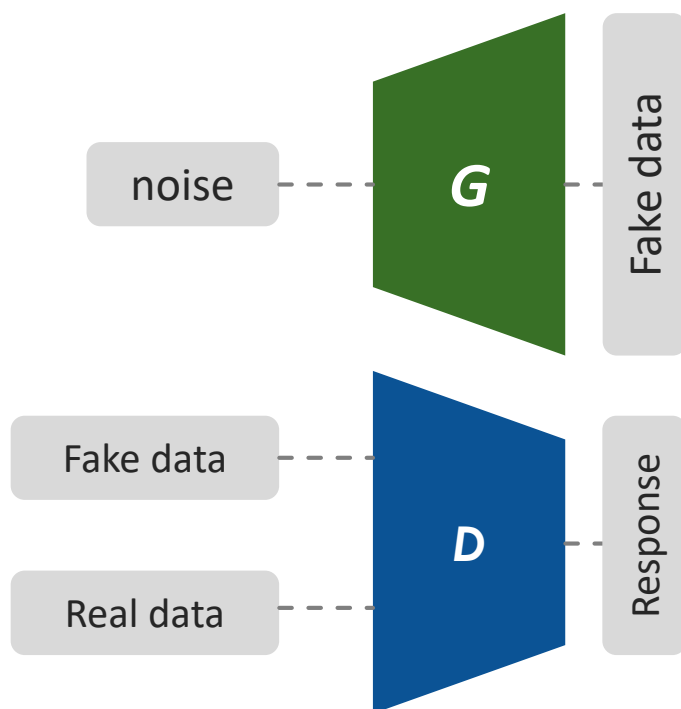


Real data

Discriminator: I can see something!

Generative Adversarial Networks

How do GANs work?



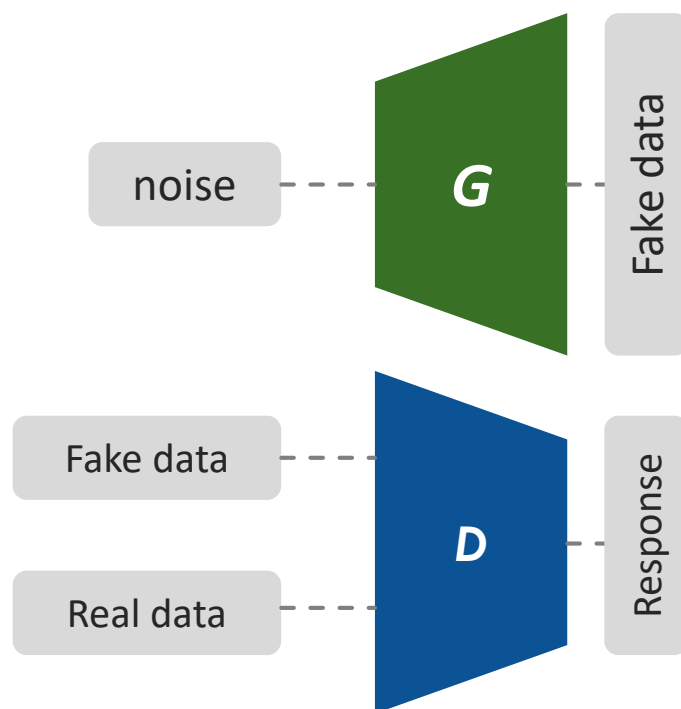
Generator: these are birds!



Real data

Generative Adversarial Networks

How do GANs work?



Fake data

Real data

Discriminator: Yes they are!

Generative Adversarial Networks



Important notes on GANs

- ▶ While training, discriminator **becomes better** to classify real and fake data, and thus, generator should try harder to trick it
- ▶ The optimum outcome is when the generator reproduces the **true data distribution**
- ▶ There are two perspectives for calculating loss function in GANs:
 - ▶ Loss (D): the max probability to correctly identify fake and real data
 - ▶ Loss (G): the min probability that D can distinguish real and fake data

$$loss = \arg \min_G \max_D E_{z,x} [\log D(G(z)) + \log(1 - D(x))]$$

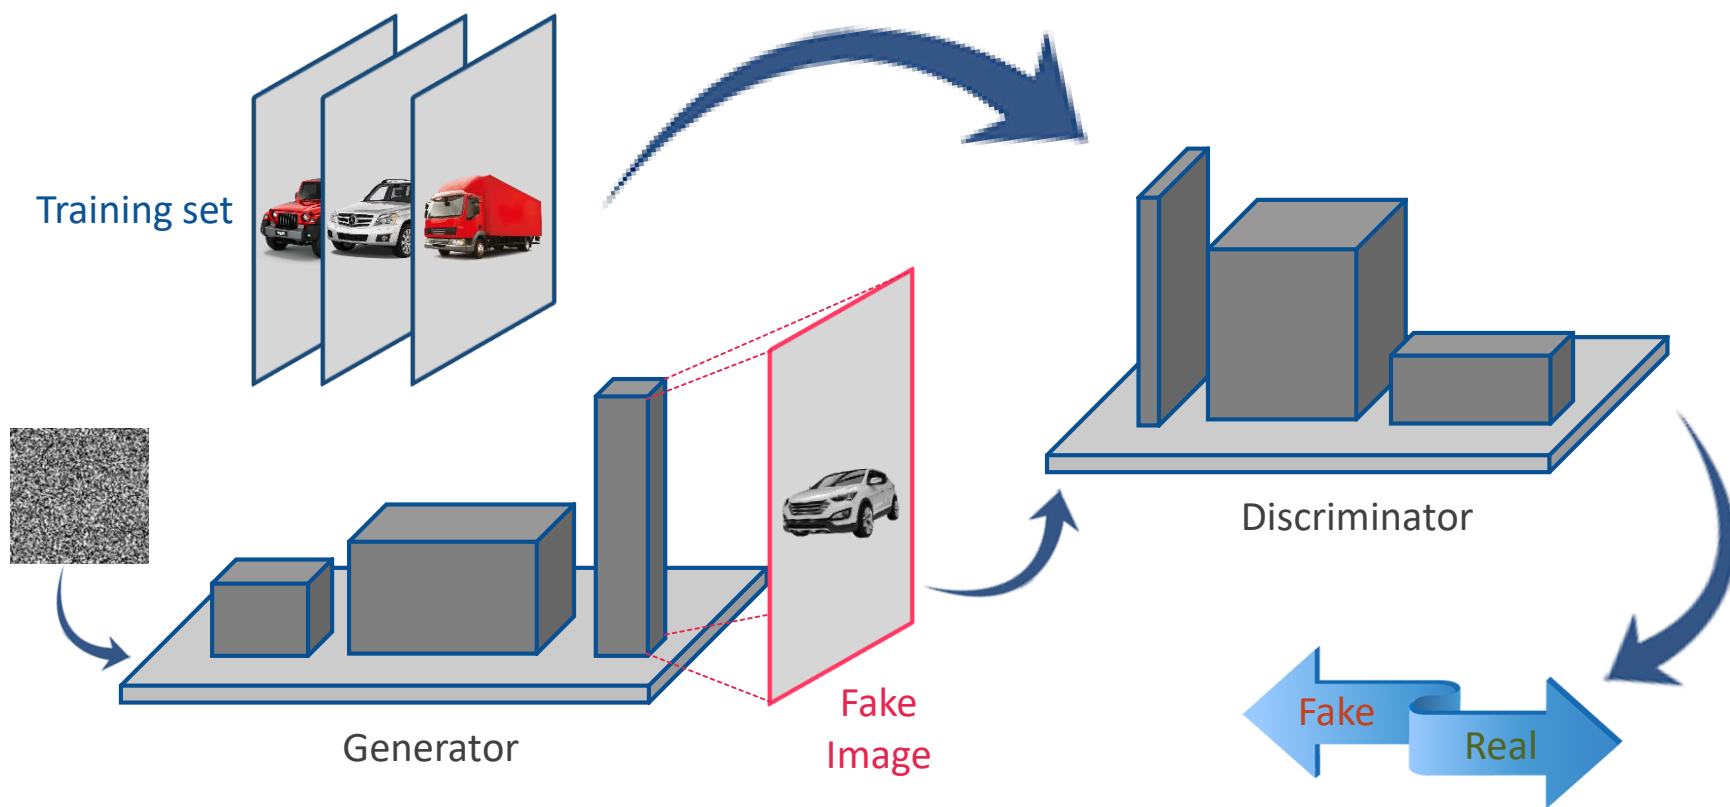
Generative Adversarial Networks



Important notes on GANs

- ▶ The output of a GAN can be a series of generated data that lies in the learned data distribution
- ▶ There are many different architectures for GANs:
 - ▶ **Conditional GAN** is a variant of GANs that can enable controlling the type of output using a conditioning factor
 - ▶ **CycleGAN** is another variant that learns a mapping for translation into another domain

Generative Adversarial Networks



Generative Adversarial Networks

Applications of GANs – Human Face Generation



Generative Adversarial Networks

Applications of GANs – Human Face Generation



Generative Adversarial Networks

Applications of GANs – Face Aging

21-30 yrs



31-40 yrs



41-50 yrs



51-60 yrs



61-70 yrs



71-80 yrs



Generative Adversarial Networks

Applications of GANs – Text-to-Image Translation





Generative Adversarial Networks

Applications of GANs – 3D Object Generation



Generative Adversarial Networks

Applications of GANs – Image Completion



Generative Adversarial Networks

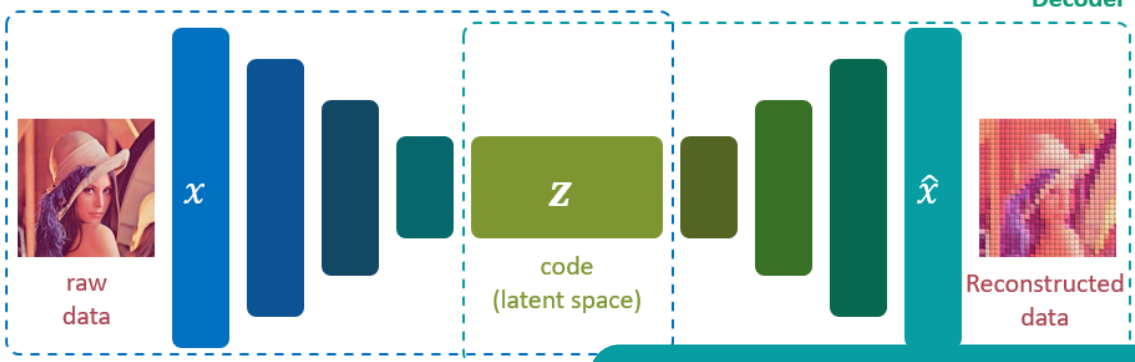
☰ README.md

✎

Autoencoders (AEs) and Variational Autoencoders (VAEs)

Autoencoders are deep neural network architectures used for **Feature (representation) Learning** in unlabeled data. The goal in an AE is to leverage NNs for representation learning. Autoencoders reconstruct the input approximately by keeping only the most relevant parts of data.

Encoder



x

raw data


z

code
(latent space)

\hat{x}


Reconstructed data

Decoder

 Codes

| # | File | Description |
|---|------------------------------------|--|
| 0 | Basic introduction | Introduction to Autoencoder codes in Keras |

Full code on GitHub



References

- ▶ <http://introtodeeplearning.com/>
- ▶ <https://www.guru99.com/supervised-vs-unsupervised-learning.html>
- ▶ <https://machinelearningmastery.com/what-are-generative-adversarial-networks-gans/>
- ▶ <https://openai.com/blog/generative-models/>
- ▶ <https://ermongroup.github.io/cs228-notes/learning/latent/>
- ▶ <https://www.jeremyjordan.me/autoencoders/>
- ▶ <https://machinelearningmastery.com/what-are-generative-adversarial-networks-gans/>
- ▶ <https://towardsdatascience.com/understanding-generative-adversarial-networks-gans-cd6e4651a29>

Questions?

